Exposure and Recovery with Persistent Randomness Reservations

Ziang Li[1*]

[1]Department of Engineering, Durham University, Durham, United Kingdom.

Corresponding author(s). E-mail(s): 2403871950@qq.com;

Abstract

Protocols that retain randomness across crashes must distinguish the creation of a secret source from the later computation that uses it. We study how persistent randomness reservations affect recovery after compromise. Our model assigns immutable origins to operations and tracks acquisition of their encrypted representations and disclosure of the keys that open them. Alias normalization preserves those origins, and an exact guarded-path theorem characterizes record exposure for fixed public key derivations. We give constructive simulations for private records and dependent authentication reports, together with candidate-local accounting for an adaptively selected compromise cut. These results connect original-operation exposure to computational confidentiality and selected-root authentication bounds for the next rekey of an established session and a selected first message. The reduction follows the complete hybrid contribution, root extraction, confirmation and traffic schedule in the stated quantum random-oracle models, retaining real verification and successor-state disclosure. An online finite-chain composition extends native confidentiality and message integrity to a target chosen after earlier-key disclosure. A shared-XOF bridge reduces the native signing-reservation interface to ordinary hedged ML-DSA unforgeability and supplies the source-attribution bound. Finite descriptor families also yield a purpose-key derivation result with exact aliases and separate owner lifecycles. A mechanized symbolic theorem establishes multi-epoch message secrecy under the provenance condition. Native executions demonstrate traffic reconstruction from pre-restart reservations, aggregate copies and delayed receipts. Independent finite checks validate the exposure semantics. Pinned reference workloads and 49 matched native/reference executions quantify reconstructed message sets jointly with admitted service, communication and persistence work.

Keywords: Post-compromise security, Persistent state, Randomness reservations, Ratcheting protocols, Record exposure, Provable security, Post-quantum cryptography

1 Introduction

Post-compromise recovery depends on the secrecy of contributions incorporated into later keys. Ratcheting protocols make this dependence precise through the evolution of key-exchange and messaging state [10, 14]. In a persistent implementation, a contribution also has a storage history: its randomness is sampled, a replay record is saved, the cryptographic operation is executed, and its result is published and confirmed. These events can be separated by a crash or a delayed message. A security experiment must account for the secret source across this entire history.

The reservation boundary.

Consider an endpoint that durably reserves the random tape of a KEM operation so that an interrupted call can resume with its original inputs. An adversary acquires this reservation, its wrapping key, and the predecessor root before the key-generation call executes. After a crash, the endpoint resumes the call and completes an honestly authenticated exchange. The captured tape determines the same private key; the public response then determines the KEM contribution and, with the predecessor root, the descendant traffic keys. The resumed computation is a later event in the lifetime of the original secret source.

Future-key exposure already has a precise place in messaging security. SCKA's vulnerable message sets include retained future-epoch keys [2], and Triple Ratchet accounts for next-epoch material sampled while its transmission is in progress [14]. The question here concerns the persistent objects from which such material can be recovered: *which facts about an original operation and the acquisition of its records must a recovery experiment preserve through restart, copying, and later key disclosure?*

There are two related difficulties. First, one reservation may occur in a pending target, a current journal image, and several retained copies. A later wrapping-key disclosure can make an earlier capture readable, even after the honest endpoint has deleted its local copy. Second, a cryptographic reduction may possess a selected operation's public outputs without possessing its random tape. It must still account for public record commitments, complete-record equality, exact retries, reached persistence states, and authenticated reports whose inputs depend on retained private values. These are semantic and simulation obligations of the recovery argument.

Original sources and durable executions.

We associate each reservation with its original owner, operation, session, epoch, and sampling event. We separately identify an immutable session-record version and a physical image carrying that version. Copying or resealing an image preserves the referenced source and record version. Acquisition records the objects obtained by the adversary; readability is evaluated against the complete acquisition and key-disclosure history. This representation makes the secrecy premise stable under changes to physical storage representation while allowing adversarial knowledge to grow.

We study this representation in Q-Periapt Continuity, a hybrid post-quantum messaging construction with persistent rekey operations and asynchronous confirmation. The construction provides a concrete setting in which the reservation, execution, and

confirmation events have distinct operational boundaries. Its journal can also aggregate several sessions, making the distinction between a new image and an existing source directly observable.

Contributions.

The paper establishes an exposure condition that follows original secret sources through persistent execution, and carries that condition into computational message security and symbolic multi-epoch recovery.

We give a provenance-preserving model of reservations and record acquisition. Alias normalization relates copied and resealed representations to the same original operation, preserving public views and the ledger of original records (Lemma 4.3). For fixed public key derivations, a guarded-path characterization exactly matches iterative decryption (Theorem 4.4). Its conjunctive guards express how captured ciphertexts and disclosed keys jointly reveal a source; alternative copies provide further routes to that source. This gives recovery experiments a stable exposure condition across physical changes to storage.

The computational connection constructs the retained private state needed by a reduction that holds an operation's public outputs. Images and intents preserve original record identities; dependent immutable reports are simulated in reverse causal order. Candidate-local accounting then localizes the original native score before the compromise cut is selected. An original-operation adapter carries this execution through the operational ML-KEM experiment, complete ContextBound tuple, root initialization, confirmation and traffic derivation. The resulting bounds establish confidentiality for a fresh marked first message after the next rekey of an established session (Corollary 5.17) and confirmation and first-message integrity under its selected recovered root (Theorem 5.25). They retain real verification, declared successor-state disclosures and explicit record, component, hash, key-schedule and AEAD costs.

An online finite-chain comparison extends this connection to a target chosen after earlier message-key disclosure (Theorem 5.12). Successive Extract and Expand comparisons isolate the selected key while preserving its disclosed successor seed and real suffix computation. The selection occurs before that step's first private computation across all endpoint aliases, including failed receive-side advances. The native composition retains the accepted send's original sources and candidate-local record stops, giving confidentiality and selected-message integrity with explicit lifetime and whole-execution budgets (Theorem 5.27).

The accepted allocation is associated with its original KEM calls by registering their existing reservation slots before private computation. A shared-XOF comparison preserves the signing randomizer and original-message ledger through replay and failed publication, reducing the native ML-DSA interface to ordinary hedged-signature unforgeability (Theorem 5.24). This instantiates the signature term in the source-attribution bound. The key-entry method also handles finite salt and information-string families, yielding a purpose-key instance with equal-input aliases and separate returned-owner lifecycles (Section 5.15).

The same provenance condition supports symbolic message secrecy when the target's original reservations follow the compromise cut and the stated authentication,

entropy and record-acquisition conditions hold (Theorem 6.1). The Tamarin model includes inherited roots, shared key identities, alternating proposers, asynchronous confirmation, and unbounded epochs. Native process-kill executions reconstruct traffic from reservations captured before restart; a delayed-receipt execution exposes epoch-two traffic while the peers' confirmed counters remain zero and one. A sibling-session write also carries an earlier target reservation into a new aggregate image. Together these executions show why the recovery condition follows original sources and acquisition histories across both persistence and asynchronous progress.

Supporting validation.

An executable finite-model supplement checks the guarded characterization and its treatment of acquisition, deletion, and late disclosure against independent expected results. A source correspondence maps the image, intent, retry, report, and acquisition operations to the declared interfaces. Separately, pinned SPQR and whole-KEM reference corpora measure reconstructible message sets across communication schedules. A further 49 native/reference executions share a list of directed transmission opportunities and charge control traffic and exact retries to that common budget. They measure reconstructed message sets jointly with first commitments, logical deliveries, emitted bytes and native persistence work. The measurements retain the originating contribution, direction, epoch and message index, together with each construction's authentication and storage interface. A one-way control records exhausted-budget suspension, its uncommitted requests and the exposure of the applications actually admitted. A paired 28-run native ablation then sets the minimum application-retry age to four public sender turns. It increases completed service in every measured bidirectional or resumed schedule while retaining the alternating current-epoch exposure counts; the complete accounting records the schedule-dependent communication cost.

Security model and organization.

Crash recovery resumes retained, bound operations under the authentication, acquisition and disclosure interface of Section 4. The symbolic theorem covers unbounded epochs. The computational results connect the stated finite record, report and online message interfaces within the common message experiment, retaining their explicit cryptographic and disclosure terms. Section 3 describes the construction; Sections 4 and 5 develop the record model and computational connection. Section 6 presents the mechanized recovery theorem and separating executions, and Section 7 validates the exposure semantics and presents the matched executions. Section 8 develops their consequences and further instantiation.

2 Related work

Ratcheting and fine-grained exposure.

Cohn-Gordon et al. model Signal's key-exchange core as a branching multi-stage authenticated key exchange and prove security for that abstraction [10]. Auerbach et al. develop SCKA with vulnerable epoch sets and derive secure messaging relative to direction- and index-sensitive vulnerable message sets [2]. Their Definition 3.2

includes current chains, already produced future-epoch keys, and pending out-of-order messages; Theorem 3.3 gives the corresponding messaging composition. The object studied here is the history of persistent records from which a contribution can be recovered. Original-reservation provenance connects that history to exposure of the underlying cryptographic operation.

Triple Ratchet combines classical and post-quantum key agreement while distributing larger post-quantum messages across transmissions [14]. Its Remark 4.3 accounts for exposure of next-epoch material already sampled during distribution. This is the relevant comparison for our delayed-receipt execution. Our additional distinction is between the original source, its durable representations, and the later computation that uses it. Each representation may have its own acquisition time and wrapping keys while referring to the same reserved operation.

Persistence, exact replay, and crash recovery.

The Signal specification discusses deferring ratchet-key generation to shorten the lifetime of sensitive state [24, Section 8.5]. A persistent reservation can fix that state before the deferred call runs. The original-reservation event therefore identifies the lifetime relevant to our freshness condition. Exact replay and a fresh sealing nonce preserve that source identity even when they change the operational or physical representation.

Crash consistency addresses the state recovered after an interrupted update. For example, GoJournal provides verified atomicity for storage applications through concurrent crash specifications [8]. This storage property explains which retained operation an endpoint resumes. Our exposure interface additionally records which copies an adversary acquires and which key disclosures enable their contents. The native reservation executions illustrate the interaction: an operation resumes with its original inputs, and those inputs can also determine the adversary's later derivation. The guarded normalization theorem expresses that relation for the stated encrypted-record fragment.

Signing-randomness leakage.

Damm et al. recover ML-DSA secret-key material from noisy bit leakage of the internal masking vector y [13]. The retained signing record studied here contains the external 32-byte randomizer `rnd`. FIPS 204 derives y through a secret-key-dependent seed from the key seed, `rnd` and the message representative [23, Algorithm 7]. These are distinct observation interfaces. Section 5.10 specifies the corresponding reservation-aware signature experiment and uses it to associate an accepted rekey with its original KEM calls. Section 5.11 then connects this interface to ordinary native hedged-signature security.

Hedging and exposed nonces.

Nonce-based cryptography treats nonce handling as an explicit security interface; Bellare and Tackmann develop this perspective and introduce hedged extractors [4]. Aranha et al. analyze hedged Fiat–Shamir signatures under fault attacks [1], and Grilo et al. give quantum random-oracle analyses of hedged Fiat–Shamir security against nonce and fault attacks [17]. Kosuge and Xagawa study the fault security of ML-DSA

[20]. Our signing connection uses this hedging perspective for persistent reservations: opening an original randomizer, selecting a reserved call, exact replay, and failed publication preserve their original message ledger. The native reduction compares the shared key-generation and hedging XOF with a private random-function interface and then returns to ordinary hedged ML-DSA security (Theorem 5.24).

Session handling and state loss.

Cremers, Jacomme, and Naska analyze how the session-handling layer determines conversation security and give a Tamarin analysis of Signal's Sesame [11]. Cremers, Medinger, and Naska prove incompatibilities between full end-user PCS and specified state-loss resilience requirements [12]. These results make the surviving state and authentication resources essential parts of a recovery claim. Our failure model retains the bound journal and resumes its original operations. Its disclosure conditions follow actual key identities, including a key shared across sessions, and acquisition of the target projection inside an aggregate record. This specifies the persistence and authentication setting of the symbolic recovery theorem.

Admissible network behavior.

Fischlin, Günther, and Janson formalize robustness for channels over unreliable networks and analyze the record layers of QUIC and DTLS 1.3 [16]. Their correctness and security notions account for supported delivery histories. Our model likewise retains delayed receipts and asymmetric progress as ordinary executions. The resulting provenance question concerns the secret inputs retained along such a history. In particular, the delayed-receipt witness is evaluated using both endpoints' state and the captured reservation, preserving their separate confirmed-epoch counters.

Mechanized ratchet analyses.

Cheval, Jacomme, and Richards give an automated Double Ratchet analysis covering out-of-order messages and key-compromise scenarios, including PCS results and forward-secrecy attacks and fixes [9]. Hérouard et al. mechanize computational PCS for a two-party asymmetric ratchet with unbounded steps and a fixed non-adaptive compromise scenario [18]. Their proof distinguishes non-deduction of intermediate roots from real-or-random security of derived keys. Our symbolic recovery theorem and quantitative interface theorems retain their respective conclusion types. The private-record simulator addresses the separate task of executing a specified storage interface when selected operation bytes are unavailable to a reduction.

Bhargavan et al. study SPQR and its implementation [7]. The authors' public account describes ProVerif protocol models and hax/F* verification of panic-freedom and selected correctness properties of the Rust core [6]. These developments motivate comparison at the level of explicit state and observation interfaces. Our source correspondence records the native constructors and transitions represented by the original-record model. The SPQR reference corpus supplies independently identified executable measurements of retained-secret exposure and communication; its observations keep their source revision and predictor interface.

Key extraction and dependent uses.

HKDF-Extract places its input key material in HMAC's message argument [21]. In the root update studied here, this input remains the secret contribution when the predecessor root is disclosed. Backendal et al. analyze HMAC's swapped PRF use and characterize its feasible input domains [3]. This is the relevant primitive-level distinction for the root-extraction step. Our computational accounting places that use alongside the record, report, confirmation, and traffic interfaces of the same recovery experiment.

3 Construction and the motivating execution

Q-Periapt Continuity combines ML-KEM-768 and X25519 in its ContextBound hybrid KEM and authenticates protocol controls with an ML-DSA-65/P-256 identity. Our analysis follows the persistent session and rekey machinery: an original operation first reserves its entropy, places the resulting token in a retained plan, and later computes or replays the corresponding KEM operation. An exact retry uses the same reservation. The protocol alternates the proposing endpoint across epochs and carries separate send, receive, and confirmed-epoch state.

Fixed cryptographic context.

The combiner hashes a length-delimited encoding of the two component secrets, both public keys and ciphertexts, and domain, suite, policy and application-context fields. On the rekey path, the application context is a domain-separated digest of the complete retained offer wire and is nested with the admitted policy digest. These fields form fixed inputs of the original operation; exact replay retains that binding and its reserved tape. The source correspondence in Appendix A.7 identifies the encoding and reservation constructors used here.

3.1 Reservations and persistent representations

A generation reservation determines the future hybrid key pair. An encapsulation reservation, together with the authenticated offer, determines the corresponding ciphertext and contribution. The implementation seals these reservations under the host wrapping key. It persists the selected target in a pending write before installing that target as the next journal image. Applying or reopening that write preserves its original contents.

A journal image aggregates state from multiple sessions. Updating a sibling session advances the global image revision and samples a new outer encryption nonce, while the target session's retained reservation can remain unchanged. Consequently, the journal revision, outer nonce, local operation identifier, and original entropy reservation represent distinct identities. We preserve these distinctions in the model and relate them to the native execution boundaries.

Event	Effect on the target operation	Provenance fact
Reserve	Draw KEM entropy and seal its replay token	Creates an original source
Persist target	Retain the original pending-write bytes	Creates an acquisition opportunity
Apply or retry	Install or replay the retained target	Preserves the source
Sibling write	Re-encode the aggregate journal with a new outer nonce	Preserves unchanged target sources
Confirm	Advance the relevant endpoint's protocol state	Records protocol progress

Table 1 The events tracked by the analysis. Several events may occur in one symbolic transition; the native execution exposes their separate boundaries.

3.2 Root and traffic dependencies

Write R_{e-1} for the preceding root, S_e for the hybrid contribution, C_e for the public response core, and O_e, P_e for the complete signed offer and response. With the construction's canonical encodings and domain-separated transcript hashes, its dependency equations are

$$R_e = \mathrm{HKDF}_{\mathrm{SHA256}}\big(R_{e-1}, S_e, d_R \| H(d_C, C_e), 32\big), \tag{1}$$

$$T_e = \mathrm{HKDF}_{\mathrm{SHA256}}\big(sid, R_e, d_T \| ctx \| e \| H(d_{\mathrm{tr}}, O_e \| P_e), 128\big). \tag{2}$$

Here HKDF arguments denote salt, input key material, information, and output length in bytes. T_e supplies two directional traffic chains and two acknowledgement keys. The actual domain strings and encodings are pinned by the source locations in Appendix A.

These equations describe a concrete derivation path. A captured generation reservation and matching wrapping key reconstruct the private key; a subsequent public ciphertext then yields S_e. A captured encapsulation reservation yields the same contribution using its retained offer. Combining S_e with a disclosed R_{e-1} and the public control wires reproduces R_e and the traffic material. Honest admission checks can still reject an operation inside the SDK while an adversary holding these inputs performs the underlying computation independently.

3.3 Asynchronous confirmation and a later exposed epoch

Consider the native execution in Figure 1. Endpoint B commits the first receipt, advances locally, and prepares an epoch-two offer before A receives that receipt. At the capture point their local observations are

$$A : (c, s, r, p) = (0, 1, 0, 1), \qquad B : (c, s, r, p) = (1, 1, 1, 2),$$

where c, s, r are confirmed, send, and receive epochs, and p is the pending target epoch. The snapshot of B contains R_1 and the original generation reservation for epoch two.

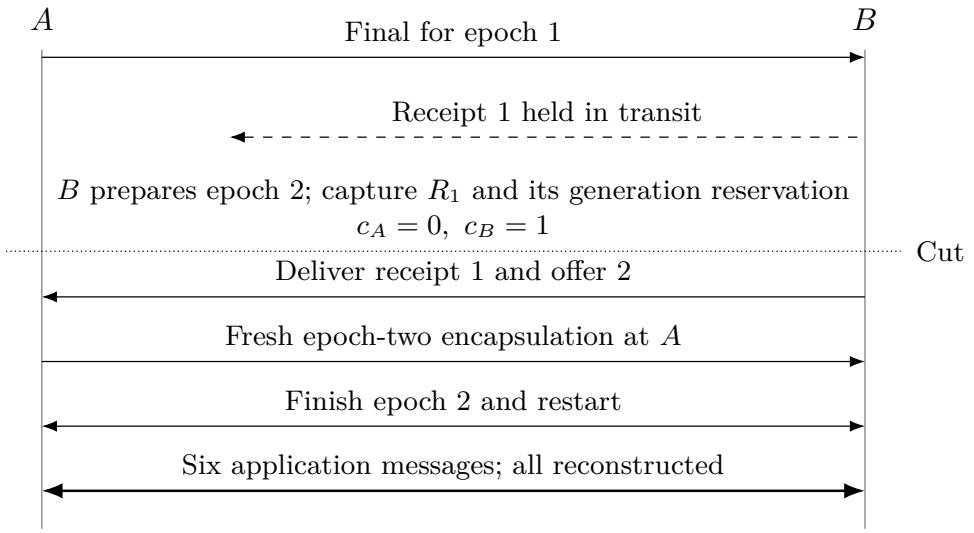


Fig. 1 A retained native disclosure execution. The smaller confirmed counter is zero at capture, while the captured reservation enables epoch-two traffic recovery after the public exchange. The dotted Cut orders the exposure and the subsequent fresh encapsulation; it does not refresh the earlier generation reservation.

After delivery resumes, those captured inputs and the public exchange reconstruct epoch-two traffic.

The resulting counterexample targets the proposed bound $\min(c_A, c_B) + 1$: its value is one, yet epoch-two traffic is exposed. The protocol permits this asymmetric progress. Recovery reasoning therefore follows the retained inputs and their descendants, with each endpoint's confirmation counter retaining its own value.

4 Reservation origins, persistent records, and recovery

An interrupted cryptographic operation can retain its random tape before it has produced a key or a ciphertext. Its later execution and confirmation therefore have to be related to the original reservation. We make this relationship explicit at two levels: the origin of a cryptographic operation, and the origin of a saved session-state version. A third identifier describes the physical image that carries that version. This distinction lets us account for a target record copied by another session's write, an unapplied pending image, and a ciphertext acquired before its wrapping key.

The model is an interface model for an asynchronous protocol machine $\mathcal{P}$. The adversary schedules public messages, process loss, honest recovery, and the disclosure operations specified below. Application deliveries are private to the honest applications. Capture grants read access to the stated assets; it does not grant an additional operation for replacing the authoritative journal. An implementation with adversarial rollback or forking requires those transitions in its storage machine. The protocol, its cryptographic primitives, and its public outputs are parameters of the record-normalization lemma.

Table 2 Principal objects in the record model. The three identities separate an original random draw, an immutable session-state version and its physical storage representation.

Notation	Meaning
o, ρ_o	Original invocation and its reserved random tape
v, X_v	Immutable session-state version and its payload
h	Physical image or pending object carrying one or more versions
κ	Identity of actual key material, shared by equal keys
$\mathsf{res}(o)$	Event at which the original random tape is drawn
$\mathsf{save}(v)$	First retention of v on the durable acquisition surface
$\mathcal{C}_t, \mathcal{L}_t$	Acquired and readable original versions at prefix t

Fix finite execution, record-size, and query budgets; the wrapper's bookkeeping is polynomial in their total representation size.

Table 2 summarizes the principal objects.

4.1 Three origins and their event order

Definition 4.1 (Reservation origin) A reservation origin is a fresh invocation index o, together with

$$\big(\mathsf{kind}_o, \mathsf{owner}_o, s_o, e_o, B_o, \mathsf{res}(o)\big).$$

Here kind_o distinguishes generation, encapsulation, and any other reserved operation; s_o, e_o identify its session and epoch; B_o is the full canonical binding, including the admitted policy and operation inputs; and $\mathsf{res}(o)$ is the event at which its random tape ρ_o is drawn. Execution or exact replay of o consumes the same ρ_o and the same binding. A new draw receives a new origin even if a failed attempt's external operation identifier is reused.

The ideal source in this model draws independent tapes of the prescribed lengths. Explicit entropy-disclosure queries reveal designated tapes. Independence is a distributional assumption about the source at sampling. Later API calls preserve the sampled tapes. The generation and encapsulation budgets G, C count original draws, including draws in attempts that subsequently fail. Exact retries consume no new slot. Independently sampled operations retain distinct origins even if some public outputs happen to coincide.

A record version v has an immutable payload X_v and semantic metadata

$$\big(\mathsf{subject}(v), \mathsf{role}(v), \mathsf{epoch}(v), \mathsf{phase}(v)\big).$$

The subject identifies the session whose state is represented. A write to the enclosing journal preserves this subject, including a write by another session. A version can refer to several reservation origins and to roots or traffic state derived from earlier origins. Preparing a version allocates its handle. The event $\mathsf{Saved}(v)$ occurs when that version first enters the durable acquisition surface, and its time is denoted $\mathsf{save}(v)$. In particular, an intact target in a durable pending-write intent is saved before the corresponding image is applied. A storage operation may have made this transition even when its caller receives an unknown-outcome error.

A physical object h has a separate creation or sealing index. It may carry one version, several sessions' versions, an exact earlier image, or a pending target. An unchanged version keeps its v and $\mathsf{save}(v)$ when it is copied, archived, applied from a pending intent, or included in a newly sealed aggregate. A new semantic session-state version receives a new v even when some of its fields are unchanged. These rules distinguish the original reservation from each session-state update that retains it.

The following transition rules specify the provenance discipline.

1. *Reserve and execute.* Reserve creates o at the actual draw. Execution uses its retained binding and tape. Unknown-outcome reconciliation keeps that origin; a definite new attempt allocates another.
2. *Prepare and retain.* Preparing v fixes X_v and its provenance. Retaining the first intact image or pending target containing v emits $\mathsf{Saved}(v)$. An attempt that retains no such object has no successful acquisition event.
3. *Copy, apply, and reseal.* An exact copy preserves the complete object and its sealing origin. Resealing creates a new sealing invocation and its actual new wrapper randomness, while preserving the referenced record versions. Applying an existing pending target copies its retained bytes.
4. *Crash and resume.* Process loss removes volatile references. Recovery reconstructs the retained handles and their original provenance. It does not turn a retained reservation into a new draw.
5. *Confirm and delete.* Confirmation changes protocol state and may authorize particular host-side deletions. Neither action changes a reservation's origin or removes an adversary's previously acquired copy.

For an original reservation actually retained in v, causality requires

$$\mathsf{res}(o) \;\leq\; \mathsf{save}(v). \tag{3}$$

A later physical object's creation time gives an additional timestamp, not a replacement for either side of (3). The interval between drawing ρ_o and first retaining it remains explicit. The symbolic reservation-and-save transition can represent that interval only through a separately justified source correspondence.

4.2 Acquisition, delayed readability, and key identity

Wrapping-key handles resolve to identities of actual key material. Two session labels that designate the same key resolve to the same identity κ. Disclosure through either label therefore enables every acquired wrapper protected by κ. A rotation creates a different identity when it creates different key material. Signing-key identities are treated in the same way across sessions and roles.

For each intact physical object h, its construction records a finite set

$$\mathsf{paths}(h) = \{(v, \gamma)\},$$

where γ is the set of wrapping-key identities on a recoverable serialization path to X_v. Multiple paths to one version are permitted. An unprotected record has $\gamma = \varnothing$.

Nesting a wrapper adds its key identity; aggregation unions its members' paths. An exact ciphertext copy keeps its paths. Resealing decoded records reconstructs the paths under the new wrapper and retains the same v values. Thus resealing decoded plaintext and wrapping an old ciphertext are distinct operations.

Only intact, recoverably embedded record payloads create these paths. A public hash or MAC of a private record is a public computation of that record, not a path to the record itself. The renderer introduced below preserves that computation and its output. This distinction prevents a public head query from being relabeled as a snapshot merely because its computation used private bytes.

Definition 4.2 (Record acquisition and readability) A successful $\mathsf{Capture}(h)$ returns the actual retained bytes of h and adds an immutable copy to the adversary's archive $\mathcal{A}_t$. Let $\mathcal{K}_t$ be the wrapping-key identities disclosed by time t. Define

$$\mathcal{C}_t = \{v : \exists h \in \mathcal{A}_t,\ \exists \gamma,\ (v, \gamma) \in \mathsf{paths}(h)\}, \tag{4}$$

$$\mathcal{L}_t = \{v : \exists h \in \mathcal{A}_t,\ \exists \gamma,\ (v, \gamma) \in \mathsf{paths}(h)\ \wedge\ \gamma \subseteq \mathcal{K}_t\}. \tag{5}$$

Each capture records $\mathsf{AcquireOriginal}(v, t)$ for every version in its support, including versions already acquired through another path. An unsuccessful request for unavailable state returns its ordinary error and adds no acquisition.

Both sets describe the specified disclosure interfaces, using provenance and key identities. They do not test whether the adversary has guessed a secret. The read set can grow without another capture: if h was obtained at t_1 and its final missing wrapping key is disclosed at $t_2 > t_1$, the corresponding versions enter $\mathcal{L}_{t_2}$. Host deletion at an intermediate time leaves h in $\mathcal{A}_{t_2}$. Acquisition, by contrast, already occurred at t_1.

An additional key that is explicitly recoverable from an opened record can be included by a declared extraction rule. In this extension, iterate (5) and the declared key-extraction rules to their least fixed point after each event. Each rule identifies its input bytes, its deterministic extraction algorithm, and the resulting actual key identity. Only rules with available inputs fire. This accounts for nested protection without assuming that inaccessible keys have already been learned.

Signer-file capture is a separately typed disclosure. For correspondence with the retained symbolic theorem, it records exposure of the actual signing key conservatively at acquisition, as does that theorem's signer-file rule. The readability ledger can additionally record when the file's contents become decryptable. In particular, a wrapping-key disclosure cannot be scoped to traffic state when the declared capture also includes the matching signer file.

Constructive exposure.

From public protocol inputs, explicit leaks, and readable record payloads, we form a labelled replay-dependency graph. Its edges are specified algorithms: unseal a reservation with its matching key; regenerate a key from its original tape; replay encapsulation on the bound offer; decapsulate a published ciphertext; derive a root; and derive the direction-specific traffic and ACK material. Directly stored roots, skipped keys, and plaintext enter through their own record fields. For the implemented root schedule, an

edge is

$$(R_{e-1}, S_e, C_e) \longmapsto R_e = \mathsf{HKDF}\big(R_{e-1}, S_e, D_{\mathrm{root}} \| H_{\mathrm{core}}(C_e), 32\big).$$

Closing these explicit edges over the nodes of the bounded execution yields a computable set of message keys with replay witnesses. It supplies positive exposure certificates; the confidentiality game below is the separate criterion for messages outside the declared disclosure conditions. Confirmation counters belong to the public progress state and do not reset this graph.

4.3 An original-record implementation of the copy interface

We now isolate a property of record management that is independent of KEM secrecy. The following two storage machines use the same protocol machine $\mathcal{P}$, random-tape source, scheduler, and cryptographic interfaces.

The *literal-copy machine* materializes the immutable payload X_v in each indicated serialization position. Its objects include current images, pending targets, and retained historical copies. It records the origin tags alongside these values as private bookkeeping.

The *original-record machine* keeps one immutable table entry $\mathcal{O}[v] = X_v$ and represents each serialization position by a reference to v. Its wrapper maintains the physical-object directory, pending states, adversary archive, origin tags, and key identities. It has the following specified interfaces.

1. $\mathsf{RegisterOriginal}(v, X_v)$ is a private write by the honest machine. It fixes one value at a fresh v. The wrapper receives the handle and the provenance metadata.
2. $\mathsf{RenderOriginal}(\theta)$ evaluates the supplied serialization recipe θ by dereferencing $\mathcal{O}$. A recipe specifies literal public fields, record references, references to already rendered objects, and each sealing, hash, MAC, or other admitted byte-processing call. A fresh sealing call uses a fresh invocation index and the same prescribed coin distribution as the literal-copy machine. An exact-object read or replay returns its cached original bytes.
3. $\mathsf{AcquireOriginal}(v, t)$ records acquisition of that version. This is an event, not an assertion of decryption. A capture's externally returned bytes come from the rendered object cache.
4. $\mathsf{RevealOriginal}(v)$ returns the original payload X_v. The wrapper uses this interface for a plaintext record-opening response only after a captured path is enabled by the disclosed keys, and caches the answer for subsequent openings.

The renderer is an explicit trusted interface in this intermediate experiment. It has access to the original payload table and evaluates the stated byte operations, including operations on undisclosed payloads. Consequently, it can produce real ciphertexts, commitment digests, and report MACs. This definition does not supply an algorithm for producing those bytes from an unknown KEM-challenge tape. It gives a precise location for the subsequent cryptographic replacement of that interface.

The operational contract has three further requirements. First, all publicly visible payload-dependent bytes and statuses are obtained through the specified renderer or through $\mathcal{P}$'s common protocol computation. Memory addresses and host timing are

not additional outputs of this interface. Second, actual byte-operation order, error responses, partial transaction outcomes, and cryptographic-oracle calls are retained. Caching an exact object does not discard a native hash-validation call, observable retry, or write attempt: such an action remains a renderer or protocol action. Third, the reference metadata and the key-identity relation are available to the wrapper without inspecting a selected hidden reservation's coins. These are checkable interface conditions; none asserts the secrecy of a record, a root, or a message.

Lemma 4.3 (Alias normalization) *For the operations and renderer contract above, every adversary $\mathcal{A}$ interacting with the literal-copy machine has an online wrapper $\mathcal{B}$ interacting with the original-record machine such that their public views have the same distribution. Their reservation-origin ledgers, record-save events, acquisition events, and read sets are identical after projection onto the original handles.*

The wrapper makes at most $|\mathcal{L}_T|$ distinct RevealOriginal *calls by terminal time T. Repeated physical copies, captures, and openings do not multiply this number. The wrapper preserves every original reservation and every fresh sealing invocation, so it does not merge their respective cryptographic resource counts.*

The statement also holds for an adversary with quantum private state and quantum access to common public cryptographic oracles, provided the storage interface is classical and both experiments retain the same oracle calls.

Proof Maintain a map from every literal record occurrence to its immutable handle v, and from every native object identifier to a rendered-object entry. Couple the protocol random tapes, sealing coins, scheduler choices, and public cryptographic oracles between the experiments. The invariant is that dereferencing every original-record recipe produces the corresponding literal-copy bytes, with the same semantic provenance and physical disposition.

For a new record, register its payload once and extend the occurrence map. For a fresh rendering, submit the same recipe and byte-operation sequence to RenderOriginal. Dereferencing gives the same inputs to every sealing, digest, MAC, validation, and output operation. Coupled coins and common oracles therefore give the same bytes, statuses, and state changes. This step retains a write that became durable before returning an error; the caller's status does not overwrite the persistence state.

An exact copy or pending-target application copies the object reference and its original cache entry. Aggregation collects the corresponding original references, and resealing evaluates the new recipe using those references. Each operation preserves the session-record handles and their save times. It follows by induction on the recipe constructors that $\mathsf{paths}(h)$ is the same in both machines. A public hash or MAC constructor contributes its common output and has an empty record-acquisition path set.

On successful capture, output the cached object bytes in both experiments and append its immutable descriptor to the adversary archive. For each version in its support, emit the same acquisition event at the actual capture time. Preserve the ordinary error for an unavailable object. When a matching wrapping key is disclosed, use the common key identity and reevaluate the enabled paths in the complete adversary archive. For an opening response, obtain X_v from RevealOriginal on its first use and from the wrapper's cache thereafter. Every such call has $v \in \mathcal{L}_T$. Shared-key aliases and later key acquisition are thus handled without another copy of X_v. The optional key-extraction extension preserves the invariant by induction on the least-fixed-point rounds.

Host deletion removes only host references. The adversary archive and already returned payloads persist. Crash and resume use the same retained object identifiers; neither operation recreates a reservation or changes a save event. Ordinary protocol actions are performed by the same $\mathcal{P}$ with the same results. This completes the induction, including adaptive choices based on any preceding response.

For quantum private state, identical classical responses and the same public oracle transformations induce the same transformation of that state at each interaction. The argument introduces no measurement of a quantum oracle query and no oracle-dependent advice. Finally, caching each plaintext opening by v proves the query bound; reservation and fresh-seal indices were preserved separately throughout. □

The lemma places multiplicity in the correct ledger. An adversary may retain many ciphertext aliases of one version, and those ciphertexts may have many different public heads. All such byte-level observations remain in the view. The repeated payload uses one original-record opening, whereas each genuinely fresh seal and each public hash or MAC computation keeps its own resource accounting.

4.4 Explicit guarded normalization of record readability

The path and readability definitions above already account for copying and later key disclosure. We now give their explicit guarded normalization and an executable realization for a record-decryption specialization. Here key derivation has fixed public edges independent of recovered plaintext. The optional plaintext-dependent key-extraction extension remains as defined above and is outside this specialization. The set Ω_{src} below contains reservation-source identities; it is distinct from saved record versions and from the original-payload table $\mathcal{O}[v]$. Likewise, $\mathcal{E}_t^{\mathrm{src}}$ projects derivability onto protected source values, whereas $\mathcal{C}_t$ and $\mathcal{L}_t$ track acquired and readable versions. The graph is a fixed representation of admitted retained objects, not a replacement for the provenance transition rules or the byte renderer.

A source identifier names one original sampling operation. An acquisition history records what the adversary obtains, including objects obtained at different times. These are distinct objects of the model: a copy preserves the former, while an additional acquisition can change the latter.

Immutable record graph.

Let $\mathcal{R}$ be a finite set of immutable encrypted record objects. Each $r \in \mathcal{R}$ has a matching decryption-key identity $\kappa(r)$, a set of embedded child records $\mathrm{child}(r)$, and a set $\mathrm{orig}(r)$ of protected source values in its plaintext. Containment is acyclic. Record identities distinguish ciphertext objects; source identities distinguish original sampling operations. Several records can contain the same source. A rewrapped record has a new record identity and may have a new key, while retaining its source identity. Key identities are equivalence classes of equal actual keys. Equal key bytes therefore have one identity across separate API objects.

Let $\mathcal{D} \subseteq \mathcal{K} \times \mathcal{K}$ be fixed public key derivation edges. An edge (k, k') means that possession of k permits deterministic computation of k'. This interface represents, for example, derivation of an operation-wrapping key from the journal key. In the theorem

below, these edges do not depend on plaintext recovered from records. Write D^* for their reflexive transitive closure.

Let $\Omega_{\mathrm{src}} = \bigcup_{r \in \mathcal{R}} \mathrm{orig}(r)$ be the declared source domain. At a trace prefix t, let $A_t \subseteq \mathcal{R}$ be all record objects ever acquired by that prefix, $K_t \subseteq \mathcal{K}$ all directly disclosed keys, and $\mathcal{S}_t^{\mathrm{dir}} \subseteq \Omega_{\mathrm{src}}$ all directly disclosed source values, represented by their source identities. These are cumulative sets. The decryption-key closure is

$$K_t^* = \{k' : \exists k \in K_t,\ (k, k') \in D^*\}.$$

The concrete derivation rules are

$$\frac{k \quad (k,k') \in \mathcal{D}}{k'}, \qquad \frac{r \quad \kappa(r)}{c}\ (c \in \mathrm{child}(r)), \qquad \frac{r \quad \kappa(r)}{o}\ (o \in \mathrm{orig}(r)).$$

Their least closure starts from $A_t, K_t, \mathcal{S}_t^{\mathrm{dir}}$. A fact o in the source sort means that the protected value of source o is derivable; it does not mean that a public identifier has become known. The source projection of this closure is denoted $\mathcal{E}_t^{\mathrm{src}}$.

Guarded routes.

For every path $p = (r_0, \ldots, r_\ell)$ in the containment graph ending at a record with $o \in \mathrm{orig}(r_\ell)$, define

$$g(p) = \{\kappa(r_0), \ldots, \kappa(r_\ell)\}.$$

The normalized table contains a route $(r_0, g(p), o, p)$ for every such path. Its exposure expression is

$$\widehat{\mathcal{E}}_t^{\mathrm{src}} = \mathcal{S}_t^{\mathrm{dir}} \ \cup\ \{o : \exists (r, g, o, p),\ r \in A_t \ \wedge\ g \subseteq K_t^*\}. \tag{GP}$$

Alternative records and paths are disjoined, while all keys on one path are conjoined. Equal key identities occur only once in the guard set.

Theorem 4.4 (Exact guarded normalization) *For every finite record graph and derivation interface above, and every acquisition prefix $(A_t, K_t, \mathcal{S}_t^{\mathrm{dir}})$, the concrete and normalized exposure sets agree: $\mathcal{E}_t^{\mathrm{src}} = \widehat{\mathcal{E}}_t^{\mathrm{src}}$. The equality applies to all prefixes of an adaptively chosen history.*

Proof Key derivations depend only on disclosed keys and the fixed relation $\mathcal{D}$, so their least closure is exactly K_t^*. For the forward inclusion, take a source $o \in \mathcal{E}_t^{\mathrm{src}}$. If it was directly disclosed then $o \in \mathcal{S}_t^{\mathrm{dir}}$. Otherwise its first derivation decrypts a record r_ℓ with the matching key. If that record was not acquired directly, its first availability follows by decrypting a parent. Follow these first derivations backwards. The finite acyclic record graph makes this process terminate at an acquired record r_0. Every intervening key lies in K_t^*, yielding a route satisfying (GP). For the reverse inclusion, take a satisfied route. Starting with its acquired first record, successively decrypt along the path. Every required key is in K_t^*, so each child and finally o is in the concrete least closure. Direct disclosures are in that closure by definition. Both inclusions hold for each fixed prefix; no independence or advance choice of the acquisition schedule was used. □

Efficient evaluation.

Equation (GP) is a declarative normal form: a compressed record DAG can have exponentially many paths. An efficient leakage evaluator instead computes the iterative decryption and key-derivation fixed point, without enumerating those paths. Let N be the total number of record, key and source atoms, and let M count rule and payload incidences. There are at most N strict-growth rounds; repeated rule scans therefore take $O(N(N + M))$ time and $O(N + M)$ space under standard indexed set and dictionary representations, with source-dictionary operations included in this accounting. The finite checker materializes the route table only for its declared small fixtures, to compare that normal form against the forward evaluator.

Corollary 4.5 (Copies, deletion and late disclosure) *Exposure is monotone under additional acquisitions, direct disclosures and key disclosures. Copying or rewrapping a source adds routes to the same source identity. Local deletion of a record previously acquired does not remove its routes from adversarial knowledge. In particular, a later matching-key disclosure can expose a source in a ciphertext acquired before that disclosure and deleted locally in the meantime.*

Proof The cumulative sets and their key closure only grow, and every predicate in (GP) is positive. A source-preserving copy changes a route's record identity or key guards but not its terminal source. Local deletion restricts future acquisition from that local store; it removes no member of A_t and therefore cannot invalidate a satisfied route. □

Interface significance.

The theorem gives an exact finite normalization of the chosen record-decryption semantics. It supports a precise original-operation leakage oracle: a query for source o can be evaluated against the complete acquisition history, with a path and its key guards as a witness. The cryptographic transition from native ciphertexts and public commitments to this interface has a separate simulation obligation. The derivation rules do not model guessing, cryptanalysis, timing leakage or a quantum decryption oracle. Accordingly, absence from $\mathcal{E}_t^{\mathrm{src}}$ is recorded as non-derivability by these rules; message indistinguishability outside an exposure set is supplied by a security theorem for the surrounding protocol.

Native shape and finite checks.

At the fixed research source, operation reservations are encrypted under a wrapping key derived from the journal key, and tokens are retained within encrypted durable images or pending targets. The corresponding route therefore includes the image key and the token key, with their actual derivation edge. A fresh aggregate image can preserve an older token and its original source. Public image-head digests are distinct observations from the record objects needed by the decryption rules. The executable supplement independently evaluates iterative decryption and guarded-route normalization, exhausts all acquisition/key/direct-source subsets for its declared finite graphs, and checks histories with late key disclosure and local deletion. Its counts and

source hash are in the accompanying machine-readable result; those checks validate the implemented finite model and supplement the general proof above.

Why exposure must retain conjunctive guards.

The following elementary separation explains the conjunction in (GP).

Proposition 4.6 (Acquisition exposure is not additive) *Let r be a retained-record atom, k a wrapping-key atom, and s a protected-seed atom associated with an origin identifier o. For $A \subseteq \{r, k\}$, let $\mathrm{Cl}(A)$ be the least set containing A and closed under the rule*

$$\{r, k\} \implies s.$$

There are no other rules in this fragment. Define the exposed-origin set by

$$E(A) = \{\, o : s \in \mathrm{Cl}(A) \,\}.$$

There is no fixed function $f : \{r, k\} \longrightarrow \mathcal{P}(\{o\})$ satisfying

$$E(A) = \bigcup_{x \in A} f(x) \qquad \text{for every } A \subseteq \{r, k\}.$$

Proof The closure rule gives

$$E(\{r\}) = E(\{k\}) = \varnothing, \qquad E(\{r, k\}) = \{o\}.$$

Exactness for the two singleton acquisitions forces $f(r) = f(k) = \varnothing$. Their union is then empty for the joint acquisition, contradicting $E(\{r, k\}) = \{o\}$. □

Meaning of exposure.

Membership of an origin o in $E(A)$ records derivability of its protected seed under the specified closure rules; the identifier itself may be public. Computational estimation is measured separately by success probability. A one-time-pad illustration makes the dependence on joint observations explicit: let S, W be independent uniform λ-bit strings and $C = S \oplus W$, with no additional observation depending on S or W. Then

$$H(S \mid C) = H(S \mid W) = \lambda, \qquad H(S \mid C, W) = 0.$$

From either observation alone, any estimator of S succeeds with probability at most $2^{-\lambda}$; from both, $S = C \oplus W$ recovers it exactly. The XOR construction serves here as an information-theoretic illustration of acquisition; the protocol's record encryption is handled by its cryptographic record interface.

Consequence for provenance routes.

A source label and an exposure decision serve different purposes. Copying a retained record preserves the source label, while subsequent acquisition of a matching key may change exposure. A guarded-route representation retains this distinction: alternative acquired records supply alternative routes to the same source, and every route retains the conjunction of its required keys. Exposure is therefore evaluated from the joint acquisition and key-derivation history. The separation applies specifically to an exact exposure representation formed by the union of fixed per-object labels.

4.5 A recovery experiment with acquisition-based eligibility

Definition 4.7 (Fixed-target recovery experiment) Fix a protocol profile $\mathcal{P}$ and finite execution budgets, together with a designated session s, recovery cut c, target epoch e, and candidate generation and encapsulation reservation handles g_e, a_e. Handles may be named by original invocation slots before those invocations occur; a missing invocation will fail the reservation-event test below. Public protocol messages may be delayed, duplicated, reordered, or replaced. The adversary can capture the declared retained objects, reveal wrapping keys, reveal actual signing keys or signer files, and reveal designated reservation tapes. The cut exposes the selected wrapping keys; preceding roots may already be public or disclosed by the stated prior captures. These disclosures have their exact scope in the event ledger.

For one fresh logical application send, the adversary supplies equal-length messages (m_0, m_1). The challenger samples $b \leftarrow \{0,1\}$ and runs the ordinary sending transition on m_b. Rejection or an unknown outcome keeps its ordinary protocol behavior. A retained accepted send is identified by (s, e, d, j), including direction d and the epoch-scoped logical index j; exact retries reuse its original ciphertext. The challenge is valid only when that logical send is accepted in the designated epoch. Subsequent protocol actions and disclosures continue until the adversary outputs b'.

Write ι_A, ι_B for the two actual signing-key identities configured for the target session. The predicate $\mathsf{ReservedFor}(s, e, g_e, a_e)$ checks the occurrence of the named honest generation and encapsulation reservation events in the target epoch's local durable plans, with their target session and roles. It checks event metadata, without asserting that their outputs form the accepted target exchange. Eligibility uses the following event conditions over the complete execution:

$$\begin{aligned}\mathsf{Eligible}_{s,e,c}(\tau) \iff\ & \mathsf{ValidTarget}(\tau) \ \wedge\ \mathsf{ReservedFor}(s, e, g_e, a_e) \\ & \wedge\ c < \mathsf{res}(g_e) \ \wedge\ c < \mathsf{res}(a_e) \\ & \wedge\ \iota_A, \iota_B \notin \mathsf{IdentityExposed}(\tau) \\ & \wedge\ \mathsf{EntropyExposed}(s, e; \tau) = \varnothing \\ & \wedge\ \nexists v \in \mathcal{C}_T : \mathsf{subject}(v) = s \ \wedge\ c < \mathsf{save}(v). \end{aligned} \tag{6}$$

Here $\mathsf{EntropyExposed}(s, e; \tau)$ records the target epoch's designated entropy disclosures. The separate event Attr states that these named reservation events form the accepted target's unique usable generation/encapsulation pair, with an adapter-computable association to the original slots. In particular, it connects the generation public key, the encapsulation input, the authenticated response, and the target's subsequent key schedule. This matching is a source/authentication obligation: a failure of Attr can occur while $\mathsf{Eligible}_{s,e,c}$ holds and is charged in the computational accounting.

The final line of (6) concerns every saved version of session s, across its roles and epochs. It uses acquisition $\mathcal{C}_T$, including protected bytes, at the time they are obtained. Readability $\mathcal{L}_T$ is evaluated separately. Other epochs' entropy disclosures retain their ordinary behavior. Both peers' confirmation is not an additional send prerequisite.

Define the unconditional centered advantage by

$$\mathsf{Adv}^{\mathrm{rec}}_{\mathcal{P}}(\mathcal{A}) = \left| \mathbb{E}\Big[\mathbf{1}_{\mathsf{Eligible}_{s,e,c}(\tau)} (-1)^{b \oplus b'} \Big] \right|. \tag{7}$$

Invalid or ineligible executions contribute zero. This definition retains failures and adaptive disclosures in the probability space and avoids conditioning a reduction on an uncontrolled eligibility probability. If every execution with $\mathsf{Eligible}_{s,e,c} = 0$ is completed by an independent fair guess, the conventional advantage over one half is

$$\mathsf{Adv}_{\mathrm{msg}} = \left|\Pr[\text{completed guess is correct}] - \tfrac{1}{2}\right| = \tfrac{1}{2}\,\mathsf{Adv}^{\mathrm{rec}}_{\mathcal{P}}(\mathcal{A}).$$

The computational accounting uses this latter convention.

The adaptive version retains one application challenge and permits the candidate selector to choose its coordinates and reservation-event handles adaptively from the observed interaction. Eligibility requires those coordinates to identify the challenge's accepted logical send; selection never transfers its test bit to another send. Eligibility and the usable-pair attribution are evaluated on that experiment's complete history. An origin-pair reduction fixes a slot pair in $[G] \times [C]$ and scores the event that the final selector attributes the target to that pair. The selector and disclosure predicates are recomputed in each comparison world. No independence between that selection and the transcript is imposed.

The record restriction in (6) is the one used by the retained multi-epoch symbolic theorem: its $\mathsf{Snapshot}$ action outputs a protected record, while Cut separately releases the wrapping keys. The theorem establishes symbolic plaintext secrecy with both original reservations after the cut, key-scoped identity exclusions, the target entropy exclusion, and the stated post-cut saved-record acquisition exclusion [26]. Its joint replay verifies the main lemma and all 24 reusable dependencies. Equation (7) specifies the computational message experiment to which the cryptographic modules are to be connected; the alias lemma establishes its record-interface normalization.

Corollary 4.8 (Safe stopping at a forbidden acquisition) *Consider the fixed-target experiment of Definition 4.7. Suppose an actual successful capture contains a target-session version v with $c < \mathsf{save}(v)$, identified by the original-record metadata. For any terminal bit D, replacing the remainder of that execution by output zero leaves*

$$\Pr[\mathsf{Eligible}_{s,e,c}(\tau) \wedge D = 1]$$

unchanged. Replacing its terminal score by zero also leaves the expectation inside (7) *unchanged.*

In particular, if a captured target-session version retains the selected reservation $o \in \{g_e, a_e\}$, then $c < \mathsf{res}(o)$ and (3) *establish this stopping condition. The wrapping key may be obtained before or after the capture.*

Proof The successful capture inserts v into $\mathcal{C}_t$ immediately. Its save time and subject are immutable, and acquisition remains recorded after host deletion, resealing, confirmation, or process loss. Thus the last conjunct of (6) is false for every continuation. Each such continuation already contributes zero to both quantities. On executions that never reach this condition, the experiment is unchanged. The final assertion follows from $c < \mathsf{res}(o) \le \mathsf{save}(v)$. □

The stopping test uses the successful acquisition event and its record metadata; it can run before the response's private bytes are serialized by a challenge adapter. It is not an SDK rejection of the capture. A public digest, an unrelated record, or an unsuccessful

capture keeps its real behavior. In particular, difficulty computing a permitted digest or report MAC does not trigger this corollary. For adaptive target selection, stopping is applied to the corresponding fixed-candidate event in the reduction. One possible candidate becoming ineligible does not justify terminating a game that can still select a different eligible target.

Two short executions illustrate why both kinds of origin are necessary. If v was saved before c and a sibling write reseals it after c, the acquisition still refers to the original pre-cut v. The new global image time alone does not falsify the final conjunct of (6). If instead v was saved after c, a later write attributed to a sibling does not change its subject or make its capture eligible. Likewise, a pre-cut generation reservation retained through a delayed receipt remains pre-cut when its KEM computation completes later. The reservation condition in (6) determines that case, with the endpoints retaining their respective confirmed counters.

4.6 Connection to the cryptographic argument

The normalization theorem fixes the original objects whose bytes, acquisition, and key identities the reduction must represent. The next section gives constructive image/intent and report interfaces for a selected KEM operation whose public outputs are available while its original random tape is opaque. Their algorithms preserve original-object caches, public commitments, exact retries, and actual durable outcomes. The fixed-target stopping rule above extends to an adaptive provenance certificate in Lemma 5.6.

Usable attribution is computed from available semantic metadata. Original generation and encapsulation slots belong to a fixed lifetime budget $[G] \times [C]$: a fresh draw receives a new slot, while copies retain their source. Each candidate induces its own protected records and report-key closure. Proposition 5.8 decomposes the score in the native experiment before comparing those local continuations. It retains both unattributable eligible executions and permitted key projections as explicit native-game terms.

The computational composition evaluates accepted operations, eligibility, and all public outputs anew in each comparison world. Exact component repeats retain one response and one correctness event. The native constructor map in Section 5.4 identifies the storage operations used by these interfaces; Section 8 specifies the joint root and message continuation to which their bounds connect.

5 Computational interfaces and composition

The origin model provides the accounting boundary for a computational experiment. Each original generation or encapsulation invocation receives a slot when its entropy is reserved, including attempts that later fail. Replaying or copying that reservation retains its slot. Fix public lifetime budgets $G, C \geq 1$ and write $J = [G] \times [C]$. Unused slots remain unused. These proof-resource budgets differ from a bound on the number of records resident at one time.

The native application experiment evaluates the selected target after the adversary's interaction. Eligibility requires both original reservations to follow the selected cut, secrecy of the actual target signing keys and target-epoch entropy, and no acquisition

Table 3 Shared dependencies of the native computational results. Each row identifies an interface and its establishing result.

Connection	Supplied interface and principal result
Record exposure	Original sources, immutable versions and acquisition history; guarded normalization (Theorem 4.4).
Private retained state	Real-to-opaque image and intent correspondence (Corollary 5.3); candidate-local stopping and original-score localization (Lemma 5.7).
Hybrid contribution	Original-operation component adapter and complete-tuple comparison; native single-root use (Lemma 5.15).
Application continuation	Joint confirmation and traffic schedule (Theorem 5.10); online prefix continuation (Theorem 5.12).
Source attribution	Executed-slot allocation provenance (Proposition 5.20), instantiated through ordinary hedged ML-DSA (Theorem 5.24).

of a selected-session version saved after that cut. It admits the asynchronous transitions of the construction, including sending after the local FinalReady transition.

The main native results are the established-session confidentiality bound (Corollary 5.17), selected-root authentication (Theorem 5.25), and their online selected-message extension (Theorem 5.27). Their proofs share three connections. Private-record and immutable-report simulations feed candidate-local accounting, whose permanent stopping certificates use original provenance. The complete ContextBound and root-entry comparisons then supply a uniform recovered root. Finally, the joint schedule and online prefix comparison carry that root to the selected message while retaining the declared verification and disclosure operations. Executed-slot attribution and the ordinary-signature bridge close the association between that message and its original KEM sources.

Proof dependencies and advantage accounting.

Table 3 identifies the interface supplied by each connection. Each fixed-candidate comparison retains the original label $j = (g, a)$ and evaluates the same terminal code on that world's execution. Eligibility, accepted-send provenance and acquisition history are recomputed from its own ledgers. The final selector may choose the compromise cut after interaction; the protected tuple, traffic entry and message step are designated at their specified first-private-allocation boundaries. This order permits adaptive target selection while preserving the simulator's input and alias tables.

The final confidentiality bounds compare a signed score: a comparison whose output-bit advantage is ε contributes at most 2ε before the final factor $1/2$ converts that score to message advantage. The operational component uses the two actual positive and negative output tests, giving $a_j^+ + a_j^- + 2\delta_j$. Authentication compares its own Boolean event, so the corresponding record cost is $\varepsilon_{\mathrm{rec},j}$ and its component cost is $a_j^{\mathrm{auth}} + \delta_j$, or $a_j^{\mathrm{msg}} + \delta_j$ for online message integrity. Each of these adapters retains its own whole-execution query budget. The source-attribution residual is localized

once in the original game. The live core supplies its complete observation interface in Section 5.8; the separate report extension carries its declared projection event.

Consequently the three native statements share the record and source connection, with the continuation selected by the terminal test: the first-message score in (21), the confirmation/first-message event in (31), and the online message score and event in (35)–(36). The native online correspondence (Lemma 5.26) is the additional source-order connection for the last pair.

The final two subsections develop the finite multiple-root and SDK purpose-key instances of the same primitive method. The common-virtual-game accounting variant is retained in Appendix D.

5.1 An explicit renderer for private images and intents

The alias lemma preserves a renderer that can read every original payload. A component reduction additionally needs to execute record management when the selected KEM challenger supplies public outputs and a shared-secret challenge without supplying the selected random tape. We construct that record-management algorithm for a finite image/intent interface. Its tables contain immutable handles, known fields and locally sampled sealing coins. The protocol's public computations are supplied by a common functional interface F, instantiated by the original-operation adapter. This section holds F fixed and isolates the image/intent contribution to the record connection in Proposition 5.8.

Typed record interface.

An opaque block is an immutable fixed-length private value, represented by a handle and its known length. Blocks may be correlated and may themselves be retained sealed reservations. A private image recipe contains at least one such block; an atom-free recipe is fully known. Local image-encryption and intent-authentication keys are available to the simulator. Fix a public budget M for all fresh image-sealing invocations, including failed attempts and atom-free images. Exact copies and exact retries retain their invocation. The following operations specify the interface.

1. *Seal.* Before sampling its nonce, invocation i fixes the recipe, its length and the non-nonce header fields. It samples independent $N_i \leftarrow \{0,1\}^n$ and privately constructs
$$B_i = \mathsf{Hdr}_i \| N_i \| C_i,$$
where C_i is the specified deterministic local authenticated encryption under that nonce. Local encryption and intent authentication use their prescribed fixed algorithms without additional shared-SHA3 calls. All admitted image formats place the nonce at the same offset. The reply consists of the image handle, length and declared non-nonce metadata; N_i, C_i and the recipe's private blocks remain private.
2. *Intent.* A complete canonical prefix P fixes the tag, binding, predecessor and the image-head value used in this execution. Its tag determines its length. The intent bytes are
$$W_{i,P} = P \| B_i \| \mathsf{MAC}_{k_i}(P \| B_i),$$

with one fixed authentication key k_i for image i. An exact repeat retains the same (i, P) entry. Image and intent tags are disjoint.

3. *Public commitments.* For a designated private image, the ideal-commitment experiment returns independent uniform 256-bit values cached by i for its head and by (d, i, P) for its intent digests. The domain d ranges over the five admitted intent-digest domains. Atom-free objects retain their real hashes and ordinary raw-byte outputs. Later validation and zero-value checks use the actual replies in that execution.
4. *Equality and storage.* Equality queries compare complete retained image or intent handles. Copy, retain, persist, delete, crash and resume update an explicit durable-state directory. Storage choices, errors and branch decisions depend on available metadata, lengths, commitments and F outputs. A failure after commit preserves that commit.
5. *Acquisition.* Atom-free objects are rendered from their known inputs. For an opaque object, the admitted successful acquisition emits a computable event that permanently falsifies the fixed-candidate eligibility before its raw bytes are returned. The stopped score is zero, as in Corollary 4.8. Unavailable-object requests retain their ordinary error result.

All other public computation uses these replies, known inputs, F, and the shared public oracles. This is a classical record interface; the adversary may have quantum private state and quantum access to those public oracles. The successful-acquisition condition is supplied by the recovery experiment's existing event rule. For a fixed target session s and cut c, acquisition of an original version v with $\mathsf{subject}(v) = s$ and $c < \mathsf{save}(v)$ supplies such an event. Its original-record provenance determines that save time, including when a later aggregate carries the same version. Lemma 5.6 gives a provenance certificate valid before the final cut is chosen. A separately permitted projection is described by its own output operation.

Table algorithm.

The simulator stores an image table indexed by fresh sealing invocation, an intent table indexed by image and complete prefix, a commitment cache, and the directory with its provenance and acquisition ledgers. Each new image entry records the recipe handles, length, known fields and a sampled nonce. Each new intent entry records (i, P) and its known length. The simulator samples designated-private commitments on first use and retrieves them on exact retry; it constructs and hashes atom-free objects normally. Equality is structural: images compare equal exactly when their indices match, and intents compare equal exactly when their pairs (i, P) match. The storage directory determines the state recovered after a crash, including an acknowledged or uncertain write that actually committed. At the stipulated permanent-failure acquisition event the computation returns zero. Every other transition executes the common public code and F. Thus the algorithm uses no selected opaque block bytes.

Lemma 5.1 (Structural equality of complete records) *If the nonces of fresh image invocations are pairwise distinct, then*

$$B_i = B_j \iff i = j, \qquad W_{i,P} = W_{j,P'} \iff (i, P) = (j, P').$$

Proof One immutable image handle retains one byte string. Two different image invocations have different nonces at the same explicit offset, so their complete wires differ. Intent tags first separate different formats. For one tag the prefix length is fixed: unequal prefixes differ before the target image, while equal prefixes with different images differ at the target's nonce. The same image and prefix give the same deterministic MAC under the fixed key. Disjoint tags also separate images from intents. This argument permits coincident private block values and ciphertext components; the explicit nonces distinguish the complete images. □

Let I be the materialized ideal-commitment experiment, which retains real private bytes and answers equality by comparing them. Let T be the table algorithm above. Both use the stipulated fixed-candidate stopping rule and the same public-oracle model and interface F.

Theorem 5.2 (Opaque image/intent simulation) *For at most M fresh seals with independent n-bit nonces, define*

$$\delta_{\text{nonce}} = \min\{1, \tbinom{M}{2} 2^{-n}\}.$$

For every efficient terminal bit D computed from the public view, available metadata and event ledger, F outputs and its own independent test coins (including its retained application challenge bit),

$$|\Pr[I : D = 1] - \Pr[T : D = 1]| \le \delta_{\text{nonce}}.$$

For a score $Z \in [-1, 1]$ computed from the same available inputs, $|\mathbb{E}_I Z - \mathbb{E}_T Z| \le 2\delta_{\text{nonce}}$.

Proof Couple the fresh sealing nonces, private-commitment replies, scheduler choices, the state of F, and the public-oracle transformations. Before the first nonce collision, maintain agreement of the directories, known fields, immutable references, public bytes, statuses and semantic ledgers. Sealing adds the same metadata and nonce to corresponding entries; its private ciphertext is confined to the materialized entry. Intent creation adds the same complete prefix and target reference. A fresh private commitment gets the coupled reply, and a repeat consults the same cache. Atom-free commitments and acquisitions use identical materialized inputs. Lemma 5.1 gives agreement of equality replies. The common public computation consequently makes the same transitions and oracle calls. Copying, committed or failed writes, deletion, crash and resume preserve the directory invariant. An opaque acquisition emits the same permanent-failure event before either computation needs to return its unknown bytes, and both stopped scores are zero.

The next nonce is independent uniform even when the decision to seal is adaptive. Across at most M invocations, a union bound charges $\binom{M}{2}2^{-n}$ for any collision. Coupled terminal bits differ only on that event; bounded scores can differ by at most two. Identical classical replies and shared-oracle transformations also preserve a quantum adversary's state at every coupled good prefix. The argument uses no measurement of a quantum query. □

Connecting real commitments.

After replacing complete-wire equality by structural equality, let ϵ_{com} bound the bit-distinguishing gap between real private record digests and the independently cached commitments of the typed interface. Both games run the same public code, ordinary known-record computations, available metadata and declared stopping rule. The comparison counts the full distinguishing computation against their shared public oracle. It is a cryptographic input to the following composition, separate from the explicit table simulation. Appendix B records the concrete multi-record bound under fixed-input injectivity, per-nonce confinement and the prescribed local encryption/authentication algorithms.

Corollary 5.3 (Real-record to table-algorithm bound) *For the declared fixed-candidate interface and the commitment conditions above, let R be the real-digest materialized experiment. With $n = 192$,*

$$|\Pr[R : D = 1] - \Pr[T : D = 1]| \leq \min\{1, \epsilon_{\mathrm{com}} + \delta_{\mathrm{nonce}}\}.$$

The corresponding $[-1, 1]$ score bound is $\min\{2, 2(\epsilon_{\mathrm{com}} + \delta_{\mathrm{nonce}})\}$. When the interface exposes no private-wire equality query, its additional δ_{nonce} term is zero.

Proof First replace real byte equality by structural equality, coupling exactly as above and charging δ_{nonce}. The resulting game has no extra nonce-dependent equality observation, so apply the stated private-commitment replacement there. Its ideal form is implemented exactly by the table algorithm. A triangle inequality proves the bit bound. Applying that bound to a test that returns one with probability $(1 + Z)/2$ proves the score bound. Without a private-wire equality operation the first replacement is absent. □

The hybrid order starts with constructible real record inputs, replaces their private commitments, and then uses the table algorithm for the KEM challenge embedding. Every changed commitment is consumed by that world's subsequent intent prefixes, checks and outputs. For an adaptive origin-pair event E_j, early stopping is justified only by a computable event making E_j false for every continuation and every still-admissible final target choice attributed to j. The disjoint-event accounting above retains this per-pair obligation.

5.2 Immutable report families and reverse-order composition

A later cleanup report can authenticate an earlier report's private key and tag. We compose the one-report HMAC replacement by creation order, retaining these dependencies. The reduction below supplies a concrete algorithm for handling a selected unknown report key: later dependent rows are already ideal, and earlier immutable transcripts precede that key.

Finite interface.

Fix a public budget N for fresh report-key draws and index the actual draws chronologically. Failed attempts retain their indices; a new draw after definite absence receives a

new index. Exact replay of a retained original report keeps its index, key, and transcript. Each K_i is an independent uniform 256-bit value, with no rejection of the zero key.

Let $\mathcal{X}$ be the source-atom handles whose bytes the subsequent KEM adapter will not possess. In this report replacement, their real bytes are still available to the surrounding source experiment. Before using K_i, the machine fixes a bounded immutable transcript recipe m_i. It can contain available literals and protocol outputs, atoms in $\mathcal{X}$, earlier raw keys K_j with $j < i$, and previously returned tags. Its own key and tag do not occur in m_i. Lengths, dependency sets, and admission errors are computable from the recipe metadata. Native code may draw the key before assembling the transcript, preserving that order and its failures, provided those inputs do not depend on the draw. Every direct or computed use of a raw key retains its origin dependency; returned tag bytes are available values.

Compute an online protection flag before using the fresh key:

$$P_i = \mathbf{1}[\mathrm{deps}_{\mathcal{X}}(m_i) \neq \varnothing \ \vee \ \exists j < i : \ P_j = 1 \ \wedge \ K_j \in \mathrm{deps}_K(m_i)] \,. \tag{8}$$

This is forward closure under raw-key dependencies. It need not propagate through a previously returned tag. The flag uses origin references, not actual key equality or the primitive challenger's bit, and is recomputed in each comparison world's history.

The following operations specify the interface. Creation computes $t_i = \mathrm{HMAC}_h(K_i, m_i)$ and retains the actual nonzero-ID check, encoding errors, and persistence outcomes. Replay or validation names the original report and uses that same m_i; supplied identifiers can be compared with its retained tag. This operation does not submit changed private bytes as a new transcript under K_i. Acknowledgement, terminal-tag comparison, copying, and crash recovery preserve their ordinary original-handle semantics. Private storage uses the explicit recipe tables of the opaque-record interface, including its already idealized commitments and its separate error accounting. Public report bodies and other outputs use available inputs or their specified application interface. Protected raw keys have no additional unlisted computation or observation channel.

The full real experiment G retains each declared fragment/key disclosure. Let $\mathsf{Bad}_{\mathrm{rep}}$ be the first successful outward projection whose recipe depends on a raw K_i with $P_i = 1$, detected by its actual origin map before returning the dependent bytes. An unavailable fragment, a terminal tag alone, or a projection without that dependency does not trigger it. Define G^{stop} to end with score zero at this event. This is an analysis game; a report-key fragment is not thereby declared to be a selected-session acquisition or a new freshness failure.

One-row bound and hybrids.

Use one public compression oracle h for all rows and quantum queries. The one-row experiment draws a uniform hidden 256-bit key K, independently of h and oracle-independent initial auxiliary state. The algorithm performs its complete counted preparation and chooses one classical transcript m of length at most L without a raw-key input. Its first row query returns either $\mathrm{HMAC}_h(K, m)$ or an independent uniform 256-bit tag; at most $T \geq 1$ row calls repeat this same transcript and cached

answer. The algorithm has no additional raw-key, midstate or key-dependent output interface. It may continue querying the same h and output a bit. The budget q counts all public compression queries, including preparation, other real rows, continuation and the terminal test. Independently sampled known keys may coincide with K; their ordinary calls to h remain allowed and counted.

Let $\epsilon_{\mathrm{row}}(T, q, L)$ be any upper bound on this experiment's bit-distinguishing gap. This operational definition is the primitive input to the composition theorem. Appendix B gives the compression-random-oracle instantiation, its padding convention and query accounting. A zero answer is retained for the protocol's actual error check.

Define H_k, for $0 \le k \le N$, by running the stopped experiment with every unprotected row real, every protected row $i \le k$ real, and every protected row $i > k$ replaced by its own independent uniform cached tag. All later recipes, flags, errors, and selections use that world's actual replies. Thus $H_N = G^{\mathrm{stop}}$ and $H_0 =: S^{\mathrm{stop}}$. A zero ideal tag takes the original error branch; it is never resampled. An absent invocation makes no row query.

Theorem 5.4 (Reverse-order report composition) *For each slot i, let L_i, T_i, q_i uniformly bound the transcript length, row calls, and complete ordinary compression work of the reduction below, including its prefix, future computation, other real rows, and final test. Set*

$$\epsilon_{\mathrm{rep},i} = \epsilon_{\mathrm{row}}(T_i, q_i, L_i), \qquad \epsilon_{\mathrm{rep},\Sigma} = \min\left\{1, \sum_{i=1}^{N} \epsilon_{\mathrm{rep},i}\right\}.$$

Under the stated interface, every common efficient terminal bit test D satisfies

$$\left|\Pr[G^{\mathrm{stop}} : D = 1] - \Pr[S^{\mathrm{stop}} : D = 1]\right| \le \epsilon_{\mathrm{rep},\Sigma}.$$

For a terminal score $Z \in [-1, 1]$ which is zero on stopping,

$$|\mathbb{E}_{G^{\mathrm{stop}}} Z - \mathbb{E}_{S^{\mathrm{stop}}} Z| \le 2\epsilon_{\mathrm{rep},\Sigma}.$$

The tests use the common available view, semantic ledger, protocol-interface outputs, and independent test coins, including the retained application challenge bit. They have no extra access to protected key or transcript bytes.

Proof Fix i and compare H_i with H_{i-1}. Run the full prefix from setup, locally sampling all earlier keys and computing their real rows. Before the selected row is used, both hybrids have the same prefix rules. If slot i is absent or unprotected, their behavior agrees and no row challenge is used. Otherwise install the one-row challenger's key as an opaque $\mathsf{KeyRef}(i)$. The transcript m_i is computable: it excludes its own key, all referenced earlier keys were locally sampled, and the source atoms are still available at this proof stage. Submit this actual transcript to the row oracle when its first MAC call is reached, retain its reply, and run the original nonzero and persistence checks.

Maintain the original-report, key-origin, recipe, durable-state, and event tables. Every available output agrees with the selected hybrid. Earlier real rows remain computable because their immutable transcripts predate K_i and their keys are known. A future unprotected row cannot depend on $\mathsf{KeyRef}(i)$: such a dependency would make it protected by (8). Compute that row with its locally sampled key and public h. A future protected row has index greater than i and is already ideal in both hybrids; sample its first tag, cache it by original invocation,

and retain its symbolic recipe without evaluating the unavailable bytes. This includes a later report which contains K_i or another earlier protected key and tag. Flags and subsequent known fields are recomputed using the actual returned tags.

For replay and validation, the selected row is called only on its original m_i; preserve each counted recomputation and compare supplied identifiers with its retained reply. Copying, failure after a durable commit, and crash/resume use the same immutable handles and the actual reached storage state. Each report retains its original key and immutable transcript. Public decisions use available metadata and common-interface outputs. Before an outward protected-key projection, the same ledger event $\mathsf{Bad}_{\mathrm{rep}}$ enters the common stopped state with application score zero; any allowed terminal-bit postprocessing uses that state. No unavailable key bytes are serialized. Unprotected-key fragments retain their known bytes. This gives an explicit online simulation throughout the continuation.

The selected real row therefore produces H_i, and its ideal counterpart produces H_{i-1}. All other real keys are sampled as actual independent byte strings. If one happens to equal the unknown challenged key, evaluating its real MAC with public h automatically preserves that equality; those queries belong in q_i. No comparison with the hidden key or distinct-key conditioning is used. Intentional aliases of the same original report instead share its single row cache.

The reduction executes the adversary's quantum queries to the same h and introduces no measurement of those queries or oracle-dependent initial advice. The one-row lemma bounds the adjacent bit difference by $\epsilon_{\mathrm{rep},i}$. Sum over the reverse-order hybrids and cap at one. For the score, use the allowed randomized terminal test which outputs one with probability $(1+Z)/2$; its difference is half the score-expectation difference, giving the displayed factor two. □

Corollary 5.5 (Original eligibility and real disclosures) *For the full real experiment G, define*

$$\eta_{\mathrm{rep}}(Z) = \mathbb{E}_G\left[|Z|\,\mathbf{1}_{\mathsf{Bad}_{\mathrm{rep}}}\right].$$

Then

$$|\mathbb{E}_G Z - \mathbb{E}_{S^{\mathrm{stop}}} Z| \le \min\left\{2, \eta_{\mathrm{rep}}(Z) + 2\epsilon_{\mathrm{rep},\Sigma}\right\},$$

and hence

$$|\mathbb{E}_G Z| \le \min\left\{1, \eta_{\mathrm{rep}}(Z) + 2\epsilon_{\mathrm{rep},\Sigma} + |\mathbb{E}_{S^{\mathrm{stop}}} Z|\right\}.$$

For $Z = \mathbf{1}_{\mathsf{Eligible}}(-1)^{b\oplus b'}$, $\eta_{\mathrm{rep}}(Z) = \Pr_G[\mathsf{Eligible} \wedge \mathsf{Bad}_{\mathrm{rep}}]$.

Proof Couple G and G^{stop} up to the first $\mathsf{Bad}_{\mathrm{rep}}$. On its complement the runs agree; on that event the stopped score is zero while G continues with the actual disclosed bytes. Therefore

$$\mathbb{E}_G Z - \mathbb{E}_{G^{\mathrm{stop}}} Z = \mathbb{E}_G[Z\mathbf{1}_{\mathsf{Bad}_{\mathrm{rep}}}].$$

Bound its magnitude by $\eta_{\mathrm{rep}}(Z)$ and apply Theorem 5.4. No conditioning on eligibility or independence of later target selection is used. □

Consequences for the record connection.

The residual is zero on a disclosure surface whose actual provenance already makes the original target permanently ineligible. A batch-only report-key fragment requires its own mapping: it is not automatically a member-session capture. This distinction is material. Given a known transcript and its later disclosed key, recomputing HMAC distinguishes the real tag from an independent tag with difference $1 - 2^{-256}$. This

example does not assert that a native batch fragment also reveals its members' complete transcript.

The report replacement precedes KEM embedding. Afterwards every protected row needs only its original cached tag, and every unprotected recipe is free of $\mathcal{X}$ and protected raw-key dependencies. The report layer can therefore use opaque source handles for tag generation and immutable tag validation. A native composition first charges the actual eligible-exposure residual, then uses a private-record correspondence proved for the explicitly stopped native interface, and finally applies the report hybrids. This record hop retains its own comparison error and acquisition rules. For key-only fragments, that correspondence requires an explicit mapping of their actual provenance and justification of every permanent-ineligibility premise used by the renderer. Direct source-atom or application-plaintext outputs retain their own adapters and acquisition rules. Immutable validation uses the original transcript and cached tag; arbitrary changed-byte import has a distinct MAC-query interface. Under these stated interfaces, the explicit reduction covers dependent immutable report rows.

5.3 Candidate-local simulation in the native experiment

A guessed origin pair can determine which records are opaque, which report rows are protected, and where the reduction stops. We therefore localize the score in one native experiment before comparing its candidate-specific continuations. This formulation preserves the original target eligibility and does not require the local real-component games to coincide.

Let $\mathsf{G}_{\mathrm{nat}}$ be the full native experiment and $J = [G] \times [C]$ its fixed reservation-slot budget. Write E for original eligibility, A for usable unique attribution, and $I \in J$ for the attributed pair on A. These variables describe the actual accepted logical challenge send; selection does not transfer its test bit to another send. For a bounded application score $S \in [-1, 1]$, set

$$Z = \mathbf{1}_E S, \qquad Z_j = \mathbf{1}_{E \wedge A \wedge (I=j)} S.$$

For each j, let U_j and B_j be efficient, cumulative prefix events computed from the available classical provenance ledger. They are noninterfering monitors: before stopping they change no native response or protocol transition and do not inspect an unknown tape or a quantum query. A *free stop* U_j permanently falsifies $E \wedge A \wedge (I = j)$ for every continuation. A *charged stop* B_j may instead be a permitted protected-report-key projection. Define $\mathsf{G}_{\mathrm{nat}}^{[j]}$ to run the native code until the first such event, then enter a stopped state with candidate score zero before returning the monitored bytes. An event for one candidate does not stop the original game or another candidate's reduction.

Lemma 5.6 (Origin-anchored adaptive stopping) *Fix candidate j and one of its original reservation slots o which has occurred. Suppose a successful capture supplies an intact version v with*

$$\mathsf{subject}(v) = \mathsf{session}(o), \qquad \mathsf{res}(o) \leq \mathsf{save}(v).$$

Assume usable attribution to j assigns its member origins to the selected session, and eligibility places the cut before both reservations. Then this capture is a free U_j event, even if the other reservation or challenge send has not occurred and the final cut has not been selected.

Proof In any continuation with $E \wedge A \wedge (I = j)$, the selected session is $s = \mathsf{session}(o)$ and the selected cut satisfies $c < \mathsf{res}(o) \leq \mathsf{save}(v)$. Acquisition has already inserted v into the cumulative archive. Thus the continuation has acquired a version of s saved after c, contradicting the original acquisition conjunct. Copies, resealing, deletion, and recovery preserve this contradiction because the subject, original save event, and acquisition are immutable. If the other reservation or challenge send never occurs, the attributed eligible event is already false. Decryption and a final cut value are unnecessary for this certificate. □

Actual retention of o in v supplies the time inequality by reservation/save causality; an independently recorded later save of the same session also suffices. More generally, let L_s be the largest original save time of an acquired intact version of session s, with $L_s = -\infty$ for an empty set. For an existing pair (g, a), the original temporal eligibility conditions are exactly

$$L_s \leq c < \min\{\mathsf{res}(g), \mathsf{res}(a)\}.$$

The remaining eligibility conditions retain their own tests. This statistic uses original version times, not the enclosing image's latest sealing time. A batch-only key fragment or a public commitment supplies no intact-session witness merely because a report depends on that session.

One native residual.

In the full native run let B_I denote the charged event for its final attributed pair, used only on A, and define

$$\eta^{\mathrm{nat}}_{\mathrm{attr}} = \Pr_{\mathsf{G}_{\mathrm{nat}}}[E \wedge \neg A], \qquad \eta^{\mathrm{nat}}_{\mathrm{proj}} = \Pr_{\mathsf{G}_{\mathrm{nat}}}[E \wedge A \wedge B_I].$$

These are disjoint events in one game. In particular,

$$\eta^{\mathrm{nat}}_{\mathrm{proj}} = \sum_{j \in J} \Pr_{\mathsf{G}_{\mathrm{nat}}}[E \wedge A \wedge (I = j) \wedge B_j].$$

A projection affecting only another candidate is not charged to the selected one. For the report construction, B_j is the actual protected-key projection event for the closure induced by that candidate's source family. Its probability is retained explicitly; the theorem does not declare it negligible or turn the fragment into a session snapshot.

Lemma 5.7 (Native score localization with local stops) *Let Ψ_j be the same terminal code for $\mathbf{1}_{E \wedge A \wedge (I=j)} S$ in each experiment, with value zero in its stopped state. Then*

$$\mathbb{E}_{\mathsf{G}^{[j]}_{\mathrm{nat}}} \Psi_j = \mathbb{E}_{\mathsf{G}_{\mathrm{nat}}}[Z_j \mathbf{1}_{\neg B_j}], \tag{9}$$

$$\mathbb{E}_{\mathsf{G}_{\mathrm{nat}}} Z = \mathbb{E}_{\mathsf{G}_{\mathrm{nat}}}[Z \mathbf{1}_{\neg A}] + \mathbb{E}_{\mathsf{G}_{\mathrm{nat}}}[Z \mathbf{1}_{A \wedge B_I}] + \sum_{j \in J} \mathbb{E}_{\mathsf{G}^{[j]}_{\mathrm{nat}}} \Psi_j. \tag{10}$$

Proof Until its first local stop, the stopped machine has the native prefix distribution. On a run with B_j, its candidate score is zero. On a run with U_j, every native continuation already

has $Z_j = 0$, whether or not B_j later occurs. On all other runs the executions and candidate scores agree. This proves (9). Partition the one native score by attribution and its unique index; pointwise, $\sum_j Z_j \mathbf{1}_{B_j} = Z\mathbf{1}_{A \wedge B_I}$. Substitution proves (10). The final index is used in this identity through retrospective attribution. For a quantum adversary the argument compares identical classical-prefix instruments and their stopped continuations, retaining the native quantum-state transformations and oracle-query interface up to the local stop. □

Local cryptographic comparisons.

For each j, let $\mathsf{V}^{[j]}$ be its specified real-component continuation after the local source/record/report replacements, and let $\widetilde{\mathsf{V}}^{[j]}$ be its corresponding random-component game. All terminal events are recomputed in their own game by the same score code; no event from a native transcript is transferred to an altered transcript. Suppose the actual stopped-interface correspondence gives

$$\left| \mathbb{E}_{\mathsf{G}^{[j]}_{\mathrm{nat}}} \Psi_j - \mathbb{E}_{\mathsf{V}^{[j]}} \Psi_j \right| \leq \rho_j.$$

This is a local score bound, including its native source correspondence. A binary record or report bound contributes twice its value to this score comparison. The stopping certificate does not itself establish the correspondence for permitted partial projections.

For the remaining component step, specialize to $S \in \{-1, +1\}$, so that $\Psi_j \in \{-1, 0, +1\}$. Suppose its actual fixed-pair adapter satisfies, for $\sigma \in \{-1, +1\}$,

$$\left| \Pr_{\mathsf{V}^{[j]}}[\Psi_j = \sigma] - \Pr_{\widetilde{\mathsf{V}}^{[j]}}[\Psi_j = \sigma] \right| \leq a_j^{\sigma} + \delta_j.$$

Here a_j^{σ} includes the complete derived distinguisher and its shared oracle work; δ_j is the selected pair's correctness term. Write $b_j = \mathbb{E}_{\widetilde{\mathsf{V}}^{[j]}} \Psi_j$.

Proposition 5.8 (Candidate-local native accounting) *Under the stated local comparisons, the native message-guessing advantage satisfies*

$$\mathrm{Adv}_{\mathrm{msg}} \leq \min\left\{ \frac{1}{2}, \frac{1}{2}\left(\eta^{\mathrm{nat}}_{\mathrm{attr}} + \eta^{\mathrm{nat}}_{\mathrm{proj}} + \sum_{j \in J} \left(\rho_j + a_j^{+} + a_j^{-} + 2\delta_j + |b_j| \right) \right) \right\}.$$

Proof The first two terms of (10) have magnitudes at most the two native residuals. The local correspondence costs ρ_j. Subtracting the positive- and negative-event comparisons bounds the local component score change by $a_j^{+} + a_j^{-} + 2\delta_j$. Sum the candidate terms, apply the triangle inequality, and use $|Z| \leq 1$. Completing invalid runs by an independent fair guess gives $\mathrm{Adv}_{\mathrm{msg}} = \frac{1}{2}|\mathbb{E}_{\mathsf{G}_{\mathrm{nat}}} Z|$. □

The ideal continuations may have different adaptive selectors. The terms $|b_j|$ therefore remain separate; they are not probabilities from one partitioned ideal world. If a genuine common intermediate game and a global comparison have separately been established, that shared error can instead be charged once before the common-game

accounting. The proposition above covers the candidate-local constructions without asserting such a global identification.

5.4 Native constructors and retained-object transitions

We map the declared record constructors to the fixed Continuity source. The accompanying audit contains 37 entries across 27 directly cited files, with exact source ranges and checked Git-blob and SHA-256 identities. It separates constructor layout, state transition, known-input computation and the protocol's additional output interfaces. The following correspondences give the native instances of the objects used in the preceding proofs.

Images, intents and commitments.

`durable::seal` encodes and bounds the payload before drawing a 24-byte nonce. The four image tags place that nonce at byte offsets $[80, 104)$, followed by the XChaCha20-Poly1305 ciphertext and authentication tag. The target-bearing `PendingWrite` formats have a canonical tag-dependent prefix, the complete original target image, and a deterministic HMAC under the journal-derived intent key. These layouts supply the fixed offset, prefix and intact-target premises of Lemma 5.1. The wrapping key for a reserved operation and the intent-authentication key use their respective derivation domains; the simulator can retain these local storage keys independently of the selected KEM random tape.

The source uses one image-head domain and five target-bearing intent-digest domains. Their callers feed the actual returned digests into fences, checkpoints and subsequent validation. Cancellation intents carry known descriptors instead of a private target image; selection rows and public archive descriptors likewise use their actual known-input computations. This variant distinction supplies the real-versus-designated-private split used by the commitment table.

Reservation and durable identity.

Both generation and encapsulation paths persist the reserved operation token before invoking the corresponding KEM computation. An original retry opens that token and consumes its retained entropy. The write-intent path first retains the complete pending wire. Applying it compares the retained original wire, checks the expected and next heads, installs its original target bytes, and removes the pending entry. An already installed exact target is handled as the same operation. A sibling-session write can instead seal a new image containing unchanged target payload bytes, preserving the target's source while creating a new physical sealing invocation.

Uncertain commit and recovery.

The native state distinguishes a caller's outcome from the durable directory. An unconfirmed commit maps to `CommitUncertain`; the owner is closed and subsequently reopened against the retained image and pending target. An actual committed write therefore persists across an error response. The explicit storage machine records this outcome and retains the original operation on resumption. The cited acquisition rule

Family	Immutable transcript and key placement	Ordinary report result
Session closure	The report MAC covers the session payload with its own 64-byte report/key tail zeroed. The pending key is retained in that tail. Earlier resolution keys can occur in the preceding traffic entries.	Session and epoch progress, counts and loss metadata.
Fanout abandonment	The MAC covers batch metadata and complete member payloads. Its own key is held in a separate batch tail, outside that metadata. Member payloads can include earlier resolution keys.	Per-member progress and loss metadata.
Closed-epoch resolution	The MAC covers the old epoch's counters, message identifiers, retained digests and skipped indices, excluding its own key. Its pending key is retained in the traffic entry.	Includes unconsumed plaintext for the authorized application.

Table 4 Source correspondence of the report families. The application's authorized plaintext output retains its own audience in the security experiment. Report-key fragments use their actual record content and origin mapping.

uses the contained session version's original save event, including when that version is carried by a later aggregate.

Report dependencies.

Table 4 identifies the three report families. The resolution key and pending tag are serialized in an epoch's traffic entry. A later closure or fanout report can authenticate that complete entry. This gives the causal raw-key edge used by the reverse-order report argument. The code's exact retry paths retain an original key and tag; a genuinely new attempt draws a new key. The nonzero check applies to the returned identifier, with a real zero-ID failure branch. Pending keys use the full 256-bit space.

The normal retained-object path reopens a registered image and revalidates its original semantic state. The interface mapping assigns those validation calls to known state, the protocol interface F, or the immutable report row. Arbitrary modified storage inputs and separately returned projections receive their corresponding parse, authentication and output operations in an extended experiment. In particular, the batch-key record and a member's session payload have different contents: their acquisition events are determined separately. These typed boundaries preserve the existing target predicate while making each native call's role in the reduction explicit.

5.5 Joint continuation from an epoch root to a message

A recovered root feeds confirmation, traffic derivation and the next epoch. We treat these uses jointly and then isolate a selected message key. The resulting continuation permits later contributions, successor roots and the next chain seed to be disclosed. The current root is an internal value. Each hybrid generates all outputs using that hybrid's root and subsequent key schedule.

Primitive experiments.

Let $F(k, m)$ be HMAC-SHA256 over one public quantum compression oracle $h : \{0,1\}^{512} \times \{0,1\}^{256} \to \{0,1\}^{256}$, in message-block, chaining-state order. All keys below have 32 bytes. Put

$$\Gamma(a) = \big(h(I(a), \mathrm{IV}), h(O(a), \mathrm{IV})\big),$$

where $I(a) = (a \| 0^{32}) \oplus \mathtt{36}^{64}$ and $O(a) = (a \| 0^{32}) \oplus \mathtt{5c}^{64}$. Initial auxiliary state is independent of h and the designated hidden uniform key. All oracle-dependent preparation is executed and counted, together with the complete continuation.

Write $\varepsilon_{\mathrm{Hjoint}}(C, T, q, L)$ for an output-bit gap bound between

$$\mathrm{Col}(a) = F(a, R), \quad \mathrm{Row}(m) = F(R, m) \quad \text{and} \quad V(\Gamma(a)),\ U(m),$$

where R is hidden uniform, U, V are independent private random functions, there are at most C classical column calls, T classical row calls of length at most L, and q ordinary compression queries, which may be quantum. Equal inputs are cached. Distinct salts with equal Γ share their column reply. The ideal here is the independent-table ideal, without a separately imposed cross-pair repair.

Write $\varepsilon_{\mathrm{Hrow}}(T, q, L)$ for the corresponding real-HMAC-row versus private-random-function bound for one uniform hidden key and *multiple adaptive messages*. This differs from the immutable report row experiment. Finally, $\varepsilon_{\mathrm{Hcol}}(q)$ bounds one distinct zero-salt evaluation $F(0, S)$, for hidden uniform S, versus an independent uniform reply; exact repeats reuse it. The retained compression-QROM arguments give admissible bounds for these operational experiments.

The finite continuation.

The initial root R is sampled independently of h and the initial auxiliary state, before any use. It may be presampled and withheld during counted preparation, which chooses the public session identifier s and transcript inputs. The previous root and old wrapping keys may be known. The current root has only the uses

$$P_0 = F(0, R), \qquad P_s = F(s, R), \qquad P_x^+ = F(R, S_x^+). \tag{11}$$

There are at most t_+ distinct 32-byte successor contributions S_x^+. They are supplied by known local computation, may be disclosed, and may depend on the available view. The controller may compute and disclose

$$R_x^+ = \mathrm{Expand}(P_x^+, I_x^+, 32)$$

and run arbitrary known subsequent computation from these values with the same public oracle. Repeated arguments remain consistent.

Let D and D_M denote the actual 40-byte rekey and 43-byte message domains. The three confirmation information strings are

$$\begin{aligned} I_r &= D\|\texttt{responder-confirmation}, \\ I_f &= D\|\texttt{final-confirmation}, \\ I_a &= D\|\texttt{receipt-confirmation}. \end{aligned}$$

Their keys are $\mathrm{Expand}(P_0, I_d, 32)$. The controller uses these keys in the real MAC, verification and signing algorithms. At most $M \geq 1$ distinct traffic information strings have the native form

$$\begin{aligned} I_T = D\|&\texttt{epoch-traffic/HKDF-SHA256/ChaCha20Poly1305/} \\ &\|\mathit{context}_{32}\|\,\mathrm{BE64}(e)\|\mathit{transcriptDigest}_{32}. \end{aligned} \tag{12}$$

It may be chosen after the response confirmation and its signature: the native transcript digest hashes the complete offer and response wires. Every repeated complete I_T reuses its entry, regardless of endpoint or wire representation. At the first allocation of one new string I^*, mark that entry and a direction $d^* \in \{1, 2\}$ before any of its blocks are produced. The choice may use the available prefix; a previously allocated entry cannot be reclassified as fresh. Later equal-info aliases inherit the same protection. This is a profile for a designated fresh traffic allocation, without a cross-endpoint wire-uniqueness premise.

The four 32-byte output blocks are

$$(s_{I,1}, s_{I,2}, a_{I,1}, a_{I,2}) = \mathrm{Expand}(P_s, I, 128).$$

The selected seed $s^* = s_{I^*,d^*}$, designated before allocation, has only the selected first-step use

$$\begin{aligned} I_M &= D_M\|\texttt{chain}\|\,\mathrm{BE64}(0), \\ P_M &= F(0, s^*), \qquad (s', K) = \mathrm{Expand}(P_M, I_M, 64). \end{aligned} \tag{13}$$

The opposite seed, confirmation keys, ACK keys, blocks at every distinct nonmarked information string and next seed s' may be disclosed. A repeated I^* remains the same protected entry. Real ACK computation and arbitrary candidate-ACK verification remain available. Subsequent chain computations use s' honestly. The selected key K has only its declared application encryption/decryption use. Current-root, PRK, old selected-seed and message-key bytes have no separate public return, comparison or hashing interface. Private persistence uses an already justified opaque interface; original candidate-local stops and projection accounting remain in force.

Lemma 5.9 (Adaptive information and fresh traffic blocks) *Replace each PRK-class HMAC row by an independent private random function, available only through HKDF-Expand. The resulting interface is exactly an independent block table indexed by the class, complete information string and block index, with consistent prefixes. This holds for adaptively chosen*

information strings. In particular, the first four traffic blocks are fresh at installation despite the confirmation dependence of I_T.

Proof For one private row W, the inputs are $I\|1$ and $T_{I,j-1}\|I\|j$ for $j \geq 2$, where each preceding block has exactly 32 bytes and the counter has one byte. Different counters separate inputs. For the same counter, deleting the counter and, when present, the fixed-width preceding block shows that equality requires equal I. This is true for every preceding output assignment. Equal (I, j) pairs retrieve equal inputs and replies inductively. New pairs therefore receive independent table blocks, even when I was chosen from prior outputs and public-oracle computation.

The confirmation information lengths are $62, 58, 60$ bytes and each traffic information length is 155, so their table entries differ even if the zero and session salts share a PRK. The marked I^* is also distinct from every earlier traffic entry. Importantly, this exact rewrite removes the internal parent-row inputs containing a preceding secret block. Because that row has no other interface, they leave no extra use of the selected traffic seed. Public quantum-oracle computation is unchanged. □

An available test uses the public view, event ledger, declared locally computable outputs and its independent test randomness. It cannot inspect the designated private roots, PRKs, seeds, message key or tables. For the application experiment below we can fix either challenge bit and then average; the test is not given a hidden primitive key.

Let G_R be this real schedule with uniform root. In G_{leaf} replace (11) by $V(\Gamma(\cdot))$ and $U(\cdot)$, replace each column-PRK Expand interface by Lemma 5.9's block table, and replace (13) by two independent cached blocks s', K. Real confirmation/ACK operations and successor-root derivations continue from the known replies; all signatures, checks and later outputs are regenerated within the current game.

Theorem 5.10 (A selected message leaf with joint continuation) *Let q_R, q_0, q_1, q_s, q_M bound the entire ordinary compression-query work of the five reductions below. Set*

$$\begin{aligned}\Delta_{\text{sched}} = \min\{1,\ & \varepsilon_{\text{Hjoint}}(2, t_+, q_R, 32) \\ & + \varepsilon_{\text{Hrow}}(3 + 4M, q_0, 188) + \varepsilon_{\text{Hrow}}(4M, q_1, 188) \\ & + \varepsilon_{\text{Hcol}}(q_s) + \varepsilon_{\text{Hrow}}(2, q_M, 89)\}.\end{aligned} \tag{14}$$

For every available terminal bit, the probability gap between G_R and G_{leaf} is at most Δ_{sched}. For a common available score in $[-1, 1]$, the expectation gap is at most $2\Delta_{\text{sched}}$. The second PRK-row term is omitted when the interface itself rules out a second class.

Proof First give every current-root column and row call to the joint primitive experiment. Their replies suffice to compute all Expands, confirmations, signatures, traffic and successor roots. In particular, disclosed successor contributions and roots require no raw current root. This gives the first term and retains the same public oracle throughout.

Next replace the at most two column-PRK rows in class-allocation order. Preallocate the zero-salt class privately if necessary; a second class exists only when $\Gamma(s) \neq \Gamma(0)$. For a designated class, install the row primitive's hidden uniform key as its first V reply, withheld until allocation. The identifier is computed from known salts and h; it does not inspect the private reply. Every equal Γ reuses this handle. Answer Expands by querying the challenge row

on the actual successive messages. Their output blocks are known, so real MACs, verifications, signatures and continuation can all be performed. Other class keys are sampled locally. All their real HMAC computations, and all preparation and future work, count as ordinary queries.

Three one-block confirmation Expands and at most M four-block traffic Expands give input lengths $63, 59, 61$ and $156, 188, 188, 188$. Thus the first class has at most $3 + 4M$ distinct row calls and a second class at most $4M$. A never-allocated class causes no challenge call. Their total is at most $3 + 4M$; the separate worst-case budgets avoid conditioning a primitive bound on the oracle-dependent class-equality event.

Independent V replies may coincide numerically. No inequality is assumed: known-key computations are performed through the same h, including accidental overlap with the challenge key, and are counted by the primitive experiment. An ideal row is the declared hybrid replacement. Intentional Γ aliases share that row in every game.

Apply Lemma 5.9 exactly. The selected traffic seed is now a fresh independent block, and its parent no longer exposes an input containing it. Presample the single-column challenge seed and install it at the newly marked entry before its blocks are produced. There is no earlier equal-info allocation; later aliases retain this protected handle. Its only raw use is $F(0, s^*)$; every other block is sampled locally. The next hop therefore makes P_M independently uniform at cost $\varepsilon_{\mathrm{Hcol}}(q_s)$.

Finally install the row challenge key at that fresh P_M. The two actual Expand inputs have lengths 57 and 89, giving the final term. The same exact block-table rewrite produces independent s', K. The next seed may be disclosed and used in all later chain computation; the reduction has every value it needs except the key reserved to the application interface. These five comparisons prove the bit bound. Randomizing a terminal bit with probability $(Z + 1)/2$ proves the score bound. Equivalently one may fix an independent application challenge bit, apply this bound, and average over that bit. □

The comparison installs an independent key and regenerates the complete continuation from it. Subsequent ciphertexts retain their dependence on that key, and the intermediate roots keep the joint distributions specified by the corresponding hybrids.

Application confidentiality.

Use a one-key, one-encryption left-or-right AEAD experiment with nonce 0^{96}, equal-length messages, common associated data and the same additional public oracle h. Its classical decryption interface returns real plaintext or failure, except that an exact challenge query returns an authenticated marker. Let $\varepsilon_{\mathrm{AEAD}}$ be its absolute guessing advantage over $1/2$. An encryption-only profile uses the corresponding one-use privacy experiment.

The honest application creates the selected send. The network adversary receives ciphertext, metadata, status, ACKs and real authentication/error outcomes, and can request plaintext-free Resume(id). The actual payload and its retained send-intent digest are opaque application handles. Delivery and any payload-bearing resolution output are private to the honest application; the public continuation has no additional payload-dependent return. An exact challenge reception can therefore update public state using its authenticated marker and known metadata. Other ciphertexts use the declared decryption interface.

The send-origin normalization caches the encryption reply at its first *private computation*, even if that computation precedes durable commit or network release. Native crash recovery may repeat encryption, but always on the same original tuple;

the reduction retrieves its cache. The cache includes the marked material's original allocation and authoritative owner/history, direction and index. Distinct material allocations are not merged merely because their epoch numbers agree.

Corollary 5.11 (First message of a recovered direction) *For an event A computable from the available view and ledger without reading the hidden AEAD challenge bit b, let $Z = \mathbf{1}_A(2\mathbf{1}[g = b] - 1)$. Under the application interface above,*

$$|\mathbb{E}_{G_R} Z| \le 2\Delta_{\text{sched}} + 2\varepsilon_{\text{AEAD}}.$$

Proof Theorem 5.10, applied for each fixed b and averaged, changes this score by at most $2\Delta_{\text{sched}}$. In G_{leaf} give the selected encryption/decryption operations to the AEAD challenger and sample every other block locally. Confirmation, ACK, successor-root and subsequent-chain computation uses known material and the same public oracle. Cache the first encryption response; exact challenge reception needs only the marker and metadata. At the end return g when A holds and a fresh fair guess otherwise. Its centered success is exactly $\mathbb{E}Z$, bounded by $2\varepsilon_{\text{AEAD}}$. □

The public application boundary is substantive. For two challenge plaintexts with different native send-intent digests, publishing the result of send_message(id, m_0, ad) after the challenge would distinguish a retained result from an intent conflict. No record acquisition is needed for that test. The network profile uses Resume while the honest application retains the plaintext-bearing API; original eligibility is unchanged.

Entry from a recovered contribution.

The uniform value above is the epoch root *after* root derivation. For a separately justified uniform combined contribution S, known previous root r and public response core, the actual derivation is

$$\begin{aligned} P = F(r, S), \qquad R = F(P, I_R \| \texttt{01}), \\ I_R = D \| \texttt{pending-root/HKDF-SHA256/} \\ \| \operatorname{hash}(\texttt{response-core}, \mathit{core}). \end{aligned}$$

Its information string has 97 bytes and its sole Expand message has 98. Suppose this is one original Extract/Expand, with exact repeats cached, and S, P have no additional raw-use or disclosure channel. The retained root-initialization lemma applies with the entire controller above as continuation, at bit cost Δ_{init}. Its ideal reply can be presampled and withheld, providing precisely the joint uniform root initialization required here. It regenerates confirmations, signatures, traffic and future epochs; it does not replace a root inside a fixed public transcript.

Let $\mathsf{W}^{[j]}_{\text{cont}}$ denote this uniform-combined-contribution continuation, with the same candidate-local score code Ψ_j, and put $b^{\text{cont}}_j = \mathbb{E}_{\mathsf{W}^{[j]}_{\text{cont}}} \Psi_j$. Consequently,

$$|b^{\text{cont}}_j| \le 2\Delta_{\text{init},j} + 2\Delta_{\text{sched},j} + 2\varepsilon_{\text{AEAD},j}.$$

The native accounting term $b_j = \mathbb{E}_{\widetilde{\mathsf{V}}^{[j]}} \Psi_j$ belongs instead to the random-component game, which may still use the real ContextBound combiner. If its actual complete comparison to $\mathsf{W}_{\mathrm{cont}}^{[j]}$ has bit gap at most $\varepsilon_{\mathrm{CB},j}$, then

$$|b_j - b_j^{\mathrm{cont}}| \le 2\varepsilon_{\mathrm{CB},j}, \qquad |b_j| \le 2\varepsilon_{\mathrm{CB},j} + |b_j^{\mathrm{cont}}|.$$

This comparison retains the controller, query interfaces and score/event code, counts all shared-oracle preparation and continuation, and recomputes events within each game's execution. Its actual tuple and disclosure conditions must hold. The randomized test with success probability $(\Psi_j + 1)/2$ gives the first inequality; the second is the triangle inequality. Native correspondence ρ_j, component and authentication terms, and original-game projection residuals remain in their established locations. This supplies the continuation bound for the specified profile, including known old roots and allowed later disclosure.

5.6 Online selection within a finite message chain

The same recovered direction supports selection after earlier messages have been processed. We retain the joint root, confirmation and traffic interfaces of Section 5.5 and give the selected direction a finite chain prefix. Earlier message keys may be disclosed; the selected message's successor seed may be disclosed after its derivation. The comparison tracks the seeds connecting those two observation boundaries.

Prefix interface and selection.

At the first allocation of a fresh traffic information string, designate its material identity and a direction as before. Fix a deterministic lifetime bound $1 \le L \le 2^{64} - 1$. With s_0 the designated traffic seed, the prefix has the native equations

$$\begin{aligned} I_i &= D_M \| \texttt{chain} \| \,\mathrm{BE64}(i), \\ P_i &= F(0, s_i), \qquad (s_{i+1}, K_i) = \mathrm{Expand}(P_i, I_i, 64), \quad 0 \le i < L. \end{aligned} \tag{15}$$

Complete material identity, direction and index identify one original step. Matching endpoint copies, failed computations and resumptions retrieve that step's original result. Distinct traffic materials keep their own histories, including when their epoch and index agree.

Immediately before each index's first private step computation, the controller either continues the prefix or marks that index. This choice uses its available history, including earlier disclosed message keys, messages, authentication outcomes and counted public-oracle computation. The mark is permanent and occurs at most once. Every computation at an alias participates in the same first-computation ledger; this includes transaction-private receive-side advancement through skipped indices. An index already computed at any such alias has passed its selection point.

Before a mark, the seeds s_i and Extract values P_i are private handles with only the uses in (15). The message keys of steps passed without a mark may be returned and used arbitrarily within the finite counted computation. At a marked step ℓ, K_ℓ has the single-original-encryption, classical-decryption and exact-retry interface of

Corollary 5.11. Its paired successor $s_{\ell+1}$ becomes an available value, from which subsequent real chain computation and disclosure may proceed. Private persistence uses the same justified opaque interface. If the execution ends without a mark, the selected application event and score are zero.

This timing separates observing earlier keys from choosing an already observed target. It also fixes when each protected seed can first be installed in a reduction. The root's other traffic entries, the opposite direction, real confirmation and ACK checks, and successor-root disclosure remain available as in the preceding joint continuation.

Let Δ_{traffic} denote the first three terms of (14), capped at one, with q_R, q_0, q_1 bounding the expanded controller's entire execution. They compare the joint root uses and the at most two traffic/confirmation PRK rows. For each $i < L$, let $q_{s,i}$ and $q_{M,i}$ be deterministic upper bounds on all ordinary compression queries in the corresponding complete column and row reductions. In particular, these budgets include the other real chain steps, known earlier keys, verification, subsequent chain computation, preparation and terminal testing. Define

$$\Delta_{\text{prefix}}(L) = \min\Bigg\{1,\, \Delta_{\text{traffic}} + \sum_{i=0}^{L-1}\Big(\varepsilon_{\text{Hcol}}(q_{s,i}) + \varepsilon_{\text{Hrow}}(2, q_{M,i}, 89)\Big)\Bigg\}. \tag{16}$$

The unused second traffic-row term is omitted under the same one-class condition as in Theorem 5.10.

Theorem 5.12 (Online selected chain leaf) *For the prefix interface above, compare the real uniform-root schedule with the experiment that replaces the joint root and traffic rows by their preceding independent tables and gives each prefix step through the mark two independent cached blocks (s_{i+1}, K_i). Computation strictly after the marked step is real and starts from its available successor seed. Every available terminal bit has probability gap at most $\Delta_{\text{prefix}}(L)$ between these experiments. Every common available score in $[-1, 1]$ has expectation gap at most $2\Delta_{\text{prefix}}(L)$.*

Proof The first three comparisons of Theorem 5.10 retain the new controller as their complete continuation. Their replies suffice to compute every chain locally, including keys returned before the mark. They therefore cost Δ_{traffic}. Apply Lemma 5.9 exactly to the traffic Expand table. The designated s_0 is an independent private block; its parent row has no remaining internal input containing that block.

Proceed in fixed index order. Before comparison i, every preceding step reached before or at the mark has its independent block-pair representation. If an earlier index was marked, step i is already in the real suffix and this comparison leaves that branch unchanged. Otherwise the preceding block table, or the traffic table for $i = 0$, allocates a private uniform s_i. At that allocation the controller has already decided whether the predecessor was marked. Hence the column reduction can install its presampled hidden seed exactly on the branch that continues the private prefix. On other branches it uses locally sampled values. It never needs to predict the later mark.

Until step i is first computed, its seed has only an opaque handle. All available observations, including K_{i-1} when $i > 0$, can be generated without its bytes. At the first computation, give $F(0, s_i)$ to the single-column challenger. Repeated computations use the same answer. Its returned P_i lets the reduction perform the real Expand, all later steps, marking decisions and the complete public continuation. If the execution never reaches this call, the challenge remains unused. Thus this comparison costs $\varepsilon_{\mathrm{Hcol}}(q_{s,i})$ without conditioning a primitive bound on either reachability or selection.

The fresh independent P_i is then installed as the hidden key of the row comparison. Its two Expand messages have lengths 57 and 89 bytes for every permitted index. The first reply supplies the input prefix of the second query, so the reduction can generate both outputs and the subsequent real continuation. This costs $\varepsilon_{\mathrm{Hrow}}(2, q_{M,i}, 89)$. Exact block-table normalization then removes the parent-row input containing s_{i+1}, giving an independent pair. If step i was marked, its successor is locally available and all later steps remain real. Otherwise that successor is the private handle for comparison $i + 1$.

This construction includes skipped-key computation and failed private attempts: their first step results are recorded before later publication or commitment. All aliases reuse the same table entry. Independent samples may coincide numerically; the comparisons preserve known-key computations through the common public oracle and charge their queries. No distinct-value conditioning is introduced.

The $2L$ comparisons use fixed worst-case budgets and recompute the mark and terminal test within each game. Their triangle inequality, together with the initial three comparisons, gives the bit bound. The randomized test with success probability $(Z + 1)/2$ gives the stated score bound. □

Corollary 5.13 (Selected message confidentiality after prefix disclosure) *Use the private-application and public-Resume boundary of Corollary 5.11, and let A be an available event independent of direct access to the hidden challenge bit b. Include existence of a mark in A, and put $Z = \mathbf{1}_A(2\mathbf{1}[g = b] - 1)$. For the declared one-key, one-encryption AEAD experiment,*

$$|\mathbb{E}Z| \le 2\Delta_{\mathrm{prefix}}(L) + 2\varepsilon_{\mathrm{AEAD}}.$$

For fresh authenticated installation under that marked message key,

$$\Pr[\mathsf{Bad}_{\text{selected msg}}] \le \min\{1, \Delta_{\mathrm{prefix}}(L) + \varepsilon_{\mathrm{INT}}\}.$$

Both primitive budgets include the entire adaptive selection and subsequent continuation.

Proof In the independent-pair experiment, a step is marked before its pair is first generated. All earlier keys and marking decisions are therefore available without K_ℓ. Install the AEAD challenger's key at that pair's message-key handle, sample its successor locally, and compute the real suffix from that successor. Previous message keys, confirmation keys and ACK keys are locally known. Exact retries reuse the original encryption response recorded before its durable outcome. The selected-key decryption interface supplies the same authenticated marker or real error/plaintext handling as in Corollary 5.11.

For confidentiality, return g on A and a fair guess otherwise; the centered success is the ideal experiment's $\mathbb{E}Z$. For integrity, a new selected-key inbox installation absent from the encryption ledger supplies the integrity experiment's winning decryption pair. Apply Theorem 5.12 to the respective score and Boolean event. This yields the two bounds. □

For $L = 1$, the comparison has the same five primitive terms as the selected first-message theorem. A fixed target index ℓ can use $L = \ell + 1$; online selection

uses a lifetime bound fixed before execution. The native chain source uses the stated 56-byte information strings, constant-width index encoding and zero AEAD nonce, advances a send index once, and reuses exact outbox bytes on Resume. Transaction-private receive advancement can compute skipped keys before commitment; failed-transaction resumption retains their first-computation identity in the proof ledger. The established native next-rekey results retain their first-message controller; the present theorem supplies the finite-chain continuation with its explicit selection and observation interface.

5.7 From an accepted ContextBound operation to the message continuation

We now discharge the ContextBound comparison left explicit in Section 5.5 for one concrete native profile. The protected component is ML-KEM-768 with a real-versus-uniform 32-byte shared-secret challenge. X25519 supplies its native output and rejection behavior. Write H for the whole SHA3-256 random oracle and h for the independent SHA-256 compression oracle of Section 5.5. Both remain available throughout every comparison. All oracle-dependent preparation, including construction of an established starting state, is executed and counted as part of the complete reduction.

Complete effective tuples.

For a byte string x, let $\mathrm{LP}(x) = \mathrm{BE64}(|x|)\|x$. The admitted native ContextBound computation is $S = H(\mathrm{enc}(T))$, where enc concatenates the LP encodings of the ten fields

$$T = (D_{\mathrm{CB}}, \mathit{suite}, \mathrm{BE32}(v), s_M, s_X, c_M, p_M, c_X, p_X, C_{\mathrm{policy}}).$$

Here

$$C_{\mathrm{policy}} = \mathrm{LP}(D_{\mathrm{policy}})\|\,\mathrm{LP}(\mathit{policyDigest})\|\,\mathrm{LP}(C_{\mathrm{app}}),$$
$$C_{\mathrm{app}} = H\big(\mathrm{LP}(D\|\texttt{response-kem})\|\,\mathrm{LP}(O)\big),$$

with O the entire retained signed offer and the policy digest that of the verified admitted policy. In the pinned suite the complete encoded input has 2639 bytes, and the 32-byte s_M field occupies offsets 76–107. The reservation's trust-root binding, original-slot identifiers, operation identifiers and journal identities are kept in the provenance ledger. The hash input consists of the ten fields above; the retained signed offer contributes through C_{app}.

Lemma 5.14 (Complete-tuple hidden-component comparison) *Let U be uniform on $\{0,1\}^{256}$, independent of the initially sampled oracles and initial auxiliary state. In a protected family replace one 32-byte component field by U, and let m contain all nine other actual fields, including the other component secret and the complete policy context. The private interface returns $H(f(U,m))$, where f is the actual ten-field encoding. Exact repeats are cached by m. Assume that the rest of the experiment accesses U only through this interface or opaque handles and respects its stated stops. For at most d distinct private inputs and Q_H ordinary queries to H, replacing the private interface by a random function $\mathcal{T}(m)$ changes every efficient output-bit probability by at most*

$$\varepsilon_{\mathrm{CB}} = \Phi_{256}(d, Q_H). \tag{17}$$

The inputs m may be adaptive. Q_H counts the complete derived experiment, including preparation, all ordinary KEM/hash work, continuation and terminal test. The same independent oracle h is retained.

Proof Length-prefix parsing recovers all ten actual fields. For fixed m and any complete string w, at most one $u \in \{0,1\}^{256}$ satisfies $f(u,m) = w$, so its point probability is at most 2^{-256}. If $m \neq m'$ as effective field tuples, then $f(u,m) \neq f(u,m')$ for every u. Cache repeats before invoking the private interface. These are precisely the point-probability and all-secret freshness conditions of the hidden-input theorem stated in Appendix B [15, Theorem 2]. It gives the displayed bound for adaptive private queries. The independent h interface is part of the unchanged auxiliary computation. Ordinary H queries, including an accidentally matching known computation, retain their actual oracle responses and are counted in Q_H. The comparison covers the unconditional experiment, including collisions and executions that end without acceptance. □

The cache key in this lemma cannot be replaced by an origin identifier. Two different original calls can have the same effective tuple: their real answers agree with probability one, whereas two independent per-origin ideal answers agree with probability 2^{-256}. Conversely, dropping the other shared secret or any effective field could merge distinct real hash inputs. Provenance and hash consistency therefore require separate tables.

Accepted original component calls.

Fix $j = (g, a)$ in the existing original-reservation ledger. The component adapter assigns its ML-KEM challenge public key only to generation origin g, retaining the locally sampled X25519 scalar, runtime admission and owner metadata. It obtains its challenge pair (c_M^*, K_{ch}) at the first actual private encapsulation computation of a on the original retained offer from g. It computes the X25519 leg locally and caches the whole result before subsequent persistence. Exact retries reuse that challenge; a fresh entropy reservation, including one after an unsuccessful attempt, receives its actual new original slot. Copies and reopenings do not create a new challenge invocation.

After challenge creation, a PQ decapsulation under the selected generation owner with ciphertext c_M^* returns the cached K_{ch}, regardless of accompanying traditional ciphertext or context. All other PQ ciphertexts are routed to the component decapsulation interface. The traditional decapsulation and its rejection path remain native; a rejected traditional leg produces no ContextBound call. The challenged PQ ciphertext is never submitted to the postchallenge CCA oracle. The operational component game first supplies its generated public key and permits an ordinary prechallenge decapsulation phase. A single challenge is requested at the specified first computation of a; it returns a fresh encapsulation ciphertext and either its real key or an independent uniform 32-byte key. Thereafter its decapsulation oracle rejects that challenge ciphertext. Before challenge creation, selected-key decapsulation follows this ordinary prechallenge interface; a query that happens to equal the later sampled c_M^* retains its original answer throughout the execution. Other original owners use their locally known keys, even when public-key bytes happen to coincide. The component costs below refer to this two-stage game.

The immutable admitted binding must show that a consumes the original offer from g. An immutable contrary binding makes attribution to j permanently impossible and gives a free stop U_j under Section 5.3. Provenance follows that original binding. Missing or unaccepted operations otherwise retain the native execution. Reveal and projection cases use the established free stops, charged stops and local record correspondence. Thus the existing component costs a_j^+, a_j^-, δ_j retain exactly their meanings. Correctness and permitted-projection costs remain at their established comparisons. More explicitly, the derived distinguisher D_j^σ runs the entire adapter and outputs $\mathbf{1}[\Psi_j = \sigma]$ for $\sigma \in \{-1, +1\}$. Its operational advantage is

$$a_j^\sigma = \left|\Pr[D_j^\sigma = 1 \mid \text{real challenge}] - \Pr[D_j^\sigma = 1 \mid \text{uniform challenge}]\right|.$$

The real branch matches $\mathsf{V}^{[j]}$ up to the selected-pair correctness cost δ_j; the uniform branch is exactly the declared $\widetilde{\mathsf{V}}^{[j]}$. This gives the two output-event inequalities used in Proposition 5.8.

In the resulting random-component game, $\widetilde{\mathsf{V}}^{[j]}$, $K_{\mathrm{ch}} = U$ is independent uniform. The reduction for Lemma 5.14 can materialize all actual native KEM tapes locally, execute their algorithms, and replace only the designated protected component returns by U. These locally generated tapes implement the adjacent pair of experiments. Ordinary SHA3 work inside those algorithms is included in Q_H; opaque records and the same stopping code retain their declared interfaces.

A concrete single-root source profile.

Use the next rekey of one established authoritative session. The two endpoints initially have the same established root r, confirmed epoch and control parent, with no pending selected rekey. The old root and wrapping keys may be known. The supported native call surface includes failed reservations and writes, process loss and exact resumption, adversarial public packets, old-epoch traffic, asynchronous final/receipt processing, the marked first-message interface of Corollary 5.11, and subsequent epochs. Selected private owners are used exclusively through this call surface; read-only copies remain passive. These are an initial state and call surface; the original eligibility predicate is unchanged.

Lemma 5.15 (Next-rekey root-use correspondence) *In this profile, fix $j = (g, a)$ whose immutable original encapsulation binding consumes the retained original offer from g for this next rekey. Let m^* be that original encapsulation's effective tuple: its challenged c_M^*, locally generated c_X^*, traditional secret, retained offer and admitted policy determine all nine fields. Every protected use of that same tuple feeds one pair (r, I_R), where I_R is the complete 97-byte native pending-root information string. They are one original Extract/Expand and exact repetitions. No raw combined-secret channel is added by retry, crash or completion. This invariant holds on every continuation of the matching binding, independently of later eligibility or attribution.*

Proof The retained generation token proceeds to a fixed signed offer. Its reconstruction checks the regenerated public material against the saved offer. At the responder the encapsulation

token is bound to the complete saved offer; a different offer conflicts, and completing or committed plans return their retained flights. While these tokens are pending, ordinary message operations do not change the established rekey root. The response prefix fixes the session, context, epochs, role, control parent and retained-offer digest. Equality of m^* fixes both complete ciphertext fields, hence the response core and I_R. Both endpoints therefore use their common old r and this same I_R.

Before exposing its response, the responder replaces the active encapsulation-token stage by the stage retaining the derived root. The proposer replaces the generation-token stage before root installation. The final/receipt installation sites are the live updates of the rekey root. A crash before a transition recovers the same token and old root; a committed transition recovers the retained result. Neither resumes an old selected token under a newly installed root. This invariant is preserved by every supported call, authoritative write outcome and recovery step. The source artifact records the corresponding pinned call sites and the complete induction. The statement concerns protected-family uses; it makes no claim that independently sampled numeric secrets never coincide. □

Fresh allocation and score semantics.

Mark this original target tuple m^* before its first actual private hash computation, at any protected alias call site and including a computation whose subsequent write fails. If a crash intervenes after component challenge creation, its cached ciphertext and locally determined traditional leg still identify m^* before any such first use. An earlier private use at any alias site already allocates this entry and cannot be ignored at a later durable acceptance or network release. Other protected tuples do not become the target. If the selected original a never supplies a usable target, the original candidate score is zero; no substitute tuple is chosen. In the private-random-function game, $S^* = \mathcal{T}(m^*)$ is a freshly allocated uniform 32-byte value; replays and aliases retrieve that same entry. It can be presampled and withheld at this allocation. This is an online allocation argument, not a claim of uniformity after conditioning on the future transcript. Other entries of $\mathcal{T}$ can be sampled by the controller. By Lemma 5.15, all selected uses of S^* are the one Extract/Expand required by Appendix C.3. The actual continuation is then the controller of Section 5.5, including its fresh traffic-entry and direction mark, confirmations, ACKs, successors, shared h work and private retained state.

Every hop runs the same efficient terminal code Ψ_j on its own actual trace. Its value is zero after the established local stops. The qualifying predicate in the message game must be computable from the available classical view and provenance ledger, as required by Corollary 5.11. A hidden challenge bit is not an additional input to that predicate. This condition matters: if the guess is always zero and a fictitious predicate accepts exactly when the hidden bit is zero, the qualified signed score has expectation 1/2 even for an otherwise bit-independent view. For the admitted predicate, making an independent fair guess on disqualified runs implements the full qualified score in the actual AEAD game; no unsupported zero-mean idealization is needed.

Theorem 5.16 (ContextBound connection for the next native rekey) *Use the concrete profile and adapters above, the local comparisons of Proposition 5.8, and the message interface of Corollary 5.11. Fix finite budgets before the experiment. For candidate j, let d_j bound its distinct protected effective tuples and $Q_{H,j}$ its complete ordinary H queries. Let $\Delta_{\mathrm{init},j}$, $\Delta_{\mathrm{sched},j}$ and*

$\varepsilon_{\mathrm{AEAD},j}$ have the existing whole-continuation meanings of Appendix C.3, Theorem 5.10 and Corollary 5.11. Define

$$B_j^{\mathrm{cont}} = \min\left\{1, 2\Phi_{256}(d_j, Q_{H,j}) + 2\Delta_{\mathrm{init},j} + 2\Delta_{\mathrm{sched},j} + 2\varepsilon_{\mathrm{AEAD},j}\right\}. \tag{18}$$

Then $|b_j| \le B_j^{\mathrm{cont}}$ and

$$\mathrm{Adv}_{\mathrm{msg}} \le \min\Bigg\{\frac{1}{2}, \frac{1}{2}\Big(\eta_{\mathrm{attr}}^{\mathrm{nat}} + \eta_{\mathrm{proj}}^{\mathrm{nat}} + \sum_{j\in J}\big(\rho_j + a_j^{+} + a_j^{-} + 2\delta_j + B_j^{\mathrm{cont}}\big)\Big)\Bigg\}. \tag{19}$$

The residuals are still evaluated in the one original native experiment. Each candidate's ideal continuation is compared separately.

Proof Apply Lemma 5.14 to $\widetilde{\mathsf{V}}^{[j]}$. Fresh allocation of m^* gives precisely the uniform-combined-contribution game $\mathsf{W}_{\mathrm{cont}}^{[j]}$ from Section 5.5; all other private table entries and the entire native controller remain present. For any $\Psi_j \in [-1, 1]$, a terminal randomized test that outputs one with probability $(\Psi_j + 1)/2$ converts a bit gap ε into a signed score gap at most 2ε. Thus

$$|b_j - b_j^{\mathrm{cont}}| \le 2\Phi_{256}(d_j, Q_{H,j}).$$

The native root-use lemma supplies the single-entry hypothesis for the existing $\Delta_{\mathrm{init},j}$ comparison. The joint schedule and AEAD continuation then give $|b_j^{\mathrm{cont}}| \le 2\Delta_{\mathrm{init},j} + 2\Delta_{\mathrm{sched},j} + 2\varepsilon_{\mathrm{AEAD},j}$. All derived distinguishers include the complete preparation and continuation; no oracle budget resets at a hop. Combine these inequalities, use $|b_j| \le 1$, and substitute in Proposition 5.8. □

The result supplies the previously explicit ContextBound comparison for the next-rekey instance. Its source requirement is expressed through the root-use descriptor, so extensions can retain the same tuple-level comparison while establishing their corresponding root-entry interface.

5.8 A native live-core instantiation

The next-rekey comparison has a concrete record cost on the authoritative live-session profile. Its application owns plaintext-bearing Send and delivery; the network interface includes exact Resume, active control and message processing, ACKs, status and public commitments. Storage follows the native immutable-image and pending-intent history through failed writes, process loss and reopening. Captures of old or unrelated records remain available. A complete post-cut target-session acquisition invokes the original eligibility stop. The selected successor-root and next-chain-seed projections use the ownership and phase checks of the joint continuation. The partial secret-field interface consists exactly of these two named, phase-checked projections; other record acquisitions return the declared complete objects. These observations and the private-application/Resume boundary define the core's complete call surface. The retained live-profile artifact records their source correspondence in Appendix A.6.

This core obtains all its replies through the record, primitive and application interfaces already constructed. Its ordinary calls expose public metadata or declared

available values. Consequently its charged projection event is empty: $B_j = \varnothing$ for each candidate and $\eta^{\mathrm{nat}}_{\mathrm{proj}} = 0$. The report composition theorem supplies a separate extension for cleanup and report interfaces with their declared projection events.

For candidate j, let M^{seal}_j bound fresh outer sealing invocations, including attempts whose subsequent commit fails. Let $d_{j,i}$ bound the distinct private image/intent digest inputs in row i, and let $Q_{j,i}$ bound the entire ordinary-H computation of its derived distinguisher. These are deterministic bounds fixed before execution and cover every corresponding comparison world. Define

$$\begin{aligned} \varepsilon_{\mathrm{rec},j} &= \min\left\{1, \sum_{i=1}^{M^{\mathrm{seal}}_j} \Phi_{192}(d_{j,i}, Q_{j,i}) + \delta_{\mathrm{seal},j}\right\}, \\ \delta_{\mathrm{seal},j} &= \min\left\{1, \binom{M^{\mathrm{seal}}_j}{2} 2^{-192}\right\}. \end{aligned} \tag{20}$$

The second term is the private-wire structural-equality comparison of Corollary 5.3; it is zero when that equality interface is absent. Copies and exact pending-target application retain their original sealing handles. Atom-free rows have $d_{j,i} = 0$.

Corollary 5.17 (Next-rekey confidentiality on the live core) *Use the native next-rekey and marked first-message interfaces of Theorem 5.16 with the live core above, its recorded source correspondence and finite whole-execution budgets. Then*

$$\begin{aligned} \mathrm{Adv}_{\mathrm{msg}} \le \min\Bigg\{ &\frac{1}{2}, \frac{1}{2}\Big(\eta^{\mathrm{nat}}_{\mathrm{attr}} \\ &+ \sum_{j \in J} \left(2\varepsilon_{\mathrm{rec},j} + a^{+}_j + a^{-}_j + 2\delta_j + B^{\mathrm{cont}}_j\right)\Big)\Bigg\}. \end{aligned} \tag{21}$$

Here B^{cont}_j is the explicit ContextBound, root-initialization, joint-schedule and AEAD score cost in (18)*. Attribution is evaluated in the original native experiment.*

Proof The live-profile representation invariant follows the original reservation, saved-version and sealing handles through each native transition. On its stopped interface, the materialized record machine has exactly the native outputs and storage outcomes. Corollary 5.3 and Appendix B therefore give $\rho_j \le 2\varepsilon_{\mathrm{rec},j}$. The core's complete captures use the existing acquisition stop, and its remaining replies use the declared interfaces, so its charged projection event is empty. Substitute these two facts into (19). □

Source selection and attribution.

The attribution term distinguishes two actual interfaces. In the general experiment, the selector may name any original reservation handles with the required session, epoch and role metadata. For example, a generation reservation g_1 may be installed in a transaction-private `Plan::Key` whose first persistence fails before a complete pending target or image becomes durable. Reopening the prior authoritative state permits a new reservation g_2, followed by an honest accepted exchange using (g_2, a_2). Choose

$c < \mathsf{res}(g_1) < \mathsf{res}(g_2)$ and $c < \mathsf{res}(a_2)$, with the target signing keys and designated entropy unexposed and with no acquisition of a post-cut target-session version. A selector naming (g_1, a_2) can pass the reservation-event and recovery-cut checks while failing usable attribution. Every control signature and transcript hash in this trace is honest. This execution is already charged by $\eta^{\mathrm{nat}}_{\mathrm{attr}}$ in the general theorem.

Section 5.10 defines an executed-slot selector that registers g, a before their first private KEM computations return. A retained-plan induction preserves those original slots, and the complete signed-body and wire bindings associate an accepted allocation with them. Its signature experiment retains original signing randomizers and exact resumptions. That selector receives the explicit signature/hash bound in (25); the general metadata selector in (21) retains its original attribution term.

5.9 Confirmation and first-message authentication

The same native rekey uses the three confirmation keys and directional traffic material of Theorem 5.10. We now retain each confirmation key behind its actual MAC and verification interfaces and derive a joint authentication consequence. Public queries to the independent oracles H and h may be quantum. Protocol requests, MAC inputs, verification requests and their replies are classical.

Native transcripts and application interface.
Write $\mathcal{H}_L(x) = H(\mathrm{LP}(D\|L)\|\,\mathrm{LP}(x))$, with LP the source's eight-byte length prefix. For $d \in \{r, f, a\}$, let C_d be the responder, final or receipt confirmation key. The tag is

$$\begin{aligned} t_d &= F(C_d, \mathcal{H}_{L_d}(X_d)), \\ (L_r, L_f, L_a) &= (\texttt{response-core}, \texttt{final-core}, \texttt{receipt-core}). \end{aligned} \tag{22}$$

The complete responder core contains the control tag, profile, context, session, current and target epochs, sender role, control parent, digest of the complete signed offer, and hybrid ciphertext. The final core contains its corresponding control prefix, the offer and response wire digests, and the old sending-chain length. The receipt additionally contains the complete final-wire digest. Their respective source lengths are 1305, 225 and 257 bytes. Every MAC input in (22) has 32 bytes.

The native verifier checks the complete prefix, body width and identity signature before the relevant MAC verification. The public result follows the actual success, failure, persistence and status transitions. The MAC comparison may therefore give the reduction the expected tag internally; the verification caller obtains the native result. Every honest tag construction enters a ghost origin ledger at its first private computation, before a subsequent write or release. Verification computations have their own query ledger. Exact repetitions and matching endpoint aliases use the same role-key identity and the same original construction entry. A different signature encoding of an already logged core retains that core's origin.

For the marked traffic allocation and direction, use the first-message interface of Corollary 5.11. The encryption nonce is 0^{96}, and the associated data is

$$D_M\|\texttt{aead}\|\mathit{header}_{93}\|\,\mathrm{BE16}(|ad|)\|ad.$$

The header binds session, direction, epoch, index, message identifier and plaintext length. The first private encryption is recorded by its exact $(AAD, ciphertext)$ pair and original send identity. Retries retrieve its cached result through Resume. Plaintext delivery belongs to the private application. A fresh authenticated receive means a successful AEAD check followed by installation of the corresponding native inbox entry. The source's `PreviouslyConsumed` result is a historical consumption status. An existing inbox entry retains its original authenticated receive-intent digest and plaintext; a cached return refers to that earlier receive.

Let $\mathsf{Bad}_{\mathrm{conf}}$ mean that, before the declared local stop, a successful native confirmation verification accepts a complete core with no earlier honest private tag construction under that role-key identity. Let $\mathsf{Bad}_{\mathrm{msg}}$ mean that a fresh authenticated receive under the selected first-message key installs a pair absent from its honest encryption ledger. These definitions include failures and unaccepted executions in the probability space.

Operational bounds.

Fix a role order r, f, a. After exact-input caching, let T_d bound the distinct 32-byte messages evaluated under C_d by both tag construction and verification, and let V_d bound verification attempts. Repeated guesses count in V_d. All budgets are fixed before execution. Define

$$\Delta_{\mathrm{conf}} = \min\left\{1, \sum_{d\in\{r,f,a\}} \varepsilon_{\mathrm{Hrow}}(T_d, q_d, 32)\right\}, \qquad V = V_r + V_f + V_a. \tag{23}$$

Here q_d counts every ordinary h query of that row's complete distinguisher. In particular, it includes preparation, the target application, all other real primitives and terminal testing. If B_c bounds all such work outside these three confirmation MAC rows, the conservative choices

$$q_d = B_c + 4\sum_{e>d} T_e$$

are valid: earlier rows are already ideal, later rows remain real, and a 32-byte HMAC evaluation uses at most four compression calls including its two pads. The row lemma separately charges the replaced row's private evaluations.

Let $\chi_H(Q)$ bound production of two distinct complete strings with equal H output by a Q-query algorithm. It counts the entire derived experiment's ordinary and honest H evaluations, including preparation and the final collision test. For the random oracle with 256-bit output, an admissible bound is

$$\chi_H(Q) = \min\{1, C_{\mathrm{col}}(Q+2)^3 2^{-256}\},$$

with the universal constant of the quantum collision theorem [28, Theorem 7]. The operational notation retains that constant.

Finally, let $\varepsilon_{\mathrm{INT}}$ bound a quantum adversary's success in the one-key, at-most-one-encryption ciphertext-integrity experiment with nonce 0^{96}, classical decryption and the same independent public oracles. A successful decryption of a pair absent from

the encryption ledger wins. All repeated decryption attempts are included in its query budget. The encryption ledger begins at the first private computation.

Proposition 5.18 (Joint confirmation and message authentication) *Use the controller of Theorem 5.10 and the above native interfaces. Each C_d is accessed through its tag and verification service; the selected message key is accessed through its single original encryption, classical decryption and exact-retry service. Then*

$$\Pr_{G_R}[\mathsf{Bad}_{\mathrm{conf}} \vee \mathsf{Bad}_{\mathrm{msg}}] \leq \min\{1, \Delta_{\mathrm{sched}} + \Delta_{\mathrm{conf}} + V2^{-256} + \chi_H(Q_H^{\mathrm{conf}}) + \varepsilon_{\mathrm{INT}}\}. \tag{24}$$

The same controller retains the qualified confidentiality bound $|\mathbb{E}_{G_R} Z| \leq 2\Delta_{\mathrm{sched}} + 2\varepsilon_{\mathrm{AEAD}}$ of Corollary 5.11. Its confirmation and verification operations are part of that complete continuation.

Proof Apply Theorem 5.10 to the union event's terminal bit. This costs Δ_{sched} once and provides the joint independent-key experiment G_{leaf}. All transcript strings and subsequent information strings are generated in the current world.

For confirmation, replace the three HMAC rows in the fixed role order by independent private random functions W_d. The adaptive classical row lemma applies to every tag and verification evaluation; its ordinary query budgets include the remaining real rows. This costs Δ_{conf}. Independent key samples may coincide numerically; the real rows retain the shared compression oracle, and the row comparisons account for those coincidences.

Consider a second execution that returns failure on every verification whose digest has no earlier honest tag-construction entry under that role. The coupling stops at the first internal MAC equality on such an unissued digest, before subsequent persistence or a public response. Up to this first hit, the two executions have the same classical replies and induce the same transformation of the adversary's quantum private state. In the second execution, the unissued value $W_d(m)$ is independent of the available prefix: its verification result is fixed to failure and its table value remains private. Each of at most V attempted tags equals this value with probability 2^{-256}. A union bound therefore bounds the first fresh-digest success by $V2^{-256}$. It also covers repeated guesses for the same digest. If a fresh complete core instead has a digest already used by an honest construction under that role, the two different framed core strings give a collision for H. The reduction records these classical inputs and outputs the pair; its probability is bounded by $\chi_H(Q_H^{\mathrm{conf}})$.

For message integrity in G_{leaf}, sample every key except the selected message key locally, preserving the actual confirmation, ACK, successor-root and subsequent-chain computation. Forward the selected original encryption and every selected-key decryption to the integrity challenger. Exact retries use the first private encryption cache. A new authenticated inbox installation without an encryption origin entails a successful new decryption pair, so its probability is at most $\varepsilon_{\mathrm{INT}}$. Existing inbox entries and historical consumed statuses use their original metadata. The union bound in G_{leaf} proves (24).

For confidentiality, the AEAD reduction samples the confirmation keys locally and executes their actual verification services. Thus Corollary 5.11 applies with its unchanged complete-continuation budget. For an independent application challenge bit, both arguments may first fix that bit and then average. □

The native composition in Section 5.12 precedes this uniform epoch-root interface with the record, component, complete-tuple and root-initialization comparisons.

Each comparison uses the selected authentication event's terminal bit and complete continuation budget, connecting the present bound to the original recovery experiment.

5.10 Executed sources and signing reservations

The attribution question can be resolved for an independently specified source selector. Use the established-session controller with one authoritative history per endpoint for the selected session. Its key is the actual signer, session and epoch, including the admitted context and control parent. Passive copies, exact pending-target reconciliation and reopening preserve that history. This premise is substantive: the journal identity enters an operation scope but is absent from the signed control prefix, so two independently active journals with identical session coordinates would require an additional ownership rule.

For each local epoch, register its original generation or encapsulation slot immediately before the first actual private KEM computation. The registers ExecG and ExecA contain the existing reservation indices in $[G]$ and $[C]$; the sampling events and their cut tests are unchanged. Exact computation retries retrieve the register. The *executed-slot selector* chooses the logical challenge send and cut as before, then uses these two registered handles. A missing register fails the existing reservation-event test. Registration reads retained plan identity and public bindings, independently of future acceptance.

Lemma 5.19 (Stability of executed source slots) *Each authoritative local epoch has at most one executed generation origin and one executed encapsulation origin. Failed later writes, process loss and exact resumptions preserve these origins and their complete bindings.*

Proof Generation follows successful retention of `Plan::Key`; encapsulation follows successful retention of the exact offer and `Stage::Encapsulation`. A first-retention error exits before the KEM call. If its target became durable, recovery obtains that token; if no target became durable, the unused reservation can be followed by a new draw. Once a KEM call is reached, its original token is already authoritative. Subsequent phases retain that token or its result and root. A failed phase update recovers an earlier phase with the same source. The responder also rejects a different complete offer while its plan exists. Completion clears the plan while advancing the epoch; terminal retirement ends the session. Induction over these transitions establishes both invariants without requiring distinct public-key or ciphertext bytes. □

An honest signing ledger records the actual key, purpose, complete body and originating retained plan at the first private signing invocation, before any subsequent write or release. The signed string is

$$M = \mathit{signatureContext} \| \mathit{purpose}_1 \| \,\mathrm{BE32}(|B|) \| B.$$

The ledger includes every honest invocation under the actual key across public purposes and exact owner aliases. Absence from this actual-key ledger implies absence from its projected owner-index query ledger. Let $\mathsf{Bad}^0_{\mathrm{sig}}$ be a successful target-key verification of a string absent from this ledger, before the first target signing-key or signer-file disclosure. An alternative valid signature of an issued body is an issued-message

alias. Let Bad_H be a collision between distinct classical framed SHA3 inputs used by the association checks. Write $\chi_H(Q_H)$ for the complete-query collision bound from Section 5.9, counting preparation, honest computation and witness extraction.

Proposition 5.20 (Source association for the executed-slot selector) *For the controller above and the unchanged recovery eligibility E, usable source association A satisfies*

$$E \wedge \neg A \implies \mathsf{Bad}^0_{\mathrm{sig}} \vee \mathsf{Bad}_H.$$

If $\varepsilon_{\mathrm{sig}}$ bounds the fresh-body event in the complete signing execution stopped before returning a target key or signer-file capture, then

$$\eta^{\mathrm{exec}}_{\mathrm{attr}} \leq \Delta_{\mathrm{attr}} := \min\{1, \varepsilon_{\mathrm{sig}} + \chi_H(Q_H)\}. \tag{25}$$

Proof Eligibility excludes both target-identity disclosures throughout the execution. A comparison stopped immediately before the first such return has the same prefix, including the adversary's quantum-state transformation. Retaining whether a fresh-body verification has already succeeded gives exactly the event $\mathsf{Bad}^0_{\mathrm{sig}}$.

Consider a first proposer traffic installation before either bad event. Its retained offer O and regenerated private owner carry the unique registered generation g. A verified response body has an honest signing invocation. Its complete context, session, epoch, role and parent locate the responder's same local epoch, hence its unique encapsulation a. The checked offer-wire digest equals the digest of the offer retained by a. Collision freedom for these two actual framed inputs makes the offered wires equal. The response therefore carries a's ciphertext on g's public key with the exact KEM context used by the proposer. Native regeneration checks the public key against the saved offer before decapsulation and root confirmation. Source attribution identifies these original calls and bindings; their component-output correctness keeps its separate δ_j term.

A different valid outer response signature can preserve the same body. It keeps the source pair and root inputs, while the presented complete response supplies its actual traffic information and final-flight binding. The joint schedule already retains these information aliases.

A responder installs traffic after checking a final against both its saved offer and ready response. The final's honest signing invocation comes from a proposer transition of the preceding kind. Its two checked wire digests give the same complete offer and response in the absence of Bad_H, and hence the same g, a. Receipt acceptance makes no additional root allocation. Cached returns and reopened plans retain their earlier source association by Lemma 5.19. All association tests use available ledger identities and public strings, so the component adapter can evaluate them with an unknown selected tape. This proves the implication; the probability bound follows by a union bound. □

The failed-first-retention trace in Section 5.8 keeps its original two draws. Its executed generation register contains g_2 because g_1 never reaches computation. Thus the new selector removes that mismatch while preserving the general metadata-selector experiment and its residual.

The signing-coin interface.

Fix a signature scheme whose internal signing procedure is deterministic given key, message and randomizer. A reservation-aware fresh-message experiment provides classical Reserve, Open, Execute and Verify calls. Reserve samples the original randomizer

and binds key, operation, purpose and complete body. Open returns that binding and randomizer at any time. Execute uses the original tuple, records the message at its first private computation and caches the exact signature; retries reuse it. Verify wins on a valid message absent from the corresponding signing-query ledger. The controlled native calls use their retained original tuples, including the proposer's recomputed-body equality check; they expose no arbitrary tuple-retarget operation through the hashed reservation binding.

The owner interface begins with provisioned keys and supplies public keys, public identity and lifecycle metadata, signing reservations and signatures. It stops before a target-key or signer-file return and supplies no earlier byte-dependent observation of that secret owner file. This keeps the owner's serialized key seed outside the simulation, while every permitted signing-reservation record is represented exactly. Additional owner-file observations require their own correspondence cost.

For U independently generated honest signing owners, use the usual per-owner-index message ledgers. Reopening or reusing an owner preserves its index. Accidental public-key coincidences remain in the distribution, and the native admission checks are performed normally. All query and time bounds cover the complete wrapper. The native pair reserves only ML-DSA's 32-byte randomizer; its P-256 component uses deterministic RFC 6979 signing and low-s canonicalization [23, 25]. The pinned implementation routes `try_sign` to RFC 6979 with empty additional data [27]; the ordinary P-256 game below uses exactly this native signing distribution.

Proposition 5.21 (Signature component adapters with retained randomizers) *For this owner interface, the native pair's reservation-aware fresh-message success probability satisfies*

$$\varepsilon_{\mathrm{pair,res}} \le \varepsilon^{\mathrm{mu}}_{\mathrm{MLDSA,res}}.$$

For classical adversaries it also satisfies

$$\varepsilon_{\mathrm{pair,res}} \le \varepsilon^{\mathrm{mu}}_{\mathrm{P256,EUF}} \le \min\{1, U\varepsilon^{\mathrm{su}}_{\mathrm{P256,EUF}}\},$$

with the complete simulation time and at most the original Execute-call budget of signing queries in the single-user reduction.

Proof For the first inequality generate the P-256 keys locally and route each ML-DSA reservation and execution to its reservation-aware challenger. Privately Open each original reservation before materializing its record, counting these calls, and disclose the known randomizer through the native observation interface when requested. Fixed-context ML-DSA signatures come from the challenger; deterministic P-256 signatures and canonicalization are local. A fresh accepted pair projects to a fresh ML-DSA message forgery.

For the second inequality generate ML-DSA keys and every reserved randomizer locally. Obtain the P-256 component from the native deterministic signing challenger at first private computation, and cache the pair by original reservation. All reservation records and openings use known bytes. A repeated body with new randomness can have a new ML-DSA signature while its deterministic component repeats. A fresh accepted pair yields a P-256 forgery on the same unqueried message. Both simulations are exact. Guessing one of the U owner indices gives the displayed single-user bound; other owners are simulated locally. □

FIPS 204, Algorithm 7, derives internal masking randomness from a secret key seed, the external randomizer and the message representative [23]. The reservation-aware ML-DSA game records this external-randomizer interface. Section 5.11 connects it to ordinary native unforgeability in the shared random-XOF model. The ordinary P-256 inequality gives the classical instantiation. The following separation explains the role of the reserved-coin interface.

Proposition 5.22 (A separation induced by signing-coin disclosure) *Let Σ_0 be a deterministic EUF-CMA signature scheme whose secret key encoding s has fixed length L. Define*

$$\mathrm{Sign}_+(s, M; r) = \big(\mathrm{Sign}_0(s, M), r \oplus s\big), \qquad r \leftarrow \{0,1\}^L,$$

and verify the first component and the appendix's fixed length. Its ordinary fresh-message game has exact transcript-preserving reductions to and from that of Σ_0, with their simulation time counted. One opened signing reservation and its signature reveal s.

Proof For every fixed secret and signing transcript, the appendix of each fresh call is independent uniform. An ordinary reduction simulates it with fresh uniform bytes and projects a fresh forgery to Σ_0; conversely a base forger ignores the appendices. In the reservation experiment, XOR of the original randomizer and its appendix recovers s, enabling a fresh signature with the base scheme's correctness probability. □

This generic witness distinguishes the experiments; its signing construction is separate from the native ML-DSA algorithm. For the executed-slot controller with the declared owner interface, (25) replaces the attribution residual in Corollary 5.17, using the appropriate component advantage from Proposition 5.21. The record, KEM-correctness and joint-continuation terms retain their existing positions. The source map, full proof and finite fault/alias checks are recorded in Appendix A.10.

5.11 From signing reservations to ordinary unforgeability

The retained randomizer can be incorporated into an ordinary-signature reduction. The simulator preserves every opened reservation and uses the complete pair of original randomizer and public message representative to identify repeated computations. For native ML-DSA, a shared-oracle comparison also accounts for the dependence created by key generation.

Use the owner interface of Proposition 5.21. A fresh owner generation receives a family index; exact references and imported copies retain that index. Independently generated owners remain separate when their public keys coincide. The family ledger records each original message at its first private invocation, including failed publication. It is contained in the native actual-key ledger. Stopping before a target-family key or signer-file return therefore gives an auxiliary fresh-message event containing the native stopped event. All routing and stopping decisions use declared lineage and public metadata, without testing hidden secret equality.

Independent hedging.
Suppose signing factors as

$$\mathrm{Sign}(s, K, M; r) = S\big(s, d_{pk}(M); F_K(r, d_{pk}(M))\big),$$

where the complete computation after the public descriptor $d_{pk}(M)$ is deterministic given its internal tape, and K is generated independently of the remaining owner state. The ordinary signer draws a fresh hidden $r \in \{0,1\}^n$ per call. The reservation interface reveals original randomizers before or after execution and retains the complete original-message ledger. The joint PRF experiment supplies the same public oracle access as the signature game, including quantum queries when present; its private function calls are classical and its ideal random function is independent of the public companion oracle.

Proposition 5.23 (An ordinary-signature bridge for independent hedging) *For one protected family and at most q distinct executed $(r, d_{pk}(M))$ pairs, there are an ordinary wrapper-EUF adversary B and joint-PRF distinguishers D_0, D_1 such that*

$$p_{\mathrm{res}}(A) \le p_{\mathrm{EUF}}(B) + \mathrm{Adv}^{\mathrm{jprf}}_F(D_0) + \mathrm{Adv}^{\mathrm{jprf}}_F(D_1) + \beta_n(q), \qquad \beta_n(q) = \min\{1, \tbinom{q}{2} 2^{-n}\}. \quad (26)$$

Each primitive interface receives at most q calls. The bounds include the complete simulation, public descriptor work and owner observations.

Proof Replace the reservation signer's hedging function by a private random function, using independence to simulate the remaining key state. This costs the first joint-PRF advantage. Simulator B samples all visible reservation randomizers and caches signatures by the full pair $(r, d_{pk}(M))$. A new pair causes one ordinary signing query on its original M; every first slot execution enters its original message in the local ledger, including a failed publication. A repeated pair reuses the cached signature.

Replace the ordinary challenger's hedge by a private random function. Until its fresh hidden input repeats, its internal tapes are independent uniform values and can be coupled with the tapes of the distinct visible pairs. The t-th hidden draw repeats a prior input with probability at most $(t-1)2^{-n}$, giving $\beta_n(q)$. This charge concerns the ordinary challenger's draws. Arbitrarily many opened reservations and adaptively selected visible collisions are represented exactly by the pair cache. Returning the ordinary challenger to its real hedge costs the second joint-PRF term. Its signing queries are a subset of the executed-message ledger, so a fresh output is an ordinary forgery. Descriptor aliases need no separate charge: factorization gives the same signature, while the ledger retains every original message. □

For several protected families with a global stop before any of their key/file disclosures, the same argument uses indexed multiuser primitive games and the birthday term $\min\{1, 2^{-n} \sum_h \binom{q_h}{2}\}$, with no extra user factor. Permitting other-owner disclosures instead uses the selected-family construction below. All displayed query caps hold across the relevant complete hybrids; an Execute-call cap is always a valid conservative choice for q.

The shared ML-DSA oracle.

FIPS 204 key generation expands a uniform 32-byte seed into consecutive blocks (ρ, ρ', K) of lengths $(32, 64, 32)$ bytes. Signing computes

$$\mu = H(tr\|M', 64), \qquad c = H(K\|r\|\mu, 64).$$

The remaining signing computation uses the other key components, μ and c [23]. The pinned native backend passes its retained 32 bytes to `try_sign_with_seed`; the checked implementation uses K only for this hedge in the signature computation. The full source map records the fixed crate archive and its upstream commit.

Model SHAKE256 as a public random XOF H, with one independent infinite output string per complete input and consistent replies to every requested prefix. Other public algorithms retain the same interfaces in the comparison and are independent of the sampled random-XOF table. For a selected ML-DSA-65 family, let q_*^{res} and q_*^{ord} bound all remaining calls that can address a 34-byte or 128-byte input in the two complete experiments, after rewriting every designated target seed expansion and hedge as in the proof. Count public adversarial queries, preparation, verification and all other-owner work. Define

$$\delta_{\mathrm{C}}(v) = \min\{1, v2^{-256}\}, \qquad \delta_{\mathrm{Q}}(v) = \min\{1, 2v2^{-128}\}.$$

Theorem 5.24 (Native ML-DSA reservation bridge) *For the declared owner interface and a selected native ML-DSA-65 family, let q bound the distinct executed target (r, μ) pairs in every relevant experiment. There is an ordinary hedged-ML-DSA EUF adversary B with at most q signing queries such that*

$$p^H_{\mathrm{res}}(A) \le p^H_{\mathrm{EUF}}(B) + \beta_{256}(q) + \delta_{\mathsf{X}}(q_*^{\mathrm{res}}) + \delta_{\mathsf{X}}(q_*^{\mathrm{ord}}). \tag{27}$$

Here $\mathsf{X} = \mathsf{C}$ denotes the classical random-oracle model and $\mathsf{X} = \mathsf{Q}$ the quantum random-oracle model. Reservation and signing calls are classical in both. Non-target owners may be generated locally and exposed; the target interface stops before key or signer-file return and has no earlier secret-byte observation. Both sides use the same ordinary native algorithm and public-oracle model, with their full simulation time and oracle budgets.

Proof Sample independent uniform ξ, ρ, ρ', K of the specified lengths, a random XOF H_0, and a private random function $R : (r, \mu) \mapsto \{0,1\}^{512}$. Let

$$g = \xi\|\,\mathrm{byte}(6)\|\,\mathrm{byte}(5), \qquad \mathcal{S}_K = \{K\|r\|\mu : \ |r| = 32,\ |\mu| = 64\},$$

where these lengths are in bytes. Define H_* by replacing the first 128 output bytes at g with $\rho\|\rho'\|K$, and the first 64 bytes at each $K\|r\|\mu$ with $R(r, \mu)$. Preserve every remaining output bit from H_0. The two input sets have disjoint lengths 34 and 128. Conditional on ξ and the keygen output, all other full XOF strings remain independent uniform, including those in the selected slice. Thus H_* and the keygen blocks have exactly the native joint distribution.

Rewrite every target-lineage seed expansion, including reopened owners, to use those sampled blocks, and every target hedge, including retries, to use R. These are exact substitutions under H_*. Change every remaining oracle call to H_0. In this ideal experiment the query-register distribution is independent of the now-unused ξ, K. The remaining base state may depend on ρ, ρ', H_0, and every non-target owner keeps its original algorithm and shared-oracle calls.

For any ideal query, its expected mass on $\mathcal{D} = \{g\} \cup \mathcal{S}_K$ is at most 2^{-256}: each relevant length sector has that hidden-point or hidden-prefix probability, and the two sector masses sum to at most one. A classical identical-until-hit argument gives $q_* 2^{-256}$. For a quantum query, the difference between the two XOR oracle actions has norm at most $2\sqrt{p_t(\mathcal{D})}$. Telescoping using the ideal prefix states and applying Jensen's inequality gives expected final-state distance at most $2q_* 2^{-128}$, which also bounds the difference in a final event probability [5]. This comparison includes verification.

Apply it first to the reservation experiment. In the resulting ideal experiment R is independent of the public oracle, so the pair-cache coupling from Proposition 5.23 gives the ordinary ideal experiment with loss $\beta_{256}(q)$. Apply the same oracle comparison in the reverse direction to the complete ordinary experiment. Its signer is now the native ordinary challenger, giving (27). The actual reduction B uses only that signing oracle and public interfaces; the statistical comparison does not require it to know or program ξ or K. □

The source inventory makes the query budget concrete. After the designated rewrites, native ML-DSA-65 uses SHAKE256 input lengths 66 for secret-vector and mask expansion, 48 for challenge-to-ball sampling, 832 for commitment hashes, and 1952 for the public-key hash. Matrix expansion uses SHAKE128. The protocol's fixed context is 42 bytes; its complete descriptor input has length $155+|B|$ bytes. These target computations therefore agree exactly under H_* and H_0, including every rejection-loop iteration. Arbitrary public SHAKE256 calls and non-target seed/hedge calls remain in q_*. Safe-domain calls are common operations in the query hybrid. A finite computation cutoff followed by a limit proves the same statistical bound for unbounded safe-domain rejection loops; full running time remains part of the ordinary-EUF term.

Adaptive owners and the pair signature.

For at most U fresh owner families, guess the target's creation index uniformly, retaining that index across aliases. Embed the single-user challenger there and generate all other owners locally, answering their permitted disclosures. This works for adaptive creation and gives

$$\varepsilon_{\mathrm{pair,res}} \le \min\left\{1, U\left(\varepsilon^{\mathrm{su}}_{\mathrm{MLDSA,EUF}} + \beta_{256}(q) + \delta_{\mathsf{X}}(q_*^{\mathrm{res}}) + \delta_{\mathsf{X}}(q_*^{\mathrm{ord}})\right)\right\}. \tag{28}$$

The ML-DSA oracle ledger is a subset of the selected native pair owner's ledger, hence of its actual-key ledger. No aggregation of distinct pair owners with coincident ML-DSA components is assumed. Their original generation and disclosure distribution is preserved. Cross-owner seed or hedge-prefix coincidences are among the remaining shared-oracle queries already counted in the selected-target comparison.

Equation (28) supplies an ordinary-ML-DSA instantiation of the signature term in Proposition 5.20, preserving its record, KEM-correctness and joint-continuation accounting.

5.12 Native authentication for the selected recovered root

The confirmation and first-message authentication bound can be transported through the native execution by comparing its own terminal event bit. This gives an end-to-end

authentication statement for the same recovered root whose first message appears in the confidentiality game.

Target event.

Fix the live-core, single-root next-rekey profile and an admitted marked controller of Theorem 5.16 and Proposition 5.18. In particular, for each fixed candidate its target contribution and marked traffic entry are designated at the allocation points required by those interfaces; the original terminal selector corresponds to those designated entries whenever it returns that eligible candidate. As in Definition 4.7 and its adaptive version, the experiment retains one application challenge. Its original accepted logical send fixes the target root schedule through the send's stored provenance handle. The terminal selector may choose coordinates, reservation handles, and the cut adaptively; eligibility requires those coordinates to identify that same accepted send. Let E be its original eligibility event, A its source-attribution event, and $I \in J = [G] \times [C]$ the attributed pair of original generation and encapsulation reservations on A.

The target root and its event history are defined by the actual accepted challenge send, including on executions where A fails. A challenge without an accepted send has $E = 0$ and we set the target event flag to zero. On A, the root handle is the schedule entry generated from the effective ContextBound tuple m^* of the attributed pair. The handle is carried through replay, recovery, exact tuple aliases, and the resulting role keys and marked first-message key. Ledger aliasing follows Proposition 5.18.

For this target, let $\mathsf{Bad}_{\mathrm{auth}}$ be the union of the following events: a successful native response, final, or receipt confirmation whose complete core is absent from the earlier honest private-tag ledger under the corresponding target role key; or a fresh authenticated first-message receive and inbox installation whose $(\mathrm{AAD}, \mathrm{ciphertext})$ pair is absent from the earlier honest private-encryption ledger under the marked target first-message key. These are the events of the uniform-root authentication proposition restricted to this selected root schedule. The honest ledgers are populated at the first private computation, before persistence or public release. A repeated delivery or a previously consumed inbox result retains its original ledger provenance. Genuine aliases of the selected schedule entry share its honest ledger, while independently allocated entries retain their own provenance.

The experiment records these event flags as execution proceeds. Its eventual evaluation reads the accepted challenge send's root history, which may include confirmation events before that send. For a fixed pair $j = (g, a)$, the original a determines m^* before its first private hash computation at any protected alias, including a computation followed by a failed write, as specified in Section 5.7. Other protected tuples retain their own entries. If the original a never supplies a usable target, the candidate event is zero. The root entry, traffic entry, and direction retain their prescribed online marks. These rules connect the recorded history to the admitted controller throughout the comparisons. The probability bounded below is the joint probability $\Pr[E \wedge \mathsf{Bad}_{\mathrm{auth}}]$ in the original native game. The attribution residual bounds its part on $E \wedge \neg A$.

Authentication-specific comparisons.

For each fixed $j \in J$, use the native stopped continuation $G_{\mathrm{nat}}^{[j]}$, its real-component record idealization $V^{[j]}$, and the uniform-component continuation $\widetilde{V}^{[j]}$ from the candidate-local accounting. Apply to each of these worlds the same terminal Boolean code

$$W_j = \mathbf{1}[E \wedge A \wedge I = j \wedge \mathsf{Bad}_{\mathrm{auth}}], \qquad W_j = 0 \text{ after a declared local stop.}$$

All eligibility, attribution, tuple selection, ledger membership, and event flags are evaluated on the execution of that world. They are computable from the classical execution and provenance ledgers available to the existing reduction interface; no test of an unknown challenge bit is used.

Let a_j^{auth} be the real-versus-uniform advantage of the operational ML-KEM adapter of Section 5.7 when its output bit is W_j. Its public key is installed at the original generation's first actual private invocation, its encapsulation challenge at the original encapsulation's first actual private invocation, and every post-challenge decapsulation of the selected PQ ciphertext under the selected generation owner returns the challenge secret, regardless of the traditional ciphertext or context. The native traditional-component validation, its rejections, the original prechallenge history, and exact-retry caches are retained. The selected-pair correctness allowance is δ_j. Consequently this adapter gives

$$\left| \Pr_{V^{[j]}}[W_j = 1] - \Pr_{\widetilde{V}^{[j]}}[W_j = 1] \right| \le a_j^{\mathrm{auth}} + \delta_j. \tag{29}$$

This uses the same two-stage component experiment as the existing adapter, with the authentication event as its output bit.

Use the whole-continuation budgets of the ContextBound, root-initialization, joint-schedule, and confirmation comparisons. Define

$$\begin{aligned} \Omega_j^{\mathrm{auth}} := \min\{1, \Phi_{256}(d_j, Q_{H,j}) + \Delta_{\mathrm{init},j} + \Delta_{\mathrm{sched},j} + \Delta_{\mathrm{conf},j} + V_j 2^{-256} \\ + \chi_H(Q_{H,j}^{\mathrm{conf}}) + \varepsilon_{\mathrm{INT},j}\}. \end{aligned} \tag{30}$$

Here d_j counts distinct effective protected tuples in the ContextBound comparison; $\Delta_{\mathrm{init},j}$ is the event-bit cost of replacing the selected contribution's single native root-initialization entry by a fresh root; $\Delta_{\mathrm{sched},j}$ is the shared-compression joint-schedule cost; $\Delta_{\mathrm{conf},j}$ replaces the three confirmation HMAC rows; V_j counts the native confirmation verification attempts; χ_H is the framed-input collision bound; and $\varepsilon_{\mathrm{INT},j}$ is the one-key, at-most-one-encryption first-message AEAD integrity advantage with the native decryption interface. Each query budget includes preparation, ordinary primitive computations, honest protocol work, later continuation, and terminal event evaluation of its derived experiment. Budgets are fixed before interaction and are not reset at a challenge or at a hybrid boundary.

Theorem 5.25 (Native selected-root authentication) *In the live-core single-root next-rekey profile, with the admitted marked controller just specified and the origin-preservation interfaces*

of Section 5.4, Lemma 5.15, and Theorem 5.10,

$$\Pr_{G_{\mathrm{nat}}}[E \wedge \mathsf{Bad}_{\mathrm{auth}}] \le \min\left\{1,\ \eta^{\mathrm{nat}}_{\mathrm{attr}} + \sum_{j\in J} \min\{1, \varepsilon_{\mathrm{rec},j} + a^{\mathrm{auth}}_j + \delta_j + \Omega^{\mathrm{auth}}_j\}\right\}. \tag{31}$$

The record cost $\varepsilon_{\mathrm{rec},j}$ is the live-core record event cost (20)*, including its original fresh-seal equality allowance. The attribution residual is measured once in the original native game. For the executed-slot selector, it may be instantiated by Proposition 5.20 and the established reservation-signature bridge. An explicitly modeled projection extension adds its original selected-candidate residual $\eta^{\mathrm{nat}}_{\mathrm{proj}}$ and uses the corresponding event-bit record/report comparison in place of $\varepsilon_{\mathrm{rec},j}$.*

Proof First localize the native event. Write $F = E \wedge \mathsf{Bad}_{\mathrm{auth}}$. On attributed native executions the events $I = j$ partition $F \wedge A$. For a general declared projection extension, split off $F \wedge A \wedge B_I$. Its probability is at most $\eta^{\mathrm{nat}}_{\mathrm{proj}} = \Pr[E \wedge A \wedge B_I]$; likewise $\Pr[F \wedge \neg A] \le \eta^{\mathrm{nat}}_{\mathrm{attr}}$. For each fixed j, couple the original execution and $G^{[j]}_{\mathrm{nat}}$ until the first free or charged local stop. The free-stop certificate permanently falsifies $E \wedge A \wedge I = j$ on every continuation. A charged stop is absent on the summand restricted to $\neg B_j$. Therefore the same identical-prefix argument as Lemma 5.7, now applied to a Boolean variable, gives the exact identity

$$\Pr_{G_{\mathrm{nat}}}[F \wedge A \wedge I = j \wedge \neg B_j] = \Pr_{G^{[j]}_{\mathrm{nat}}}[W_j = 1].$$

This identity remains valid when a recorded authentication event occurs before the challenge send or terminal selection: the stopped predicate is the conjunction with final eligibility and attribution, and its free-stop certificate concerns every possible continuation. Consequently

$$\Pr_{G_{\mathrm{nat}}}[F] \le \eta^{\mathrm{nat}}_{\mathrm{attr}} + \eta^{\mathrm{nat}}_{\mathrm{proj}} + \sum_j \Pr_{G^{[j]}_{\mathrm{nat}}}[W_j = 1]. \tag{32}$$

The live-core interface has $B_j = \varnothing$ and $\eta^{\mathrm{nat}}_{\mathrm{proj}} = 0$.

Next compare a fixed continuation. The native representation invariant and immutable-history simulator preserve the selected event's classical execution, role-key provenance, private-tag and private-encryption ledgers, and original allocation indices. The real-to-opaque record comparison therefore applies with output W_j. Because it is a bit, its probability changes by at most $\varepsilon_{\mathrm{rec},j}$. This gives

$$\Pr_{G^{[j]}_{\mathrm{nat}}}[W_j = 1] \le \varepsilon_{\mathrm{rec},j} + \Pr_{V^{[j]}}[W_j = 1] \le \varepsilon_{\mathrm{rec},j} + a^{\mathrm{auth}}_j + \delta_j + \Pr_{\widetilde{V}^{[j]}}[W_j = 1],$$

where the second inequality is the operational component comparison (29). All role-key and honest-origin ledgers remain private reduction bookkeeping. Membership is checked by recording the actual complete core or encryption pair at its native invocation; it requires no inspection of quantum queries or unknown component tape.

In $\widetilde{V}^{[j]}$, the protected PQ component is an independent uniform 256-bit value. Apply Lemma 5.14 to the same Boolean continuation, with the complete effective tuple cache. It replaces each protected ContextBound answer by a private random-function entry at cost $\Phi_{256}(d_j, Q_{H,j})$. Exact aliases, including aliases occurring before the eventual selected send, share their entry. The target entry is marked at its first actual private use as prescribed by the existing controller interface, so its fresh value is sampled without conditioning on a future selection. If this target is never allocated or never becomes eligible, $W_j = 0$.

Lemma 5.15 places every use of this target contribution in one old-root and complete-info initialization entry. Its event-bit comparison (47) costs $\Delta_{\mathrm{init},j}$ and leaves the uniform selected root controller of Proposition 5.18. All other tuple entries, numerical coincidences, asynchronous confirmations, failed writes, recoveries, and permitted successor computations are simulated as in that controller. The event $W_j = 1$ implies an instance of its selected-root confirmation or marked first-message integrity event. The existing uniform-root authentication bound hence gives

$$\Pr_{\widetilde{V}^{[j]}}[W_j = 1] \le \Omega_j^{\mathrm{auth}}.$$

In particular the shared-compression schedule, the confirmation rows, fresh-tag guesses, framed-input collisions, and first-message integrity contribute exactly the event costs displayed in (30). Additional eligibility and source requirements can only restrict this event in the final world.

Each fixed candidate is thus bounded by $\min\{1, \varepsilon_{\mathrm{rec},j} + a_j^{\mathrm{auth}} + \delta_j + \Omega_j^{\mathrm{auth}}\}$. Substitute these bounds in (32) and cap the result by one. The partition was used only in the original native game. Ideal continuations retain their own selectors and event flags throughout, so no partition of different hybrid worlds is required. □

What the composition establishes.

The theorem connects an eligible native recovery execution to origin authentication of its selected confirmation schedule and marked first message. The comparison uses a single probability event at every stage; record and primitive reductions therefore enter with coefficient one. The operational component advantage a_j^{auth} is attached to this event's derived distinguisher. Its value is independently controlled by the same component security assumption as any output bit admitted by that experiment. The target scope, final eligibility, and whole-continuation budgets above specify exactly the executions to which the bound applies.

5.13 Native recovery with an online selected message

The finite-chain comparison can now be connected to the original native recovery experiment. The connection preserves its acquisition-based eligibility and executed source slots while allowing earlier message-key disclosures to influence which later application send is challenged. The root, traffic material and message step have separate allocation boundaries; keeping those boundaries explicit makes the selection implementable by the same native execution ledger.

The online native controller.

Retain the established-session live core of Section 5.8, with one authoritative history per endpoint and passive copies, and fix $1 \le L \le 2^{64} - 1$ before interaction. For each fixed candidate $j = (g, a)$, designate its original effective tuple m^* immediately before its first private ContextBound hash computation at any protected alias, as in Section 5.7. Designate the complete traffic entry I^* and direction before that entry's first block allocation. These designations are retained through later acceptance, recovery and equal-input aliases.

The controller maintains a private bookkeeping table indexed by $(\mathcal{M}, d, i)$, where $\mathcal{M}$ is the complete original traffic material handle, d its direction and i its chain index. Equal-info aliases of the same root schedule share $\mathcal{M}$; an independently allocated

material has its own handle even when its epoch, index or numeric key bytes coincide. Immediately before a previously unseen selected-direction step, the controller chooses either Continue or Mark, using its available history and the public descriptor of that step. Mark occurs at most once and only for $i < L$. A receive-side advance, including a skipped step followed by failed authentication, is a first computation for this purpose. Repeating a computation after a failed write retrieves the existing ledger entry and offers no new selection point. This ledger is reduction bookkeeping and changes no native state transition or response.

For each step passed with Continue before Mark, the observer may obtain its message key; the prefix seed and Extract value retain private handles. At a marked step ℓ, the observer may obtain the paired successor seed $s_{\ell+1}$ after its derivation. This named projection is bound to $(\mathcal{M}, d, \ell)$, irrespective of subsequent progress at either endpoint. Later chain steps are real computations from that seed. The selected current root, private prefix seeds, Extract values and marked message key have only the uses in the preceding interfaces. The already admitted successor-root, opposite-direction and other traffic-entry observations remain available.

These typed projections are the complete additional observation interface. Acquiring a complete session version still has its original scope and save time. In particular, a version containing an available successor seed also contains an acquired session version and remains subject to the original acquisition predicate. Public commitments and failed captures retain their original behavior.

There is one left-or-right application operation, tied to the marked step's original send. Its messages have equal length and common associated data. A native rejection, uncertain outcome or exact retry does not move that operation to another step. The terminal selector chooses the cut and original reservation handles as before and must identify this same accepted logical send. If no mark or accepted send is reached, the original valid-target condition is false. Application payloads and payload-bearing delivery remain private; the adversary has the existing plaintext-free Resume and public protocol interfaces. Thus E is exactly the original eligibility predicate evaluated on this controller's execution.

A mark may first occur during receive-side advancement, before the target send is reached. Accordingly the one-key AEAD experiments below allow classical decryption queries before their at-most-one encryption call. Such earlier queries retain their actual replies. After the left-or-right call, an exact challenge decryption returns the existing authenticated marker. This operational convention preserves the native receive-before-send order and is included in the stated AEAD advantages.

Lemma 5.26 (Native online-chain correspondence) *For this controller and a fixed candidate j, the stopped native live core realizes the prefix interface of Section 5.6 after the existing record, component, ContextBound and root-entry comparisons. Its original send and receive computations use one cached step per $(\mathcal{M}, d, i)$. The marked step has at most one original encryption tuple, and the live-core charged projection event remains empty.*

Proof At traffic installation, the exact root schedule and complete traffic information string determine the material handle. The role-dependent assignment in the native traffic constructor

makes one endpoint's send seed the matching endpoint's receive seed; the opposite direction has its other handle. Copying, reopening and exact pending-target reconciliation preserve this association. The source-attribution induction fixes the KEM origins of each installed schedule independently of which message index is later marked.

Induct on actual private chain calls across both endpoints. For index zero, the constructor supplies the corresponding traffic block. A later call uses the predecessor's successor with the fixed native domain and big-endian index. A repeated call under the same material, direction and index therefore repeats the same Extract/Expand inputs. Record its result at first computation, including in transaction-private receive state. If receive authentication or a later durable update fails, the native state may return to an earlier counter, while this proof ledger retains the original result. A retried advance obtains that result again. Taking an already retained skipped key likewise introduces no new computation or marking opportunity. This is an exact representation of the deterministic native computation in each comparison world.

The sender stores its pending message plan before reaching the chain step and encryption. A first-retention failure exits before encryption; once encryption is reached, the authoritative plan binds the message identifier, plaintext and associated data. A failed later write or process loss recovers that same plan or its committed outbox. Exact resumption uses its original encryption tuple, and successful commitment advances the send index. Hence the first encryption answer can be cached before its durable outcome, with at most one original tuple under the marked material and index. This also covers a mark whose first step computation occurred at the receiver.

The representation invariant for images and intents carries the same material and step handles through their immutable versions. Acquiring an intact version of an occurred candidate reservation's session, saved at or after that reservation, certifies the existing free stop. Lemma 5.6 applies before a final cut is chosen: every eligible cut precedes both reservations and therefore that acquired version's save. This certificate is independent of message index or marking time. Complete old or unrelated versions retain their ordinary correspondence. The new individual projections return exactly earlier message blocks and the marked pair's successor, which the prefix reduction already supplies. They expose no additional protected record or report key. Thus the live core retains $B_j = \varnothing$ and its established real-to-opaque record comparison.

After the ContextBound and root-entry comparisons, the complete native continuation therefore has the three online designations and the private prefix uses of Theorem 5.12. Real control verification, ACKs, other materials and later epochs remain in that continuation. The whole-execution budgets include their computations and all newly permitted prefix and suffix observations. □

Original-game costs.

Use the original native E, A, I and its residual $\eta^{\mathrm{nat}}_{\mathrm{attr}} = \Pr[E \wedge \neg A]$. Let $\varepsilon_{\mathrm{rec},j}$ be the record event cost in (20), now with the complete online controller's budgets. The operational component adapters retain their definitions from Section 5.7, with their output events evaluated in this expanded experiment. For candidate j, set

$$B_j^{\mathrm{chain}} = \min\bigl\{1,\ 2\Phi_{256}(d_j, Q_{H,j}) + 2\Delta_{\mathrm{init},j} + 2\Delta_{\mathrm{prefix},j}(L) + 2\varepsilon_{\mathrm{AEAD},j}\bigr\}. \tag{33}$$

Every budget is fixed before interaction and covers preparation, failed attempts, all endpoint aliases, the adaptive mark, real suffix computation and the terminal test

in its corresponding comparison. The prefix sum in (16) gives the explicit index-wise comparison cost. No additional combinatorial message-index guessing factor is introduced; all other whole-execution budgets may also grow with L.

Let $\mathsf{Bad}_{\mathrm{msg}}^{\mathrm{chain}}$ record a fresh authenticated receive and inbox installation under the marked step whose complete $(\mathrm{AAD}, \mathrm{ciphertext})$ pair is absent from the earlier honest encryption ledger for that material, direction and index. The flag is recorded at that installation, including before the target send, and later encryption does not erase it. Accepted-send provenance identifies the event at the end of the original experiment. For the component adapter with Boolean output $W_j = \mathbf{1}[E \wedge A \wedge I = j \wedge \mathsf{Bad}_{\mathrm{msg}}^{\mathrm{chain}}]$, write its operational advantage as a_j^{msg} and define

$$\Omega_j^{\mathrm{chain}} = \min\bigl\{1,\ \Phi_{256}(d_j, Q_{H,j}) + \Delta_{\mathrm{init},j} + \Delta_{\mathrm{prefix},j}(L) + \varepsilon_{\mathrm{INT},j}\bigr\}. \tag{34}$$

The costs for the two terminal tests use their own complete-query bounds; sharing notation does not identify their distinguishers.

Theorem 5.27 (Native online selected-message security) *In the online native live-core experiment above,*

$$\begin{aligned}\mathrm{Adv}_{\mathrm{msg}}^{\mathrm{chain}} \le \min\Biggl\{\frac{1}{2}, \frac{1}{2}\Bigl[\eta_{\mathrm{attr}}^{\mathrm{nat}} \\ + \sum_{j\in J}\bigl(2\varepsilon_{\mathrm{rec},j} + a_j^{+} + a_j^{-} + 2\delta_j + B_j^{\mathrm{chain}}\bigr)\Bigr]\Biggr\},\end{aligned} \tag{35}$$

and

$$\begin{aligned}\Pr_{G_{\mathrm{nat}}}[E \wedge \mathsf{Bad}_{\mathrm{msg}}^{\mathrm{chain}}] \le \min\Biggl\{1, \eta_{\mathrm{attr}}^{\mathrm{nat}} \\ + \sum_{j\in J}\min\{1, \varepsilon_{\mathrm{rec},j} + a_j^{\mathrm{msg}} + \delta_j + \Omega_j^{\mathrm{chain}}\}\Biggr\}.\end{aligned} \tag{36}$$

For the executed-slot selector, the corresponding original-game residual in each bound is at most Δ_{attr} *from* (25)*, using the complete expanded signing experiment and its established reservation-signature bridge.*

Proof For confidentiality, put $S = 2\mathbf{1}[g = b] - 1$ and run the same available candidate terminal code $\Psi_j = \mathbf{1}[E \wedge A \wedge I = j]S$ in every world. The identity in Lemma 5.7 localizes the original score before any candidate-specific comparison. The native free stops remain valid by Lemma 5.26. Its record correspondence gives $\rho_j \le 2\varepsilon_{\mathrm{rec},j}$; the two operational component-event comparisons cost $a_j^{+} + a_j^{-} + 2\delta_j$.

In the random-component world, apply Lemma 5.14 with the unchanged complete tuple cache. The target contribution is selected at its own first private allocation, before the later traffic and message designations. Its single root-use entry is still that of Lemma 5.15, since later chain work does not reuse the combined contribution or alter its original pending-root inputs. The two score comparisons therefore cost $2\Phi_{256}(d_j, Q_{H,j}) + 2\Delta_{\mathrm{init},j}$. The resulting uniform-root continuation is exactly the online prefix interface by the new correspondence lemma. Corollary 5.13 bounds its qualified signed score by $2\Delta_{\mathrm{prefix},j}(L) + 2\varepsilon_{\mathrm{AEAD},j}$. Every

selector, disclosure, eligibility check and terminal event is recomputed within that world. In particular, qualifying a run uses the available provenance and view rather than direct access to b. Thus the random-component score is bounded by B_j^{chain}. Substitution in Proposition 5.8 proves (35).

For integrity, apply the Boolean localization identity from (32) to the selected-message event W_j. The same free-stop certificate makes a stopped candidate's final conjunction false, even if its recorded inbox event preceded the challenge send. The record comparison costs $\varepsilon_{\mathrm{rec},j}$, and the operational component comparison costs $a_j^{\mathrm{msg}} + \delta_j$. Apply the complete-tuple and root-entry comparisons to this Boolean continuation. The remaining event is a subset of the selected-key integrity event in Corollary 5.13, so its probability is bounded by $\Omega_j^{\mathrm{chain}}$. A fresh successful decryption before the honest encryption is already a valid integrity witness; the reduction may return it immediately. This justifies the same bound when Mark was first reached by the receiver. Summing the native candidate partition and capping by one proves (36).

Finally, the executed-slot association proof is an induction over traffic installation and its originating KEM plans. Its signed-body and full-wire checks precede the message-index choice. Advancing, retaining or reopening the installed traffic entry preserves those origins, while the earlier-key projections only enlarge the complete adversarial continuation. Proposition 5.20 therefore gives the stated original-game attribution bound with its updated whole-execution resources. This term is charged once, and the root, record and component costs retain their displayed coefficients; no additional guess of the marked message index is used. □

For $L = 1$, the prefix cost reduces to the existing first-message schedule cost. The enlarged controller adds the choice of a later uncomputed index after observing earlier keys, while preserving the native source and acquisition semantics. An explicitly modeled record/report projection extension keeps its original $\eta_{\mathrm{proj}}^{\mathrm{nat}}$ and corresponding local comparison costs, as in the preceding native theorems.

5.14 Several root entries from one recovered contribution

One recovered contribution can feed several available old roots and complete root-information strings. The comparison below retains these uses jointly. Let S be uniform on 256 bits, initially independent of the public compression oracle h and auxiliary state, and retain F and Γ from Section 5.5. Each root request supplies a 32-byte salt r and a complete native 97-byte information string I. Its real value is

$$P_r = F(r, S), \qquad R_{r,I} = F(P_r, I \| \mathtt{01}).$$

The parent contribution is used through these Extract calls; each resulting PRK is used through its root-Expand interface. The controller retains the original source ledger and any established local stops.

Descriptors and available data.

The root-use descriptor is $e(r, I) = (\Gamma(r), I)$, with exact byte equality. Equal descriptors retrieve the same root entry across callers and resumptions. Indeed, (43) shows that equal Γ values give equal Extract values for every S, and equal complete information strings then give equal Expand values. Assign each newly encountered Γ class its next

lifetime index i, preserving this index after a restart. Numerical output coincidences between distinct classes keep their ordinary shared-oracle distribution.

The controller computes salts, information strings and its selection from its available state: executed preparation, independent coins, public-oracle responses, unmarked roots, and the declared outputs of the marked schedule. At most one fresh descriptor is marked before its first root evaluation. Its root is installed in one retained schedule handle; every descriptor alias retrieves that handle. All its uses belong to the same interface of Section 5.5, including its session salt, confirmation and successor calls, complete traffic-information table and budgets. Other root entries may be disclosed and used in later available computations, including as salts. The retained handle supplies exactly that schedule's operations. This is a rule about available computation; a locally known salt can numerically coincide with a protected value.

Fix finite bounds $D, C, N_1, \ldots, N_C$ before execution: at most D distinct salt strings, at most $C \le D$ classes, and at most N_i distinct complete information strings in the ith class. These bounds hold in every comparison experiment. All oracle-dependent preparation is executed and counted. Exact repeats are cached. Write q_S for the complete ordinary compression-query budget in the column reduction, and q_i for that in the ith row reduction. These budgets include the whole controller, all other real primitives, continuation and terminal test. An additional independent whole-hash oracle may remain available throughout.

Theorem 5.28 (Finite multiple-root entry) *Replace the real root interface by a private table $\mathcal{W}(\Gamma(r), I)$ of independent uniform 32-byte entries, preserving all descriptor aliases and the controller above. Every efficient available output-bit test changes by at most*

$$\Delta_{\mathrm{mr}} = \min\left\{1, \varepsilon_{\mathrm{Hjoint}}(C, 0, q_S, 32) + \sum_{i=1}^{C} \varepsilon_{\mathrm{Hrow}}(N_i, q_i, 98)\right\}. \tag{37}$$

In the table experiment, an online marked fresh descriptor supplies the independent uniform root required by Theorem 5.10, together with the complete interactive controller. The marked descriptor may belong to a previously used class.

Proof Compute $\Gamma(r)$ from available salt bytes and retain its class index. Use one column call for the first salt in each new class; reuse its reply for later equal classes. The descriptor identity makes this an exact rendering of the real Extract interface. Lemma C.1 replaces these calls by $P_i = V(\Gamma(r))$ at the first cost in (37). The remaining root computation consists of HMAC rows under the independent private keys P_i.

Replace these rows in first-class-allocation order. At step i, presample the hidden uniform row-challenge key and assign it to the ith new class when that class is first encountered. Before that event the key has no use. Classes already replaced use their private tables; later classes have locally sampled independent keys and are evaluated with the same h. The controller accesses the selected PRK through its row alone. Thus the real row branch matches the preceding hybrid and the ideal row branch matches the next one. Lemma C.2 bounds this step by $\varepsilon_{\mathrm{Hrow}}(N_i, q_i, 98)$. If the class is never allocated, the hybrids agree. Summing over the fixed class budget establishes the bound.

Distinct classes remain distinct reduction slots even when their sampled PRK bytes coincide. In a real row branch, local evaluation under another class's key uses the same h and

automatically preserves any resulting coincidence; these ordinary queries are included in q_i. Thus the argument uses the full independent-key distribution.

For each ideal row, $I \mapsto I\|\texttt{01}$ is injective. Its root interface is therefore exactly the table $\mathcal{W}(\Gamma(r), I)$. At a fresh descriptor, its value is independent of all previously used table entries and the oracle. Presample the designated uniform root, withhold it during the executed prefix, and assign it when the controller marks that fresh descriptor. All other entries retain their independent sampling and all aliases retrieve the retained handle. The selection uses the available prefix, so this deferred assignment has exactly the table distribution. It embeds the whole controller in the existing root experiment. The row hybrids already covered every class, so target selection adds no class guess. □

Explicit query accounting.

Let B bound all ordinary compression work outside this entry service, including its entire downstream schedule and test. Cache the two public Γ computations per distinct salt. With $N = \sum_{k=1}^{C} N_k$, valid conservative choices are

$$q_S = B + 2D + 5N, \qquad q_i = B + 2D + 5 \sum_{k=i+1}^{C} N_k. \tag{38}$$

A 98-byte HMAC evaluation uses five compression calls, including its two key pads. In the column comparison every root row is real; at row step i, the later rows are real and the earlier rows are tables. External class scheduling accounts for $2D$, including new salts that share an existing class. The column primitive's own public pad work remains included in its established bound. B is a uniform bound over these complete derived experiments.

Corollary 5.29 (First message after a marked root entry) *For the available controller above and its marked schedule, use the first-message interface and available qualification predicate of Corollary 5.11. If Z is its qualified signed message score, with zero on runs having no marked challenge, then*

$$|\mathbb{E}Z| \le \min\{1,\, 2\Delta_{\mathrm{mr}} + 2\Delta_{\mathrm{sched}} + 2\varepsilon_{\mathrm{AEAD}}\}. \tag{39}$$

The schedule and AEAD reductions include the complete entry controller and its shared-oracle work. If the contribution itself comes from the protected ContextBound table of Lemma 5.14, its additional signed-score cost is $2\Phi_{256}(d, Q_H)$. For a candidate whose stopped execution realizes this entry controller, it replaces the single-entry $2\Delta_{\mathrm{init}}$ term by $2\Delta_{\mathrm{mr}}$ in the corresponding continuation bound; the original accounting terms keep their established definitions.

Proof Fix each application challenge bit for the entry comparison and average afterwards, as in the retained continuation. The randomized output test with success probability $(Z+1)/2$ converts Theorem 5.28's bit gap into a score gap of at most $2\Delta_{\mathrm{mr}}$. Its ideal marked root initializes the entire controller of Corollary 5.11, which supplies the remaining two terms. Apply the same conversion to the ContextBound comparison when present. □

5.15 An owned SDK purpose-key table

The native SDK already derives several application keys from one combined secret. Its `SharedSecret::derive_key` operation takes a purpose, protocol label and context, and returns a new `DerivedKey` owner. The same source owner supports repeated calls. The pinned source exercises all five purposes at both peers and checks their corresponding outputs and lifecycle behavior. This supplies a concrete application of the finite entry-table argument with the SDK's own information encoding.

Lemma 5.30 (Information-length form of the entry comparison) *For the finite service $R_{r,I} = F(F(r,S), I\|\mathtt{01})$, let the complete available information strings in class i have length at most A_i bytes. Put $L_i = A_i + 1$ and $\lambda_i = \ell(L_i)$, within the retained SHA-256 padding domain. Use the descriptor classes and finite budgets of Theorem 5.28. With the parent contribution used through the Extract interface and each PRK through its Expand row, the whole interactive service is within output-bit gap*

$$\Delta_{\mathrm{ent}} = \min\left\{1, \varepsilon_{\mathrm{Hjoint}}(C, 0, q_S, 32) + \sum_{i=1}^{C} \varepsilon_{\mathrm{Hrow}}(N_i, q_i, L_i)\right\} \tag{40}$$

of the consistent independent table indexed by $(\Gamma(r), I)$. The continuation may process each returned value through its declared interface. Conservative complete ordinary-query budgets are

$$q_S = B + 2D + \sum_i (2+\lambda_i)N_i, \qquad q_i = B + 2D + \sum_{k>i} (2+\lambda_k)N_k.$$

Proof Use the column replacement and first-class-allocation row hybrids of Theorem 5.28. At row i, every queried message has length at most L_i. The map $I \mapsto I\|\mathtt{01}$ is injective also for variable-length strings. Hence the final private row gives exactly the full-information table. Each actual Expand costs at most $2 + \ell(L_i)$ compression calls, yielding the displayed budgets. The reductions receive the root-row replies and run the entire declared continuation from them. This proves the table comparison for that continuation. The native 97-byte case has $L_i = 98$ and $\lambda_i = 3$. □

The actual SDK descriptor.

Let $\sigma = \mathtt{QPeriapt\text{-}SDK\text{-}HKDF\text{-}SHA256\text{-}v1}$, a 27-byte string, and $r_0 = \sigma\|\mathtt{00}^5$, its 32-byte zero extension. HMAC pads both σ and r_0 to the same 64-byte key block [22], so $F(\sigma, S) = F(r_0, S)$ for every S and every h. The complete SDK information string is

$$\begin{aligned} I = {} & D_{\mathrm{sdk}}\|\,\mathrm{BE32}(1)\|\,\mathrm{BE32}(2)\|\mathit{policyState}_{36}\|\mathit{rootDigest}_{32} \\ & \|\,\mathrm{BE32}(p)\|\,\mathrm{BE16}(|a|)\|a\|\,\mathrm{BE32}(|c|)\|c\|\,\mathrm{BE16}(32), \end{aligned} \tag{41}$$

where $D_{\mathrm{sdk}} = \mathtt{QPeriapt\text{-}SDK\text{-}Key\text{-}v1}$ has 19 bytes. The fixed suite code 1 denotes ML-KEM-768+X25519 and profile code 2 denotes ContextBound. Here p is one of the five purpose codes, a is the protocol label, and c is the context. The policy state is its 4-byte version followed by its exact-document 32-byte digest; the root digest is

SHA-256 of the runtime's pinned trust root. Labels contain 1–255 printable non-space ASCII bytes and contexts have at most 65536 bytes. Thus

$$|I| = 107 + |a| + |c| \in [108, 65898], \qquad |I\|\texttt{01}| \le 65899.$$

Length parsing recovers every effective field. The descriptor is $(\Gamma(r_0), I)$ across all owners of the same protected contribution. Runtime identities and owner identifiers retain their separate lifecycle roles. The effective digest bytes determine their respective info fields.

Owner correspondence.

Fix a protected contribution family with common value S and preserve the actual metadata of each source and derived owner. Represent a derived owner by $(\mathit{runtime}, \mathit{live}, I)$ and a reference to the common key entry K_I. Every accepted derivation creates a new owner; equal I values reference the same entry. Closing one derived owner changes its own live flag. Closing a source owner changes its future derivation access; existing derived owners retain their stored key. Runtime closure makes later open checks fail. An admitted derivation or an export whose open check already succeeded follows its actual completion path. Previously exported copies remain available to their recipients.

These transitions follow directly from the source's separate `Option` fields, shared runtime state and admission checks. The inspected FFI derivation path additionally retains reservation, publication and parent-revocation metadata. A first private KDF computation allocates the descriptor entry even if its later handle publication fails; a retry retains that key entry while obtaining its actual new owner. The Rust in-flight bound and the FFI live-handle bound are preserved as metadata. The proof's lifetime bound N counts distinct complete infos.

Proposition 5.31 (Native SDK key-table correspondence) *For one initially independent uniform protected contribution, let the family's source owners use the Rust derivation interface above, together with derived-owner export, close and runtime lifecycle operations. The parent contribution and internal Extract value remain within their stated KDF interfaces. Let N bound distinct infos before execution and L bound their Expand-message length. Then the entire owned-API experiment differs from one using independent uniform keys K_I, with all aliases and owner states preserved, by output-bit gap at most*

$$\begin{aligned}\Delta_{\mathrm{sdk}} = \min\{1, \varepsilon_{\mathrm{Hjoint}}(1, 0, B + 2 + (2 + \ell(L))N, 32) \\ + \varepsilon_{\mathrm{Hrow}}(N, B + 2, L)\}.\end{aligned} \tag{42}$$

The budgets cover complete preparation, other real computations and the terminal test. At full SDK input limits, $L = 65899$ and $\ell(L) = 1031$.

Proof Normalize the fixed salt as above. The first-block HKDF definition [21] gives $P = F(r_0, S)$ followed by $F(P, I\|\texttt{01})$; the source keeps its PRK internal. An induction over admitted calls preserves each owner state and every public status: validation and admission use available metadata; derivation creates a fresh owner referring to its descriptor's deterministic value; exports apply that owner's current guards; close and publication operate on their existing owner

identities. Repeated key values can therefore be memoized while all native owner transitions continue to execute. Apply Lemma 5.30 with $D = C = 1$. The class/row comparison changes only these internal key values. All field-dependent choices, exports and lifecycle responses are recomputed in their own experiment by the same controller. □

For the existing five-purpose source workload, the label has 13 bytes and its context 24, giving $|I| = 144$, $L = 145$ and $\ell(L) = 4$. Its two peer owners use five complete descriptors when their effective policy bindings agree. Thus $N = 5$ gives column ordinary budget $B+32$ and row budget $B+2$. The table comparison permits disclosure of all derived entries and establishes their joint pseudorandom separation. In particular, the probability of a key collision among exports from distinct allocated infos, with at most N total infos, is at most $\min\{1, \binom{N}{2}2^{-256} + \Delta_{\mathrm{sdk}}\}$.

A selected key for one application encryption.

An application may mark a new complete descriptor before its first private KDF evaluation, retain its key privately, and use it in the one-key, one-encryption AEAD experiment defined in Section 5.5, with that experiment's nonce, equal-length messages, associated data and decryption restriction. Every source owner and derived alias of that descriptor routes to the same key-use interface, with its own native liveness checks. Other derived entries may be exported. Let Z be the qualified signed score, with qualification computed from the available classical view and ledger and value zero when the marked challenge is absent. Fresh table allocation then supplies the uniform AEAD key, giving

$$|\mathbb{E}Z| \le \min\{1,\, 2\Delta_{\mathrm{sdk}} + 2\varepsilon_{\mathrm{AEAD}}\}.$$

Fix each application challenge bit for the KDF comparison and average the two score bounds. In the final AEAD game, use the adversary's guess on qualified runs and an independent fair guess otherwise. A preceding complete-tuple ContextBound comparison contributes its established $2\Phi_{256}(d, Q_H)$ term when the stopped execution realizes that source family and available controller.

This consequence concerns the exported purpose key's declared application use. The durable rekey driver retains its own 97-byte pending-root information and established source invariant; the SDK table provides a separate concrete application of the finite entry proof.

6 Mechanized recovery and separating executions

6.1 The multi-epoch symbolic theorem

The checked symbolic model includes inherited roots, alternating proposers, asymmetric confirmation, initial traffic, and unbounded epochs. Signing and wrapping keys may be shared across their modeled uses. Its Cut event exposes the selected wrapping keys; exposure of other epochs' entropy is represented separately. We state the main theorem in terms of its recorded events, preserving its original target and disclosure predicates.

Let τ be a trace of that model. Write $\mathsf{SessionKeys}(s, a, b)$ for the actual identity keys associated with session s, $\mathsf{KeyReserved}(s, e, P)$ and $\mathsf{EncapReserved}(s, e, Q)$ for the target reservations, and $\mathsf{Message}(s, e, r, k, m)$ for an application-message event.

Define $\mathsf{AcquiredAfterSave}_c(s)$ to hold when there exist a role, an epoch, a revision, and matching Saved/Snapshot events for session s, with the *save* event strictly after c. The matching record's session projection is the relevant saved object.

Theorem 6.1 (Checked multi-epoch symbolic secrecy) *Suppose τ contains the corresponding SessionKeys, Cut at c, target reservation events at g and h, and $\mathsf{Message}(s, e, r, k, m)$. If*

$$c < g, \qquad c < h,$$

neither actual target identity key is disclosed anywhere in τ, no reserved entropy of target epoch e is disclosed, and $\mathsf{AcquiredAfterSave}_c(s)$ is false, then τ contains no adversarial knowledge event $\mathsf{K}(m)$.

The theorem was checked with all 24 reusable dependencies in one Tamarin process. An independent replay produced a byte-identical proof certificate. Appendix A identifies the exact query, source, toolchain, replay receipt and complete auxiliary-query inventory.

The theorem uses the modeled cryptographic constructors and record interface. Its key scope follows actual identity ownership: a shared signer exposed through another session is still a disclosure of that target key. Likewise, a retained target projection saved after Cut remains subject to the acquisition condition when a sibling write carries it into a later aggregate image.

6.2 Executability and the freshness boundary

A finite two-round construction exercises the operational rules with an exposed ancestor root and satisfies the complete original positive existential formula. A separate source-derived evaluator checks its rule instances and global quantifiers. This provides a concrete execution for the modeled recovery premises.

The delayed-receipt construction instead contains a pre-Cut epoch-two generation reservation and a post-Cut encapsulation. It derives the target payload while satisfying the remaining main-theorem antecedents. Its point audit identifies $c < g$ as the sole false antecedent. Seven countertraces are rejected by the witness checker. Together, the positive and separating traces connect the recovery condition to executable schedules.

6.3 Native reservation disclosure

The native experiment stops a child process after committing a reservation, kills and reaps it, reopens the journal, and inspects the retained phase. It exercises two distinct boundaries:

- a 277-byte key-generation token, before key computation or offer publication;
- a 245-byte encapsulation token, after offer authentication and before encapsulation computation or response publication.

The predictor obtains the previous root, journal identity, wrapping key and captured token. Thereafter it uses the public policy configuration and later authenticated control wires. Its prediction interface has no journal or filesystem input, and the experiment withholds signer files and future private owners. A wrong wrapping key is rejected.

For each cut, both honest peers complete the rekey and reopen their journals. The driver then creates fresh 128-byte application payloads, sends three in each direction, and checks honest delivery. Independently reconstructed traffic owners authenticate and decrypt all six intercepted frames. The two cuts therefore yield twelve actual post-restart plaintext recoveries. The delayed-receipt extension repeats the six-message check at epoch two using one responder snapshot.

Proposition 6.2 (Execution-time and counter-based separation) *In the native executions, computing a KEM after restart, completing the target epoch, and subsequently reopening the journal do not make a previously disclosed reservation an independent secret contribution. Moreover,* $\min(c_A, c_B) + 1$ *is not a universal upper bound on exposed epochs.*

Proof At both process-cut boundaries the captured reservation reconstructs the later honest KEM result exactly. Equations (1)–(2) then reconstruct traffic keys from the captured predecessor root and public wires; authenticated decryption verifies the six corresponding plaintexts after completion and restart. In the delayed-receipt execution, $(c_A, c_B) = (0, 1)$ at capture and six epoch-two plaintexts are recovered. Thus the proposed counter bound equals one while an exposed epoch equals two. □

6.4 An aggregate write preserves the target origin

A separate two-session native execution changes a sibling session and writes a new aggregate image. It checks exact equality of the retained target record, reopens and retries the original operation, and decrypts six subsequent target messages in each build profile. The new global revision and outer encryption nonce coexist with the old target reservation. This execution instantiates the alias/provenance distinction in the record model.

The native executions isolate the reservation and acquisition conditions used in Theorem 6.1. Section 5 gives the record simulations and common-experiment accounting for the corresponding computational interfaces.

7 Validation of exposure semantics

The validation asks how the formal exposure conditions appear in executable state and communication histories. We use four complementary designs. First, the native executions in Section 6 vary the time of reservation use, acknowledgement delivery and aggregate rewriting while retaining the original secret source. Second, two implementations of record readability are checked against each other and against independently stated expected results. Third, a reference corpus records the message keys derivable from actual component snapshots under fixed communication schedules. Fourth, native and reference executions share a directed opportunity budget and report exposure jointly with admitted service and communication work. A paired native ablation within this design isolates application-retry readiness while retaining the protocol and the opportunity list. Table 5 relates each design to the question it answers.

Table 5 Research questions and their validation objects. Native plaintext recovery, finite-model equality and component key derivation are recorded as separate outcomes.

Question	Design and observation
Which event fixes a contribution's origin?	Reservation-boundary captures, delayed receipts and sibling aggregate writes; reconstruction of subsequent traffic from the captured original inputs.
Which acquired objects jointly expose a source?	Iterative decryption and guarded routes on the same finite graphs, with explicit expected cases for key dependencies and later disclosures.
How does retained state interact with communication opportunities?	Pinned component snapshots across seven schedules, with reconstructed keys indexed by direction, epoch and message number.
How does a recovery budget trade exposure for service?	Native and reference runs under common directed opportunities, with separate first commitments, logical deliveries, derived sets and full-frame costs.

7.1 Finite-record semantics

The specialization in Section 4.4 has two executable realizations: iterative decryption of available record objects, and evaluation of normalized paths after computing key reachability. Their common inputs are a finite immutable record graph, canonical key identities, fixed public key-derivation edges, acquired records, directly disclosed keys, and directly disclosed source values. The evaluator checks equality of their exposed-source sets for every subset assignment in each declared graph.

Table 6 gives the complete finite inventory. The 5,360 assignments include direct-source subsets; they are separate from the 420 reference endpoint-snapshot cases described below, including the isolated chunk-width variant. Cyclic key derivation requires an initially disclosed key to seed reachability. Cyclic record containment is rejected. The source-preserving copy and rewrapping relations are inputs to these graphs, while the native executions in Section 6 provide their separately identified implementation evidence.

The history driver additionally evaluates all 360 legal permutations of a six-event acquisition, disclosure and deletion schedule, checking equality and knowledge monotonicity at all 2,160 prefixes. Ten directed histories exercise late key disclosure, capture before local deletion, retained pending targets, sibling copies, independent inner guards, and direct entropy disclosure. All records in these finite histories are initially available; record-creation events and operation-context validation remain in the surrounding provenance interface.

An independent review driver adds 17 cases with explicitly specified expected outputs. These cover seeded and unseeded key cycles, a self-loop, a reversed key-derivation chain, shared-child paths, direct child capture, partial keys from distinct paths, and local deletion of an embedded record's standalone copy. Expected outputs are checked independently of the agreement between the two algorithms. All 17 cases pass.

Record and key configuration	Assignments
Journal-derived operation-wrapping key	512
Independent inner and outer wrapping keys	128
Rewrapped representations of one origin	512
Nested aggregate with derived keys	4,096
One actual key reused across wrappers	16
Empty record and unrelated source	64
Key-derivation cycle requiring a disclosed seed	32
Total	5,360

Table 6 Exhaustive acquisition, key-disclosure and direct-source assignments for the declared finite graphs. The general equality theorem is proved in Section 4.4; these checks validate its executable realizations on the stated graph inventory.

Five semantic controls exhibit the effect of discarding an adversary's copy on local deletion, fixing exposure at capture despite later key disclosure, resetting a source to the newest image identifier, dropping an inner-key guard, and treating a public digest as record acquisition. Each produces a specified disagreement with the declared semantics. They are controls on interpretations of the record model.

Appendix A identifies the source, result files and independent expected-result checks. The reproduction entry point rejects execution modes that disable its assertions.

7.2 Reference snapshots and communication schedules

The pinned SPQR source is `f2589fef855c10f39d72634dab3d14654dd410bf`. Seven schedules each contain 2,048 sends: alternating, loss, reordering, duplication, 9:1 asymmetry, offline-then-exchange, and one-way traffic. Six capture points for each endpoint produce 84 passive snapshot cases. An authenticated whole-KEM control runs those schedules at local-send intervals 1, 32 and 64, producing 252 corresponding cases. Together these are 28 configuration–schedule combinations and 336 snapshots. An isolated 64-byte polynomial-chunk SPQR variant supplies another seven schedules and 84 snapshots with its own source identity. The combined comparison therefore has 35 configuration–schedule combinations and 420 snapshots. The 32-byte and 64-byte labels describe the polynomial chunk width.

7.3 Exposure and communication observations

The snapshot predictor receives one endpoint's serialized state and the full public transcript, including packets lost before peer delivery. A separate harness compares every derived key with the actual sender's key. Public deterministic seeds are test fixtures; the measurement establishes causal key derivation through the declared predictor interface. The output label `not_derived` denotes that predictor's result. It is separate from the security predicate of a computational experiment.

Table 7 reports the alternating-schedule observations. The snapshot is taken at A immediately after its first send, leaving 2,047 future sends. The pending KEM key

Component/profile	Mean bytes/send	Peak frame	Epochs A/B	Derived keys
SPQR, 32-byte chunks	32.150	37	24/23	172
SPQR, 64-byte chunks	58.971	69	46/45	88
Whole KEM, interval 1	819.658	1,225	683/683	4
Whole KEM, interval 32	44.047	1,223	32/32	66
Whole KEM, interval 64	25.023	1,223	16/16	130

Table 7 Component-level alternating workload; SPQR byte labels denote polynomial chunk width. Byte counts are the component wire contribution; epoch counts are each component's root/output observations. The constructions have different authentication and retention graphs. Outer application AEAD, transport, durable storage and energy are outside these measurements.

Schedule	SPQR 32	SPQR 64	KEM 1	KEM 32	KEM 64
Reconstructed keys among 2,047 future sends					
Alternating	172	88	4	66	130
Loss	198	102	6	66	130
Reordering	219	121	37	93	163
Duplication	172	88	4	66	130
9:1	748	388	28	338	658
Offline then exchange	424	342	260	322	386
One-way	2047	2047	2047	2047	2047
Mean component bytes per send					
Alternating	32.150	58.971	819.658	44.047	25.023
Loss	31.957	58.336	819.564	44.047	23.520
Reordering	31.876	61.795	817.999	38.927	23.058
Duplication	32.150	58.971	819.658	44.047	25.023
9:1	23.508	40.531	819.381	21.031	13.417
Offline then exchange	32.957	60.376	869.973	44.125	25.118
One-way	38.875	70.875	1223.938	44.969	25.954

Table 8 The seven retained schedules at the same snapshot coordinate: endpoint A immediately after its first send. SPQR 32 and 64 identify polynomial-chunk widths in bytes; KEM 1, 32 and 64 identify whole-KEM local-send intervals. Each column uses its recorded construction, authentication graph and retained-state interface. Bytes cover the component wire; reconstructed keys are exact predictor matches.

accounts for 87 of the 172 derived keys in the pinned 32-byte SPQR case, demonstrating why the currently emitted chain epoch alone does not describe exposure.

Table 8 extends the comparison to every retained schedule at the same capture coordinate. Increasing the SPQR chunk width from 32 to 64 bytes changes the alternating trace from 172 to 88 reconstructed keys, while mean component traffic rises from 32.150 to 58.971 bytes per send. Under 9:1 traffic, those reconstructed sets contain 748 and 388 keys. The whole-KEM interval-32 trace similarly changes from 66 keys under alternating traffic to 338 under 9:1 traffic, while its mean bytes per send fall from 44.047 to 21.031. These paired observations show how communication direction

and control progress shape exposure alongside the amount of transmitted material. The retained construction and disclosure interface identify the comparison in each column.

The offline prefix leaves substantial exposure even for the whole-KEM interval-1 control: its predictor reconstructs 260 future keys at this cut, compared with four in the alternating trace. Exact duplicate delivery leaves the emitted-byte and reconstructed-key measurements unchanged in all five profiles. The accompanying deterministic analysis checks all 420 snapshot inventories and reproduces the displayed rows directly from four hash-bound corpus files. It preserves the original execution identities and the exact derived-key interpretation.

The one-way schedule produces no fresh PQ epoch in these references and the predictor derives every post-snapshot key. The whole-KEM active-control experiment addresses a different property: one initial endpoint-state capture allows two attacker forks to replace subsequent exchanges while honest confirmed epochs increase. All 1,536 sender keys across three 512-send profiles match the attacker-derived keys. This supplies a concrete distinction between fresh KEM execution and authenticated recovery when the ongoing authentication authority is itself derived from captured session state.

7.4 Matched directed opportunities and admitted service

The matched experiment executes native Continuity with per-direction send budgets $b \in \{32, 64\}$ and the five reference profiles under seven common schedules. Its 49 configuration–schedule pairs place exposure, admitted application work and communication cost on the same opportunity budget. The protocol source is fixed at `358b921a`; the reference SPQR revision and its isolated chunk-width variant retain the identities stated above.

Workload and capture.

Each schedule contains 2,048 directed transmission opportunities, indexed 0–2,047. The designated sender receives one 32-byte application request at each opportunity and may emit at most one frame of at most 8,192 bytes. Uncommitted requests remain in its FIFO queue. A response queued by receipt uses a later opportunity in the response direction. Protocol retries consume opportunities and retain their exact bytes; network duplication adds only a delivery. Terminal draining delivers already emitted frames.

Both capture experiments use the cut after opportunity 0's scheduled local and network actions and its checkpoint, and before opportunity 1 begins local computation. They separately acquire the complete current session state of A or B, together with every public emission, including packets lost before delivery. For Continuity this is the selected session record and its matching wrapping key; its independently owned identity signer remains separate. The references authenticate through their session MAC state. One application has been committed before the cut in every profile, so the post-capture denominator is the total first-commitment count minus one.

Native admission is a DeviceJournal-committed ciphertext; reference admission is a returned component state, with a separately executed 32-byte-payload AEAD audit. The native predictor derives candidate keys from the captured epoch-zero chain seeds and checks them against the actual ciphertext, with altered-seed rejection controls. Reference predictors recover component message keys, and the execution audits their

Profile	Committed / delivered	Capture A derived / admitted	Capture B derived / admitted	Emitted bytes
Continuity, $b = 32$	688 / 688	63 / 687	63 / 687	565,288
Continuity, $b = 64$	718 / 718	127 / 717	127 / 717	407,596
SPQR, chunk 32	2,048 / 2,048	172 / 2,047	85 / 2,047	172,339
SPQR, chunk 64	2,048 / 2,048	88 / 2,047	43 / 2,047	227,268
Whole KEM, interval 1	2,048 / 2,048	4 / 2,047	1 / 2,047	1,785,156
Whole KEM, interval 32	2,048 / 2,048	66 / 2,047	1 / 2,047	196,704
Whole KEM, interval 64	2,048 / 2,048	130 / 2,047	1 / 2,047	157,744

Table 9 Alternating transport with 2,048 directed opportunities and 2,048 offered 32-byte requests. Each exposure denominator counts first commitments after the capture cut; A/B identify the separately captured endpoint. Emitted bytes include actual native frames or the reference component with its executed payload envelope. Uncommitted requests equal 2,048 minus the committed count.

Budget	First app.	ACK	Control	App. retries	All exact retries	Image revisions	`sync_data` calls
32	688	676	80	604	636	3,554	14,219
64	718	711	40	579	596	3,644	14,579

Table 10 Native alternating communication and persistence. First application, ACK, control and application-retry columns partition the 2,048 emitted frames. All exact retries also count repeated control frames and overlap that partition. Image revisions and physical sync calls sum both endpoints.

envelopes with the actual keys and wrong-key controls. The displayed derived sets are the outputs of these specified predictors. The cryptographic recovery results use the eligibility and observation interfaces of Sections 4 and 5.

Current-epoch capacity and renewed exchange.

At this cut, the first application has consumed one of the two directions' epoch-zero allocation indices, and the captured native records contain no pending rekey entropy. Each direction admits at most b first allocations under that epoch. Thus at most $2b - 1$ further epoch-zero allocations remain; exact retransmissions reuse their original allocation. Equality requires both directions to use their remaining epoch-zero capacity. This elementary allocation count supplies the exposure envelope tested by the native trace predictor.

In the alternating schedule, reducing b from 64 to 32 changes the derived set from 127 of 717 post-capture commitments to 63 of 687 (Table 9). Both counts attain $2b - 1$. The same change delivers 688 rather than 718 applications, a reduction of 30, or 4.18%, under the common 2,048 opportunities. Both endpoints finish at confirmed epoch 10 for $b = 32$ and epoch 5 for $b = 64$; the corresponding emitted-byte costs are 565,288 and 407,596. The budget therefore has a measured exposure–service–communication tradeoff under this workload. The references give their own captured-state outcomes

at the same cut. For example, SPQR with 32-byte chunks derives 172 of 2,047 later commitments from capture A and 85 of 2,047 from capture B, reflecting the different pending material present in the two endpoint snapshots.

Communication policy and durability.

The native driver uses fair cyclic arbitration among queued control traffic, ACKs, exact application retries and new applications. An unacknowledged frame is immediately eligible for retry. Table 10 exposes that policy's cost: the alternating $b = 32$ execution spends 604 opportunities on application retries and the $b = 64$ execution spends 579. Independent trace analysis places 592 and 572 of those transmissions after the request's first logical delivery, accounting for 83,472 and 80,652 bytes. This uses the analyst's delivery history; the sender is still awaiting its ACK. The four frame classes partition each run's 2,048 emissions. Image-revision and `sync_data` counts measure the real native journal. Each alternating reference execution returns component state 4,096 times; its physical persistence fields are `null`, recording the component-serialization boundary. The communication totals include all native control, ACK and retry frames, and the references' complete component-plus-payload envelopes.

Transport controls.

The one-way schedule supplies only $A \to B$ opportunities, including for control traffic. Continuity commits and delivers 32 or 64 applications and then retains 2,016 or 1,984 offered requests in its queue. Both endpoints remain at confirmed epoch zero; all 31 or 63 post-capture commitments are derived from either captured state. The references continue committing 2,048 requests, with 2,047 of 2,047 post-capture keys derived. This control identifies budget-induced admission suspension and its backlog explicitly. Recovery during renewed exchange is assessed through the later admitted allocations and their original sources.

The lossy schedule drops the 205 common opportunities whose index is congruent to 4 modulo 10. Continuity delivers all 665 or 681 admitted applications, charging retransmissions to the opportunity budget. Each reference commits 2,048 component applications and delivers 1,843 under its declared workload without application retransmission. Network duplication leaves native first commitments, logical deliveries, emitted bytes, final epochs and derived sets equal to the alternating execution. Reordering, 9:1 asymmetry and a 256-opportunity offline prefix complete the seven-schedule inventory; their request queues, exact derived sets and work counters are retained in the complete 49-row artifact.

The manuscript tables are regenerated from the hash-bound executed report and checked against its independently recomputed public-evidence projection. Appendix A.13 identifies the source, execution and validation records.

7.5 Retry readiness under the common opportunity budget

The frame accounting above identifies application-retry readiness as a concrete scheduling parameter. A paired native experiment changes this parameter while retaining the protocol source, send budgets, capture cut, request stream, transport schedules and 2,048-opportunity budget of Section 7.4. It executes both the original immediate

Schedule	Delivered: immediate → age four		Δ bytes/delivery	
	$b = 32$	$b = 64$	$b = 32$	$b = 64$
Alternating	688 → 971	718 → 999	-11.41%	-19.70%
Lossy	665 → 905	681 → 965	+14.37%	-3.99%
Reordered	671 → 784	700 → 837	+3.94%	-12.07%
Duplicates	688 → 971	718 → 999	-11.41%	-19.70%
9:1 asymmetric	517 → 648	720 → 857	-32.12%	-35.76%
Offline then exchange	613 → 879	653 → 935	-22.57%	-25.39%
One-way	32 → 32	64 → 64	+0.09%	+0.10%

Table 11 Application-retry readiness at the same 2,048 directed opportunities. Each delivery column compares the freshly repeated immediate policy with minimum retry age four in the sender's own opportunities. Byte cost divides all emitted control, ACK, application and retry bytes by actual logical deliveries; positive changes indicate higher cost per delivery.

policy and an age-four policy for each of the two budgets and seven schedules, giving 28 traces. All 14 fresh immediate-policy rows agree with their retained counterparts. This ablation measures completed service and communication work in the native driver at fixed cryptographic and durability interfaces.

A public retry clock.

Each endpoint counts its own designated transmission opportunities. The counter advances after the opportunity and request are recorded and before local control planning. The driver records this counter on every first application emission and exact application retry. An unacknowledged request becomes retry-ready when the current counter exceeds its last emission counter by at least four. Within the existing fair cyclic arbitration, the retry class selects the oldest eligible unacknowledged request. Control planning, cumulative ACK processing, first-application admission and the remaining arbitration rules retain their original code. Retries obtain the retained ciphertext through the same resume operation. The readiness rule uses public scheduling and the sender's ACK state; delivery history is used only by the independent analyst when classifying transmissions.

Completed service and frame use.

Age four increases logical deliveries in all twelve budget–schedule pairs with bidirectional or resumed communication, by 16.84–43.39% (Table 11). Alternating traffic delivers 971 applications at $b = 32$ and 999 at $b = 64$, compared with 688 and 718 under immediate readiness: increases of 41.13% and 39.14%. Application retries fall from 604 to 16 and from 579 to 8. The respective 592 and 572 retries emitted after a request's first logical delivery fall to zero. These runs emit 706,764 and 455,388 total bytes; dividing the complete frame traffic by logical deliveries gives 11.41% and 19.70% lower byte cost per delivery. The two captured endpoints still each derive 63 or 127 post-capture keys, now among 970 or 998 post-capture first commitments. Thus the additional alternating service retains the observed current-epoch exposure count. Physical `sync_data` calls

rise from 14,219 to 20,015 and from 14,579 to 20,255; the respective counts per delivery change from 20.67 to 20.61 and from 20.31 to 20.28, retaining the real durability work in the service accounting.

Schedule-dependent accounting.

The common opportunity budget also exposes differences in communication cost. At $b = 32$, the lossy and reordered schedules deliver 36.09% and 16.84% more applications while their bytes per delivery rise by 14.37% and 3.94%. Age four leaves one committed application undelivered at the end of each lossy run and the $b = 64$ asymmetric run; the corresponding committed/delivered counts are 906/905, 966/965 and 858/857. All exposure denominators continue to count post-capture first commitments. Under the offline prefix, the derived counts change from 32 to 33 and from 64 to 65, with the respective denominators changing from 612 to 878 and from 652 to 934. These counts remain within the same $2b - 1$ allocation envelope while the actual allocation histories change.

The one-way control retains exactly 32 or 64 commitments and deliveries, the same queued demand, and confirmed epoch zero at both endpoints. Each run emits one additional 4,746-byte control frame and one fewer 141-byte application retry, adding 4,605 bytes. Retry readiness therefore improves the measured service under renewed exchange while the one-way admission boundary remains fixed by the send budget.

The artifact retains all fourteen paired rows, their original counts and denominators, the twenty-eight public traces' identities, and an independent recount of retry ages and frame outcomes. Appendix A.14 identifies the executed driver and the checks used to regenerate the table.

8 Consequences and further instantiation

A recovery condition attached to original operations.

The model distinguishes creation of an operation's random input, its first durable reservation, the acquisition of a retained representation, and the later availability of the matching decryption material. This distinction makes exact replay compatible with precise recovery statements: retries and aggregate writes preserve the original source, while fresh operations obtain new origins. The guarded characterization explains how records and keys combine to expose that source. The native executions demonstrate why both this conjunction and the original event history matter at the persistence and asynchronous-confirmation boundaries.

From a recovered epoch to application security.

The computational argument follows a selected original contribution through the complete tuple, root initialization, confirmation, traffic material and message chain. The joint comparison preserves real confirmation and ACK verification, successor-root disclosure and the selected message's next-chain-seed disclosure (Theorem 5.10). The online extension permits earlier message keys to influence the selected index, with selection fixed before its first private computation across all aliases (Theorem 5.12).

Exact retries reuse that original computation, and real suffix processing continues from the disclosed successor seed.

The native composition assigns these primitive interfaces to the actual send and receive order, including receive-side advances followed by failed authentication or persistence. Candidate-local stops retain the complete-image acquisition condition; the permitted individual key observations enter through their named interfaces. The resulting confidentiality and selected-message integrity bounds (Theorem 5.27) therefore apply to the original eligible application send. The selected-root authentication theorem also accounts for its earlier confirmation history (Theorem 5.25). Each terminal event is recomputed in the compared execution, with its own component adapter and whole-execution budgets. Table 3 records the shared dependencies.

Primitive instantiation and reusable interfaces.

For root derivation, the placement of a surviving secret determines the required comparison. Equation (1) gives

$$p_e = \mathrm{HMAC}_{R_{e-1}}(S_e).$$

Disclosure of the parent root leaves the new contribution in HMAC's message input, as studied by extractor and swap-PRF formulations [3, 21]. The shared-oracle entry proof in Appendix C.3 supplies the uniform-root interface after the complete-tuple comparison of Section 5.7. The finite-family extension preserves exact pad-state and information-string aliases (Section 5.14). Its purpose-key instance permits all outputs of the declared derivation table to be exported while retaining a separate lifecycle for each returned owner (Section 5.15). A marked private entry supports the corresponding one-key application experiment.

The other connection concerns the origin of an accepted allocation. Executed-slot registration ties it to the original KEM calls through retained plans and complete signed-body bindings (Proposition 5.20). The native signing adapter carries original randomizers and exact retries into a shared-XOF comparison, expressing the native reservation-interface bound through ordinary hedged ML-DSA unforgeability, with explicit replay-coupling and oracle-query losses (Theorem 5.24). These two connections instantiate different parts of the final bound: the key-entry argument supplies recovered secret material, and the signature bridge supplies its original-source attribution.

Exposure budgets and completed service.

The matched executions make the two-chain current-epoch capacity visible under a common communication budget. At the declared capture cut, budgets 32 and 64 yield 63 and 127 reconstructed future applications in alternating traffic, while completing 688 and 718 applications. Reporting the admitted denominator alongside the reconstructed set exposes the cost of more frequent rekey work. The one-way control reaches its admission cap with both confirmed epochs still zero and with every admitted post-capture key reconstructed. These observations connect the provenance condition to a concrete accounting of recovery, service and queued demand. The frame partition further identifies the selected immediate-retry policy's contribution to communication

cost. A paired native ablation changes application-retry readiness using a public sender-turn clock. Under the same 2,048 opportunities, it raises alternating deliveries to 971 and 999 while the derived counts remain 63 and 127. All twelve bidirectional or resumed budget–schedule pairs deliver more applications, with their full byte accounting reporting both reductions and increases in cost per delivery. This separates a controllable scheduling cost from the allocation envelope and the one-way admission boundary.

Conclusion.

Persistent recovery is governed by original secret sources and their acquisition histories. An exact record-exposure semantics and constructive record/report simulations connect an eligible native rekey to confidentiality and authentication of its marked message. Original-operation and shared-XOF reductions preserve source attribution through retained randomizers and exact replay. The finite key-entry comparison supplies a purpose-key instance, while the mechanized multi-epoch theorem and concrete executions establish the corresponding provenance-based recovery statements. These results preserve the meaning of fresh randomness across durable retries, copies and asynchronous progress.

Data and code availability

The protocol source and recorded executions are identified by immutable repository revisions in Appendix A. The companion artifact contains the finite-record checker, expected-result cases, source maps, proof notes and complete manuscript sources. Each recorded result retains its source identity and execution receipt.

AI-assisted preparation

OpenAI ChatGPT was used for language editing, exploratory proof development, source-code inspection and LaTeX preparation. The explicit proofs, mechanized certificates and executable checks are identified separately so that their respective claims and verification records can be inspected.

References

[1] Aranha DF, Orlandi C, Takahashi A, Zaverucha G (2020) Security of hedged Fiat–Shamir signatures under fault attacks. In: Advances in Cryptology—EUROCRYPT 2020, https://doi.org/10.1007/978-3-030-45721-1_23, URL https://eprint.iacr.org/2019/956

[2] Auerbach B, Dodis Y, Jost D, Katsumata S, Schmidt R (2025) How to compare bandwidth constrained two-party secure messaging protocols: A quest for a more efficient and secure post-quantum protocol. In: 34th USENIX Security Symposium (USENIX Security 25). USENIX Association, pp 6717–6736, URL https://www.usenix.org/conference/usenixsecurity25/presentation/auerbach

[3] Backendal M, Bellare M, Günther F, Scarlata M (2023) When messages are keys: Is HMAC a dual-PRF? In: Advances in Cryptology—CRYPTO 2023, Lecture Notes in Computer Science, vol 14083. Springer, pp 661–693, https://doi.org/10.1007/978-3-031-38548-3_22, URL https://eprint.iacr.org/2023/861

[4] Bellare M, Tackmann B (2016) Nonce-based cryptography: Retaining security when randomness fails. Cryptology ePrint Archive, Report 2016/290, URL https://eprint.iacr.org/2016/290, author version of the EUROCRYPT 2016 paper

[5] Bennett CH, Bernstein E, Brassard G, Vazirani U (1997) Strengths and weaknesses of quantum computing. SIAM Journal on Computing 26(5):1510–1523. https://doi.org/10.1137/S0097539796300933, URL https://arxiv.org/abs/quant-ph/9701001

[6] Bhargavan K (2025) Helping secure Signal's post-quantum transition. Cryspen, URL https://cryspen.com/post/signal-spqr-verification/, october 2, 2025

[7] Bhargavan K, Buyse M, Connell G, Kiefer F, Schmidt R (2026) A formal analysis of Signal's sparse post-quantum ratchet protocol and its implementation. In: Proceedings of the 2026 ACM SIGSAC Conference on Computer and Communications Security. ACM, URL https://www.sigsac.org/ccs/CCS2026/program/schedule.html, to appear; listed in the official conference program

[8] Chajed T, Tassarotti J, Theng M, Jung R, Kaashoek MF, Zeldovich N (2021) GoJournal: A verified, concurrent, crash-safe journaling system. In: 15th USENIX Symposium on Operating Systems Design and Implementation (OSDI 21). USENIX Association, pp 423–439, URL https://www.usenix.org/conference/osdi21/presentation/chajed

[9] Cheval V, Jacomme C, Richards J (2026) Automated formal analysis of Signal's double ratchet: Attacks, fixes and security proofs. In: 2026 IEEE Symposium on Security and Privacy (SP). IEEE, pp 1522–1538, https://doi.org/10.1109/SP63933.2026.00120, URL https://mrg.cs.ox.ac.uk/publications/automated-formal-analysis-of-signals-double-ratchet-attacks-fixes-and-security-proofs/

[10] Cohn-Gordon K, Cremers C, Dowling B, Garratt L, Stebila D (2020) A formal security analysis of the Signal messaging protocol. Journal of Cryptology 33(4):1914–1983. https://doi.org/10.1007/s00145-020-09360-1, URL https://www.douglas.stebila.ca/research/papers/JC-CCDGS20/

[11] Cremers C, Jacomme C, Naska A (2023) Formal analysis of session-handling in secure messaging: Lifting security from sessions to conversations. In: 32nd USENIX Security Symposium (USENIX Security 23). USENIX Association, pp 1235–1252, URL https://www.usenix.org/conference/usenixsecurity23/presentation/cremers-session-handling

[12] Cremers C, Medinger N, Naska A (2025) Impossibility results for post-compromise security in real-world communication systems. In: 2025 IEEE Symposium on

Security and Privacy (SP). IEEE, URL https://eprint.iacr.org/2024/1886

[13] Damm S, Kraus N, May A, Nowakowski J, Thietke J (2026) One (noisy) bit to rule them all: Key recovery from randomness leakage in ML-DSA. Journal of Cryptology 39(4):28:1–28:45. https://doi.org/10.1007/s00145-026-09584-7, URL https://ir.cwi.nl/pub/36609/36609.pdf

[14] Dodis Y, Jost D, Katsumata S, Prest T, Schmidt R (2025) Triple ratchet: A bandwidth efficient hybrid-secure Signal protocol. In: Advances in Cryptology—EUROCRYPT 2025, Lecture Notes in Computer Science, vol 15608. Springer, pp 302–331, https://doi.org/10.1007/978-3-031-91101-9_11, URL https://cs.nyu.edu/~dodis/ps/triple-ratchet.pdf, extended version, Cryptology ePrint Archive, Paper 2025/078

[15] Don J, Fehr S, Huang YH (2022) Adaptive versus static multi-oracle algorithms, and quantum security of a split-key PRF. In: Kiltz E, Vaikuntanathan V (eds) Theory of Cryptography—TCC 2022, Proceedings, Part I, Lecture Notes in Computer Science, vol 13747. Springer, pp 33–51, https://doi.org/10.1007/978-3-031-22318-1_2, URL https://arxiv.org/abs/2206.08132, author manuscript available as arXiv:2206.08132

[16] Fischlin M, Günther F, Janson C (2024) Robust channels: Handling unreliable networks in the record layers of QUIC and DTLS 1.3. Journal of Cryptology 37. https://doi.org/10.1007/s00145-023-09489-9, URL https://link.springer.com/article/10.1007/s00145-023-09489-9, article 9

[17] Grilo AB, Hövelmanns K, Hülsing A, Majenz C (2021) Tight adaptive reprogramming in the QROM. In: Advances in Cryptology—ASIACRYPT 2021, https://doi.org/10.1007/978-3-030-92062-3_22, URL https://eprint.iacr.org/2020/1361

[18] Hérouard C, Jacomme C, Koutsos A, Lallemand J (2026) Secrecy in Squirrel and the post-compromise security of a ratchet. In: Proceedings of the 2026 ACM SIGSAC Conference on Computer and Communications Security. ACM, URL https://eprint.iacr.org/2026/1912, to appear; full version in Cryptology ePrint Archive, Paper 2026/1912

[19] Hosoyamada A, Iwata T (2021) On tight quantum security of HMAC and NMAC in the quantum random oracle model. In: Malkin T, Peikert C (eds) Advances in Cryptology—CRYPTO 2021, Proceedings, Part I, Lecture Notes in Computer Science, vol 12825. Springer, pp 585–615, https://doi.org/10.1007/978-3-030-84242-0_21, URL https://eprint.iacr.org/2021/774, full version: Cryptology ePrint Archive, Paper 2021/774

[20] Kosuge H, Xagawa K (2025) The security of ML-DSA against fault-injection attacks. In: Advances in Cryptology—ASIACRYPT 2025, https://doi.org/10.1007/978-981-95-5096-8_20, URL https://eprint.iacr.org/2025/904

[21] Krawczyk H, Eronen P (2010) HMAC-based extract-and-expand key derivation function (HKDF). RFC 5869, Internet Engineering Task Force, URL https://www.rfc-editor.org/rfc/rfc5869.html

[22] Krawczyk H, Bellare M, Canetti R (1997) HMAC: Keyed-hashing for message authentication. RFC 2104, Internet Engineering Task Force, URL https://www.rfc-editor.org/rfc/rfc2104

[23] National Institute of Standards and Technology (2024) Module-lattice-based digital signature standard. Federal Information Processing Standards Publication 204, U.S. Department of Commerce, https://doi.org/10.6028/NIST.FIPS.204, URL https://nvlpubs.nist.gov/nistpubs/fips/nist.fips.204.pdf

[24] Perrin T, Marlinspike M, Schmidt R (2025) The double ratchet algorithm. Signal specification, URL https://signal.org/docs/specifications/doubleratchet/, revision 4, November 4, 2025

[25] Pornin T (2013) Deterministic usage of the digital signature algorithm (DSA) and elliptic curve digital signature algorithm (ECDSA). RFC 6979, RFC Editor, https://doi.org/10.17487/RFC6979, URL https://www.rfc-editor.org/rfc/rfc6979.html

[26] Q-Periapt project (2026) Independent joint replay of multi-epoch symbolic message secrecy. Q-Periapt research artifact, URL https://github.com/billlza/q-periapt/tree/228c944fc476e3e57d6c08063cacb19bca5b360e/research/sdk-alpha1/evidence/20261104-multi-epoch-replay, november 4 replay; fixed repository checkpoint 228c944fc476e3e57d6c08063cacb19bca5b360e

[27] RustCrypto Developers (2026) ecdsa 0.17.0: Signing implementation. URL https://docs.rs/ecdsa/0.17.0/src/ecdsa/signing.rs.html, version pinned by the inspected native dependency lockfile

[28] Zhandry M (2015) A note on the quantum collision and set equality problems. Quantum Information and Computation 15(7–8):557–567. https://doi.org/10.26421/QIC15.7-8-2, URL https://www.rintonpress.com/xxqic15/qic-15-78/0557-0567.pdf

A Source and evidence map

The source snapshot for this manuscript's established protocol and experiment records is commit `228c944fc476e3e57d6c08063cacb19bca5b360e` of `billlza/q-periapt`. Each retained execution keeps its original source and tool identity. The local record model and accounting proofs in this manuscript are mathematical arguments; the following entries identify the independently recorded mechanization and native observations.

Result	Reproduction source
Joint symbolic replay	20261104-multi-epoch-replay. The original archive is `CAPTURE.zip`; the checked certificate member is `replay-01/01-baseline-replay.proof.spthy`.
Positive two-round trace	20261105-two-round-ground-trace.
Delayed-receipt ground trace	20261105-delayed-receipt-ground-trace.
Native reservation disclosure	RESERVATION_DISCLOSURE.md and `src/durable/messages/tests/reservation_disclosure.rs` in the identity candidate.
Aggregate-copy execution	20261105-aggregate-reservation-copy.
Root, traffic and phase mapping	RECOVERY_DEPENDENCIES_V1.md; native `rekey/response.rs` and `rekey/completion.rs`.
SPQR snapshot corpus	continuity-spqr-reference: `TRACE_CORPUS.json`, `COMPROMISE_CORPUS.json`, and `COMPROMISE_EXPERIMENT.md`.
Whole-KEM corpus	continuity-whole-kem-reference: `TRACE_CORPUS.json`, `README.md`, and `ACTIVE_COMPROMISE.md`.
Computational component arguments	COMPUTATIONAL_RECOVERY_V1.md and its linked KEM adapter, hidden-record, shared-HMAC, root-extraction and adaptive-target notes.

The complete seven-schedule reference comparison is reproduced by `artifacts/reference-comparison/analyze_reference.py` with `-check`. Its four retained corpus files are bound to the source snapshot above by Git blob and SHA-256 identities in `artifacts/reference-comparison/source_identity.json`. `analysis.json` in that directory records all 420 snapshot cases and the 35 configuration–schedule rows; the generated `reference-schedule-table.tex` supplies Table 8. The independent review checks every displayed count and decimal mean using separate integer arithmetic and rejects eight altered-input controls.

The jointly checked certificate has SHA-256

`dde5d06a071b1560da0d4c28c3f1fd81a3c39708f9fd8564ee391ddc362bcf7f`.

Its source-base receipt identifies `6084f3bfb09333f6b66291b38809ed512296a27b`. The fixed 4 November replay retained all 40 queries and its exact tool, resource and output receipts. Reading and checking this retained artifact while preparing the manuscript is distinct from executing another Tamarin run.

A.1 Mechanization and replay record

The theorem is the retained query `multi_epoch_fresh_reservations_recover_symbolic_message_secrecy`. Its certificate was replayed with all 24 reusable dependencies

in the same Tamarin process. The independent 4 November replay used Tamarin 1.12.0 and Maude 3.5.1, one worker, a 12-GiB heap bound and the unchanged 900-second limit; it completed in 816.903 seconds. The output certificate was byte-identical to the retained input. Appendix A gives its digest and archive member, making the theorem's exact formula and proof available for inspection.

The complete inventory contains 40 queries: 27 verified and 13 incomplete. The latter comprise eleven multi-round existence queries and two historical all-trace queries outside the main theorem's dependency group. Two subsequent finite ground-trace constructions supply the concrete executions discussed below. Their independently checked formulas and original certificate inventory remain separately identifiable.

A.2 Guarded-normalization supplement

The complete manuscript extends the previously retained source at commit

`1c68aa666db22d36ee136be1ec712545a9e74f03`

under `paper/continuity`. The manuscript and its directly used evidence are maintained in the repository's continuous manuscript PR. The filenames below are relative to `paper/continuity` in that distribution.

The explicit normalization theorem and its proof are in `guarded-provenance.tex`; the conjunctive-exposure proposition and proof are in `acquisition-nonadditivity.tex`. The executable supplement is in `artifacts/guarded-provenance/`:

- `check_exposure.py` evaluates both record semantics; `results.json` records the finite enumeration and history checks.
- `review_exposure.py` checks 17 independently stated expected results; `review_results.json` records their outcomes.
- `review.md` documents the theorem and implementation review; `review_optimized_guard.json` records fail-closed optimized-mode rejection. The separately named `review_optimized_run_before_guard.json` retains the earlier execution that motivated the guard.
- `native_source_map.json` records source correspondence for the journal-to-wrapping-key derivation, durable image serialization, pending targets and the retained aggregate-copy test. It pins five source files at commit `228c944`; this inspection records no new native execution.

From that artifact directory, the reproduction commands are

```
python3 check_exposure.py --output results.json
python3 review_exposure.py
```

These commands exercise the finite record model and its independent expected cases. The retained native and Tamarin evidence remains at the source checkpoints listed in the preceding map.

A.3 Constructive image/intent simulation

Section 5.1 integrates the complete finite-interface algorithm and proof from `artifacts/record-rendering/opaque-renderer-v1.2.md`. The corresponding mathematical review is in `artifacts/record-rendering/independent-review.md`. The retained source modules defining the commitment experiment are HIDDEN_RECORD_COMMITMENTS_V1.md and MULTI_RECORD_COMMITMENTS_V1.md. Their actual domains and complete oracle budgets define the ϵ_{com} term; the new table construction supplies the image/intent implementation after that replacement.

The source correspondence of Section 5.4 is recorded in `artifacts/native-interface/native_interface_map.md` and `artifacts/native-interface/map.json`. The machine-readable map binds the source ranges and file identities to the fixed protocol checkpoint. `artifacts/native-interface/native_map_validation.json` records the source checks. These are source observations associated with the declared interfaces; historical execution receipts retain their own dates and source identities.

A.4 Immutable report composition

Section 5.2 gives the finite-interface algorithm, reverse-order theorem and real-disclosure corollary. Its complete proof note is `artifacts/report-composition/multi_report_theorem.md`; the independent review is `artifacts/report-composition/multi_report_review.md`. The fixed-source CLEANUP_REPORT_MAC_V1.md records the fixed-source one-row interface. The paper states its operational game in Section 5.2 and its primitive instantiation in Appendix B. SHARED_ROOT_HMAC_V1.md specifies the compression-oracle functions and their query regimes. In the paper's notation, Φ_{256} bounds the two secret HMAC key-pad replacements, and E_{NMAC} bounds the ensuing row replacement with the stated length and query budgets. The composition theorem accepts the one-row distinguishing bound as a modular input and adds the complete ordinary-query work of every derived reduction.

A.5 Candidate-local computational connection

Section 5.3 incorporates the frozen proof fragment paired with `artifacts/candidate-local/candidate_local_composition.md`. The independent derivation and its artifact identities are recorded in `artifacts/candidate-local/independent_review.md` and `artifacts/candidate-local/independent_review.json`. They check the all-continuations stopping certificate, the exact native-score identity, the selected-candidate residual, and the local component comparisons. The score identity applies to bounded scores; the positive/negative-event conversion is specialized to the message experiment's three-valued score.

A.6 Joint schedule and live recovery profile

Section 5.5 and Appendix C give the joint schedule and its shared-oracle primitive derivations. The complete proof record is `artifacts/joint-continuation/joint_continuation.md`; its separate review and frozen source identities are retained in the same directory. The live recovery profile in `artifacts/live-profile/source_profile.md`

maps the actual root, confirmation, traffic, message and ACK calls, together with the private application and persistent-record interfaces. The map and encoding inventory record the exact full information strings and source locations. Its 24 source files are byte-identical between the protocol pin `228c944f` and the inspected development revision `1c68aa66`; this identity check is separate from the retained native execution receipts.

The live-core instantiation and joint authentication consequence are retained in `artifacts/confirmation-traffic/`. Their proof and separate reviews record the complete observation surface, the reservation-selector witness, the native confirmation cores and the original encryption ledger. The source map fixes each inspected file at the protocol checkpoint. The authentication argument starts at the uniform-root interface and assigns its own terminal event to any preceding comparison.

A.7 ContextBound and original-operation inputs

The core `combine_policy_bound` and `combine_validated` functions in `crates/q-periapt-core/src/lib.rs` fix the length-delimited component and context encoding. The SDK replay constructor's `reserve`, `open` and `binding` operations in `crates/q-periapt-sdk/src/expert/replay.rs` retain the admitted operation and its original tape. The rekey response module's `kem_context` in `research/continuity-identity-candidate/src/durable/messages/rekey/response.rs` binds the complete retained offer into the application context. The artifact also retains an earlier Q-Periapt manuscript on ContextBound binding and implementation assurance. Those materials document the reused cryptographic substrate; the original reservation, acquisition and persistent simulation results are developed in the present paper.

The connection artifact, `artifacts/contextbound-connection/`, contains the complete proof, independent review and source induction for the next rekey of one established session. Its source map fixes fourteen files at `228c944f` and the inspected development checkpoint. The self-contained checker recomputes six retained core vectors and exercises complete-tuple aliases, original-operation retries and component-challenge routing, with explicit incorrect-cache controls. The source audit is a separate command requiring the declared native revisions; its receipt records source identity and static assertions.

Relation to earlier project material.

Earlier project work on SkyBridge Compass studies policy-governed authenticated establishment, complete-transcript binding, and a control-only recovery phase before application traffic (unpublished manuscript, revision `d49d2a9`). It supplies the operation-admission and traffic-release perspective. Here we follow the admitted operation's original randomness through persistence and acquisition, and develop record and report simulations. The artifact's `notes/skybridge-scope.md` fixes the source versions and allocates material between these research questions.

A.8 Finite multiple-root entry

Section 5.14 gives the finite controller, joint root-entry comparison and first-message corollary. The complete argument, including the class-allocation reduction, query

budgets and a constructive controller, is in `artifacts/multiple-root-entry/multiple_root_entry_proof.md`. The same directory contains the frozen TeX module, independent derivation, 25-point mathematical review and exact file identities. The controller permits available unmarked roots to become later salts; the marked root and all its aliases use the retained schedule handle. The review checks this interface and its composition with the preceding primitive bounds.

The prose correspondence record in `artifacts/editorial-review/` links the current presentation to the retained mathematical inputs. Historical proof snapshots keep their exact identities; the semantic review checks the current definitions, hypotheses, interfaces and budgets against them.

A.9 SDK purpose-key correspondence

Section 5.15 uses the owned purpose-key API in `crates/q-periapt-sdk/src/purpose.rs`. Its complete proof and source map are retained in `artifacts/sdk-purpose-connection/`. The map checks the effective information fields, HKDF calls, owner guards and retained native tests at the protocol pin and development checkpoint. The accompanying review checks the variable-length compression budgets, complete-input aliases, owner transitions and selected-key experiment.

The Rust check job at development revision `1c68aa66` passed all five purpose-key tests on 7 November 2026. The retained receipt lists their outcomes, covering the RFC 5869 vector, domain separation, dual-peer derivation, runtime closure, input admission and trust-root binding. The source correspondence and mathematical review record the additional manuscript claims.

A.10 Executed sources and signing-reservation interfaces

Section 5.10 has its complete proof, source identities and independent mathematical review in `artifacts/source-attribution/`. The directly inspected native revision is `72acaaccde1ca9ea27e5af022da5fbce932ed4e9`; five of its seven inspected files additionally match the earlier native-pin receipt. The source map records the exact reservation/persistence/computation order, complete-wire bindings, original-signing-randomizer encoding and the separate signer-seed-file interface.

The executable model explores 34 recovered states for generation and 34 for encapsulation, with up to three original draws and two epochs. Its 486 call-boundary edges include failed writes with absent or retained targets, private-computation failure, process loss, release denial and completion. Six deliberate source-tracking mutations have explicit counterexamples. An independent expected matrix covers offer, response and final signature aliases; the response alias permits proposer traffic allocation with unchanged sources and changed complete traffic information. The signature checks preserve original randomizers and private-query ledgers across unpublished results, and enumerate the generic separation's 64 small key/randomizer assignments. These are finite interface checks and exact source audits. Their receipts keep the native Rust and mechanized-proof execution status separate.

A.11 Reservation-to-ordinary signature connection

Section 5.11 has its full reduction, source identities, finite controls and independent internal argument review in `artifacts/signature-bridge/`. The native correspondence uses revision `358b921ab3469ce49f166633fc9a762ff6ed7e3d` and the lockfile-authenticated `fips204` 0.4.6 source archive. Its upstream revision is `8d456d989dce681233ef963d3c3306da8f76d7c7`. The retained packet includes the inspected sources and their licenses, so its identity and call-length checks run from the manuscript checkout.

The controls distinguish exact pair caching from per-slot oracle queries, preserve original-message freshness under descriptor aliases, and record private invocations whose publication fails. Independent finite enumeration checks the shared-XOF construction's joint distribution and its disjoint hidden-input sectors. Owner controls distinguish declared lineage, actual signing-state identity and public-key equality, and exercise non-target disclosure. The mathematical review covers the complete two-comparison ROM/QROM argument and its ordinary native signing endpoint; executable controls validate the specified finite distributions and source interfaces.

A.12 Native authentication composition

Section 5.12 gives the complete selected-root authentication proof. Its argument review and executable accounting controls are retained in `artifacts/native-authentication/`. The controls check the native Boolean decomposition on 160 compatible assignments, including free and charged stops. Separate finite examples distinguish the selected authentication event from the confidentiality score events, joint eligibility from conditional probability, and the selected root from other root histories. The proof uses the source and primitive correspondences identified in the preceding supplements; its event-bit comparisons retain their own complete continuation budgets.

A.13 Matched directed-opportunity executions

Section 7.4 uses research revision `fec8472060728cc1bb35cf15042c117dda00e336`, executed by GitHub run 42537214047. Its checkout is the synthetic merge `0d12e77d`; preparation archives protocol source `358b921a` into an isolated execution tree. The retained preparation and executable receipts identify the native driver, SPQR and its 64-byte chunk variant, and the whole-KEM reference. The run completed all 49 profile–schedule pairs.

`artifacts/matched-comparison/` retains the complete executed report, its 49-row JSON/CSV projection, original public-audit and qualification receipts, source digests and deterministic manuscript-table generator. The compressed report preserves the exact original JSON bytes. The generator checks complete profile/schedule and capture-owner inventories, admission and exposure partitions, frame budgets, full-envelope bytes and the distinction between native physical persistence and reference state returns. Input controls exercise altered denominators, duplicate rows, inconsistent request sets, fabricated physical counters and omitted frame-envelope work.

The original public audit recomputed 14 native traces, 35 reference trace/snapshot cases and 35 public envelope schedules, and checked all 126 public-evidence manifest entries. The hosted run performed the real native plaintext and reference AEAD key checks. The retained qualification receipt binds archive 13641008832, its SHA-256, the exact source and execution, and these verification outcomes. The manuscript projection preserves those execution identities and can be regenerated from the retained packet alone:

```
python3 artifacts/matched-comparison/analyze_matched.py --check
python3 artifacts/matched-comparison/check_input_controls.py
```

A.14 Retry-readiness ablation

Section 7.5 uses research revision `29f1c362b67bcd06123ea1bef847af82e9f9e100`, executed by GitHub run 42551086845 at synthetic checkout `23f600f5`. The fixed protocol source remains `358b921a`. Its source comparison covers 2,687 files and identifies only the addition of the two experiment modules to the native test file; the original immediate-policy driver retains its exact bytes. The executed test binary is hash-bound before and after both policy runs.

`artifacts/retry-readiness/` retains the exact compressed execution report, source and executable identities, fourteen paired result rows, the source/readiness control log and the independent public-trace audit. The source-bound provenance receipt identifies job 128946601334 and archive 13650034930, checks all 42 archive members, and records fourteen passing source/readiness controls. The trace audit checks 29 public evidence files, recounts all 28 executions, and matches all fourteen fresh immediate-policy rows to the retained matched comparison.

The table generator checks the two policies' complete profile/schedule inventories, application and exposure partitions, emitted-frame classes and bytes, public retry ages and physical persistence counters. It compares the executed report with the independent recount and the fourteen-pair result projection before emitting the manuscript table. The complete archived counts distinguish first commitments from deliveries and retain each post-capture denominator. Reproduction from the compact packet is:

```
python3 artifacts/retry-readiness/analyze_retry.py --check
```

A.15 Online finite-chain selection

The complete proof and source correspondence for Theorem 5.12 are retained in `artifacts/chain-prefix/`. Three native files retain identical source identities at the original pin `228c944f` and inspected integration checkpoint `6a0ac0bb`. They identify the fixed-width chain derivation, zero nonce, exact outbox resumption and transaction-private skipped-key advancement.

Exact finite enumeration covers all 128 online marking policies over a binary three-step ideal chain, together with eight four-valued policies. The selected key remains uniform conditional on earlier returned keys, the mark and the disclosed successor. Separate controls exhibit retrospective-selection bias, reconstruct later keys from an exposed prefix seed, and reject marking after an alias's first private

computation. These controls validate the selection and source interfaces used by the mathematical comparison.

```
python3 artifacts/chain-prefix/check_prefix.py
```

A.16 Native online-chain composition

The original-experiment composition of Theorem 5.27 is retained in `artifacts/native-online-chain/`. Its source map binds six native files to the original evidence pin `228c944f` and inspected integration checkpoint `a4c08957`. The identical source blobs fix pending-plan retention before encryption, directional aliases, transaction-private receive advancement, skipped-key reuse and the complete originating KEM plans.

Finite controls enumerate 19,440 traces with marking choices and 77,760 first-computation decisions. Independent temporal and algebraic enumerations check 450 reservation/cut/version assignments and 256 signed-score or Boolean-event partition identities. Counterexample controls exercise a rolled-back receive ledger, collapsed material identities, complete-image acquisition, candidate-local stopping and a fresh receive event followed by later honest encryption. Eighteen source-order anchors bind these interfaces to the retained source. The complete native bounds are proved in the manuscript; the controls check the finite accounting and source interfaces used by that proof.

```
python3 artifacts/native-online-chain/check_native_chain.py
```

B Primitive inputs to the record and report bounds

The composition theorems use operational distinguishing bounds: a record commitment comparison and a single immutable report row. This appendix fixes the concrete parameter conventions. The public hash and compression oracles are idealized in the stated games; the source construction uses the specified SHA3 and HMAC-SHA256 algorithms.

B.1 Hidden-input random-oracle replacement

For a hidden-input entropy parameter n, define

$$\Phi_n(d, Q) = \min\left\{1, 4\sqrt{2d^2Q\,2^{-n}} + 4\sqrt{2Q^2d\,2^{-n}}\right\}.$$

Don, Fehr and Huang's Theorem 2 [15] supplies this bound when private oracle inputs have point probability at most 2^{-n} and distinct effective private queries encode distinct inputs for every secret value. Exact repeats are cached before that interface is queried. The two budgets count private classical queries and all ordinary public-oracle queries; public queries may be quantum.

For image i, let K_i^{hdr} count original intent prefixes and let $d_i \le 1 + 5K_i^{\mathrm{hdr}}$ count its distinct private digest inputs. There is one image-head domain and five intent-digest domains. With an independent 192-bit seal nonce, the fixed-input encoding and

per-nonce confinement conditions give the retained record bound

$$\epsilon_{\mathrm{com}} \le \min\left\{1, \sum_{i=1}^{M} \Phi_{192}(d_i, Q_i)\right\}.$$

Here $d_i = 0$ for known, atom-free rows. The budget Q_i includes the entire derived distinguisher, all other real rows and the final test. The nonce remains private before the declared acquisition stop. Local image encryption and intent authentication introduce no additional shared SHA3 calls. This comparison is made after private-wire equality has been replaced by structural equality; its separate collision term is the one in Corollary 5.3. The exact encoding domains and fixed-source arguments are identified in Appendix A.3.

B.2 One immutable HMAC row

For the row experiment in Section 5.2, let

$$h : \{0,1\}^{512} \times \{0,1\}^{256} \longrightarrow \{0,1\}^{256}$$

be the public compression oracle, with arguments in message-block, chaining-state order. A 32-byte key K determines the two pad blocks

$$I(K) = (K\|0^{32}) \oplus \mathtt{36}^{64}, \qquad O(K) = (K\|0^{32}) \oplus \mathtt{5c}^{64}.$$

Zero exponents in these encodings count bytes. The ranges are disjoint and each encoding is injective in K. Cache the two evaluations at the fixed SHA-256 initial state. All oracle-dependent preparation is executed and counted; no oracle-correlated advice is supplied for free.

Let L be the maximum message length in bytes and put

$$L \le 2^{61} - 65, \qquad \ell = 1 + \left\lceil \frac{L+9}{64} \right\rceil.$$

The SHA-256 padding encodes $8(64 + |m|)$ bits, retaining the already consumed 64-byte key-pad block. Its suffix is determined by $|m|$ and it is injective within the stated length domain. The parameter ℓ bounds the remaining row compression work after the two key-pad calls.

Write $E_{\mathrm{NMAC}}(T, q, \ell)$ for a valid NMAC distinguishing bound with two independent 256-bit inner/outer keys, the prescribed injective padding, at most T keyed calls and q public compression calls. Hosoyamada and Iwata's Proposition 6 [19] gives the asymptotic form

$$E_{\mathrm{NMAC}}(T, q, \ell) = O\left(\sqrt{\frac{(q+T)^3 \ell^5}{2^{256}}}\right)$$

under its stated padding and query conditions. The implicit constant remains part of that published bound. Our message-length convention retains HMAC's consumed-block offset.

The two key-pad replacements use $\Phi_{256}(2, q + \ell T)$. The subsequent NMAC replacement gives an admissible row bound

$$\epsilon_{\mathrm{row}}(T, q, L) = \min\{1, \Phi_{256}(2, q + \ell T) + E_{\mathrm{NMAC}}(T, q, \ell)\}\,.$$

This is a restricted single-transcript application of the primitive results. The ideal NMAC function is queried on one original transcript only, so its first reply is a uniform tag and exact repeats retrieve the same value. Report nonzero-ID checks act on that reply in both games. All other real HMAC work remains in q; the multi-report theorem supplies the dependency order and its actual disclosure residual.

C Shared-oracle primitives for the joint continuation

This appendix supplies the primitive reductions used in Section 5.5. They use the same idealized SHA-256 compression oracle throughout. Each experiment begins with independent oracle, hidden uniform 256-bit input and auxiliary state. All oracle-dependent preparation is executed within the counted algorithm. The designated hidden value is accessed through its stated interfaces. It and the private random tables remain internal throughout terminal testing. Ordinary public queries may be quantum; private function queries are classical and exact repeats are cached.

We use Φ_{256} and E_{NMAC} from Appendix B, with

$$\ell(L) = 1 + \left\lceil \frac{L + 9}{64} \right\rceil, \qquad L \le 2^{61} - 65.$$

The HMAC padding includes the consumed 64-byte key pad: its encoded length is $8(64 + |m|)$. Equal padded strings have equal encoded lengths and hence equal message prefixes within this domain. Thus the required padding is injective, and $\ell(L)$ bounds the ordinary row compression work after its two key-pad evaluations.

C.1 A shared HMAC row and column

For a 32-byte value x, define the 64-byte blocks

$$\begin{aligned} I(x) &= (x \| 0^{32}) \oplus \mathtt{36}^{64}, \\ O(x) &= (x \| 0^{32}) \oplus \mathtt{5c}^{64}, \\ B(x) &= x \| P, \qquad P = \mathtt{80} \| 0^{23} \| \,\mathrm{BE64}(768). \end{aligned}$$

The final 32 bytes of $I(x), O(x), B(x)$ are respectively $\mathtt{36}^{32}, \mathtt{5c}^{32}, P$. These three ranges are disjoint, and each map is injective. Write

$$\alpha(a) = h(I(a), \mathrm{IV}), \qquad \beta(a) = h(O(a), \mathrm{IV}), \qquad \Gamma(a) = (\alpha(a), \beta(a)).$$

The actual 32-byte-message column is

$$F(a, R) = h\big(B(h(B(R), \alpha(a))), \beta(a)\big). \tag{43}$$

The actual row is the NMAC-form expression

$$F(R,m) = \mathrm{Iter}_h\big(\beta(R), \mathrm{pad}_{64}\big(\mathrm{Iter}_h(\alpha(R), \mathrm{pad}_{64}(m))\big)\big)\,,$$

where pad_{64} retains the consumed-key-pad length offset.

For completeness, $E_{F_1}(C,q)$ denotes a valid distinguishing bound for

$$F_1(u,v) = \widetilde{h}(v, f(u))$$

versus a private random function of the complete pair (u,v). Here f is an independent private random function from 256 to 256 bits; $\widetilde{h}$ is a public random function on 512-bit v and 256-bit final input; there are C classical F_1 calls and q ordinary quantum queries. Hosoyamada and Iwata's Proposition 4 [19] gives

$$E_{F_1}(C,q) = O\left(\sqrt{\frac{q_*^3}{2^{256}}}\right), \qquad q_* = \max\{C,q\} = o(2^{256/3}),$$

with the published constants and regime retained.

Lemma C.1 (Independent-table shared-root bound) *An admissible operational bound is*

$$\begin{aligned}\varepsilon_{\mathrm{Hjoint}}(C,T,q,L) \le \min\{1,\ &\Phi_{256}(C+2, q+3C+\ell(L)T)\\ &+ E_{\mathrm{NMAC}}(T, q+3C, \ell(L))\\ &+ E_{F_1}(C, q+2C)\}.\end{aligned} \tag{44}$$

The ideal row and column are independent private functions $U(m)$ and $V(\Gamma(a))$, independent of the public oracle.

Proof First isolate every compression call with a raw occurrence of R:

$$\mathcal{F}_R(t,s) = h(\tau_t(R), s), \qquad \tau_{\mathrm{in}} = I, \quad \tau_{\mathrm{out}} = O, \quad \tau_{\mathrm{msg}} = B.$$

For a fixed (t,s), each encoded point has probability at most 2^{-256}. Different effective pairs encode different points for every R: different tags have disjoint block ranges, and equal tags retain the chaining state. Cache by the complete pair (t,s), not just by the salt. This matters when two salts have the same α value.

There are at most $C+2$ distinct private calls: the two root key-pad evaluations and at most C message-block evaluations at $\alpha(a)$. The remaining ordinary work is at most $q+3C+\ell(L)T$, including each column's two public pad states and outer call, and each row's continuation. Don, Fehr and Huang's Theorem 2 [15] therefore replaces $\mathcal{F}_R$ by a private random function at the first cost in (44).

Its restrictions

$$A = \mathcal{F}(\mathrm{in}, \mathrm{IV}), \quad D = \mathcal{F}(\mathrm{out}, \mathrm{IV}), \quad f(s) = \mathcal{F}(\mathrm{msg}, s)$$

are independent of one another and h. The row is now NMAC with independent inner and outer keys A, D. The simultaneous column is $h(B(f(\alpha(a))), \beta(a))$. To replace the row, the NMAC reduction samples its own f table and answers a column with three ordinary compression

calls. Proposition 6 of Hosoyamada and Iwata [19], with the prescribed offset padding, gives the second term and a private row function U independent of f, h.

For the column, reindex the *same* oracle by

$$\widetilde{h}((t,s),y) = h(y\|t,s), \qquad |t| = |s| = |y| = 256.$$

This is a fixed bijection on all 768 input bits. It is implemented coherently by wire permutation and one oracle call, not by introducing an independent hash. The remaining column is exactly

$$\widetilde{h}\big((P,\beta(a)), f(\alpha(a))\big) = F_1\big(\alpha(a), (P,\beta(a))\big).$$

Computing its public states costs two ordinary calls per column. The independent row table U needs no compression queries. The F_1 bound therefore applies with C private calls and at most $q + 2C$ ordinary calls. Its ideal reply is indexed by $(\alpha(a), P, \beta(a))$, equivalently by $\Gamma(a)$ since P is fixed. Equal Γ values must share a reply; their equality is never assumed to imply equal salts. Summing the three replacements proves the result. □

The lemma compares against the independent-table ideal. In particular it does not assert pointwise real equality at an exceptional cross-interface pair where both classical arguments equal the hidden root. Such guesses are admitted by the real experiment and covered by the three distinguishing steps. A different ideal that explicitly repairs this cross-pair needs its own comparison; it is not used here.

C.2 The multiple-message row and a single Extract column

For $\varepsilon_{\mathrm{Hrow}}(T, q, L)$, the adversary has an adaptive classical row oracle. It can choose later messages using earlier replies and its quantum computation with the common public oracle. Equal messages retrieve equal values. There is no requirement that the full transcript be fixed before the first reply, and there is no raw-key disclosure channel.

Lemma C.2 (Row and single-column instances) *The following bounds instantiate the operational games:*

$$\varepsilon_{\mathrm{Hrow}}(T, q, L) \le \min\{1,\ \Phi_{256}(2, q + \ell(L)T) + E_{\mathrm{NMAC}}(T, q, \ell(L))\}, \tag{45}$$

$$\varepsilon_{\mathrm{Hcol}}(q) \le \min\{1,\ \Phi_{256}(3, q+3) + E_{F_1}(1, q+2)\}. \tag{46}$$

Proof For the row, only $I(K)$ and $O(K)$ contain the hidden key. The same disjoint-encoding replacement uses two private calls and at most $q + \ell(L)T$ ordinary calls. The resulting independent-key NMAC receives all T adaptive messages, and its replacement is precisely the private row function. No conversion to a one-transcript game is made.

For one distinct zero-salt column, omit the NMAC step from Lemma C.1, put $C = 1$, and use its conservative $C + 2 = 3$ private-call budget. The only visited $\Gamma(0)$ entry gives one uniform reply independent of h and of the executed prefix. The extra two private-call slots need not actually be used; retaining this upper bound agrees with the existing shared-root accounting. □

C.3 Root initialization at the actual entry boundary

Let S be an independent uniform 32-byte contribution. Executed preparation may choose a known 32-byte old root r and the response core using the public oracle, but has no other raw use of S. Consider one original Extract/Expand,

$$P = F(r, S), \qquad R = F(P, I_R \| \mathtt{01}),$$

with the actual 97-byte information string from Section 5.5. Repeats reuse this result. The value P has only this Expand use; S, P have no other output channel. The entire later protocol is generated from R.

First apply the single hidden-message column argument to salt r. Its ideal value is $V(\Gamma(r))$. Even when r depends on the public oracle prefix, this first table value is fresh independent uniform: conditional on each oracle and prepared state, its index is fixed and has not been visited. Presampling a fresh PRK and withholding it until this call gives exactly this joint distribution.

Embed a hidden row key at that PRK and answer the one 98-byte Expand message. Its ideal reply is a fresh uniform root. It can again be presampled and withheld until creation, independent of the common oracle and all earlier state. All later confirmations, signatures, traffic and successor operations are recomputed from this root. The proof therefore supplies the initialization for Theorem 5.10 to the entire interactive continuation.

If B bounds all ordinary compression work other than this derivation and the displayed public salt-state computations, conservative budgets are $q_{\mathrm{ext}} \le B + 5$ and $q_{\mathrm{exp}} \le B + 2$. The former includes the five actual root-Expand compression calls; the latter includes the two public states in $\Gamma(r)$. Since $\ell(98) = 3$, the resulting bit error is

$$\begin{aligned} \Delta_{\mathrm{init}} \le \min\{1,\, & \Phi_{256}(3, B+8) + E_{F_1}(1, B+7) \\ & + \Phi_{256}(2, B+5) + E_{\mathrm{NMAC}}(1, B+2, 3)\}. \end{aligned} \tag{47}$$

All terms count the complete continuation in their respective reductions. The bound retains the published asymptotic constants and their stated parameter regimes.

The downstream length budgets are equally explicit: confirmation information lengths $62, 58, 60$ give row messages $63, 59, 61$; the 155-byte traffic information gives $156, 188, 188, 188$; and the 56-byte first-step information gives $57, 89$. Thus the shared root row has $\ell(32) = 2$, the traffic rows have maximum $\ell(188) = 5$, and the selected chain row has $\ell(89) = 3$. These counts include the previous HKDF output block where it is part of the actual HMAC message.

D Accounting through one common virtual experiment

This appendix records the alternative accounting route when one common real-component virtual experiment and a global native-to-virtual comparison have been established. In that setting the shared comparison cost is charged once; uniform resource bounds follow by direct substitution. Section 5.3 gives the candidate-local route used by the record and report interfaces.

D.1 A selected-target experiment

Let V be the virtual original-record experiment with the real KEM component. The adversary chooses its target after its observed interaction. Eligibility E retains the selected recovery conditions: both original target reservations follow Cut, the actual target signing keys and target-epoch entropy remain unexposed, and no selected-session record saved after Cut is acquired. The experiment follows allowed asynchronous control and application transitions, including sending after the local FinalReady transition.

Let Attr mean that the selected target has a unique usable pair of original slots, and let $I \in J$ denote that pair. The mapping and every event used by the reduction must be computable from the adapter's known semantic state and observations. Let

$$\eta \;\geq\; \Pr_{V}[E \wedge \neg\mathsf{Attr}].$$

For $j \in J$, define V_j as the full experiment with the adapter's random component override at pair j. Its transcript, eligibility and target choice are recomputed in that world. The remaining protocol and oracle calls use the same common interfaces.

The application test bit is sampled by the surrounding simulator. Write $S = +1$ for a correct guess and $S = -1$ otherwise; invalid tests contribute zero through E. Every $E = 0$ execution is completed by an independent fair guess. Thus the usual message-guessing advantage is $\frac{1}{2}|\mathbb{E}[ES]|$. The simulator can evaluate this score without observing the hidden bit of its component KEM challenger.

D.2 Centered-score composition

For each j, put $E_j = E \wedge \mathsf{Attr} \wedge (I = j)$. Suppose the fixed-pair component adapter gives, for $\sigma \in \{+1, -1\}$,

$$\left|\Pr_{V}[E_j \wedge S = \sigma] - \Pr_{V_j}[E_j \wedge S = \sigma]\right| \leq a_j^{\sigma} + \delta_j, \tag{48}$$

where a_j^{σ} is the advantage of the actual derived component distinguisher and δ_j is the corresponding correctness term. The distinguisher includes target selection and all relevant oracle queries. Define

$$b_j = \mathbb{E}_{V_j}[E_j S].$$

Finally, let ϵ_{rec} bound the binary distinguishing gap between the native and virtual experiments for their efficiently evaluable positive-score and negative-score events. This is the concrete record-interface connection, including public outputs and accepted operation histories.

Proposition D.1 (Adaptive origin-pair accounting) *Under these stated interface and fixed-pair bounds, the native application's guessing advantage satisfies*

$$\mathrm{Adv}_{\mathrm{msg}} \leq \min\left\{\frac{1}{2},\; \epsilon_{\mathrm{rec}} + \frac{1}{2}\left(\eta + \sum_{j\in J}|b_j| + \sum_{j\in J}(a_j^{+} + a_j^{-} + 2\delta_j)\right)\right\}. \tag{49}$$

Proof In the single experiment V, the attributed events E_j are disjoint. Decomposing its score yields

$$\mathbb{E}_V[ES] = \mathbb{E}_V[ES\,\mathbf{1}_{\neg\mathsf{Attr}}] + \sum_{j\in J} \mathbb{E}_V[E_j S].$$

The first term has magnitude at most η. Each attributed expectation is the positive-event probability minus the negative-event probability. Applying (48) to those two events and then the triangle inequality bounds its distance from b_j by $a_j^+ + a_j^- + 2\delta_j$. Summing gives the virtual score bound. The two native/virtual binary-event comparisons contribute at most $2\epsilon_{\mathrm{rec}}$ to the score difference. Dividing by two and using the maximum bit-guessing advantage proves (49). □

The selector may depend adaptively on its entire observed transcript. The absolute values in $\sum_j |b_j|$ retain each override world's own target selection. In particular, the override probabilities are not normalized as though they belonged to one common ideal distribution. The score decomposition also retains failed and ineligible executions in the experiment, avoiding an uncontrolled division by the probability of eligibility.