

Proof-to-Byte: Assumption-Minimal, Cross-Substrate Binding Assurance for PQ/T Hybrid KEMs

Anonymous Author(s)

Abstract—Post-quantum hybrid key-encapsulation (PQ/T KEMs, combining a lattice KEM with an elliptic-curve KEM) is being deployed across TLS, Signal, and MLS faster than its *deployment safety* can be assured: the same period has seen side-channel breaks (KyberSlash), binding / re-encapsulation breaks (PQXDH; the ML-KEM malicious-key binding failures of Schmieg), and a widening gap between what a hybrid’s *specification* guarantees and what its compiled, multi-platform *implementation* actually does. We present Q-Periapt, an assurance suite for PQ/T hybrid KEMs organised around one principle: *a security property is only as trustworthy as its weakest realization*. Its combiner, ContextBound, achieves binding (MAL-BIND-K-CT and MAL-BIND-K-PK) reducing to collision-resistance of SHA3 *alone*—with no binding assumption on the component KEMs—and we machine-check this in EasyCrypt. The injective encoding is a proved lemma (not an axiom), and—stated honestly—the result is in substance a *commitment / transcript-binding* theorem (the session key commits to the full transcript; the component decapsulation is inert in the argument, which is precisely why no KEM assumption is needed), instantiated in the Cremers–Dax–Medinger Figure 6 game for both rejection styles. Crucially, the *same* proven combiner is realized byte-identically across five language faces and three operating systems—*executed* on four instruction-set architectures (three native, riscv64 under qemu-user emulation) and cross-compiled, with by-construction identity, on one more—verified by a six-method conformance matrix and a self-validating source→binary constant-time probe (configured as a CI gate; measured locally) that we show *discriminates* a clean primitive (libcrux ML-KEM: zero secret-dependent branches after compilation) from a leaky one (HQC’s PQClean constant-weight sampler). We integrate the combiner into a production TLS 1.3 path (a rustls CryptoProvider) and evaluate handshake tail latency under real tc netem. Beyond attaining the ceiling, we prove a machine-checked *separation map*: in one EasyCrypt model we delete each combiner input in turn and show exactly which binding notion it costs—omitting the PQ public key loses MAL-BIND-K-PK over the FIPS-203 expanded decapsulation-key serialization (a concrete adversary wins with probability one) but not over the seed-dk form; omitting the context loses MAL-BIND-K-CTX unconditionally; whereas omitting the PQ ciphertext is *harmless*. The asymmetry is the contribution: the implicit-rejection shared secret transitively rebinds the *ciphertext* but *not* the *public key* (which it leaves key-independent under expanded-dk seed-substitution). So public-key and context absorption are *necessary* while ciphertext absorption is redundant, and the deployed X-Wing shape (omitting all three) provably loses K-PK and K-CTX while keeping K-CT—a field-resolved account of exactly where ContextBound’s extra binding pays off. We are explicit about what is and is not novel: the binding *fact* is a known consequence of recent results; our contributions are this separation theorem and the *conjunction*—an assumption-minimal, machine-checked binding result whose proven object is demonstrably identical across a heterogeneous deployment

surface, with reproducible gates that catch the failure classes that have broken deployed hybrids.

Index Terms—Post-quantum cryptography, hybrid key encapsulation, binding/committing security, formal verification, constant-time, cross-platform assurance, TLS.

I. INTRODUCTION

THE migration to post-quantum (PQ) key establishment is being done conservatively, through *hybrids*: a PQ KEM (ML-KEM [1]) run alongside a classical KEM (X25519), with the two shared secrets combined so that the session key is secure if *either* component holds. Hybrid KEMs are now shipping in TLS 1.3 (X25519MLKEM768), in Signal’s PQXDH, and in MLS, and are specified in CFRG drafts such as X-Wing [5].

The gap this paper addresses: Hybrids are being deployed faster than their *deployment safety* can be assured, and the failures of the last two years were realization-level, not failures of the underlying hard problems. (i) *Side channels*. KyberSlash [6] showed that secret-dependent timings in widely-used Kyber/ML-KEM implementations leaked the secret key—a property of compiled code that a source-level constant-time argument does not by itself settle. (ii) *Binding*. The PQXDH protocol admitted a re-encapsulation issue, and Schmieg [3] proved that ML-KEM, in the FIPS-203 *expanded* decapsulation-key form many libraries store, is neither MAL-BIND-K-CT nor MAL-BIND-K-PK—so whether a hybrid binds its transcript can depend on a serialization choice made three layers down. (iii) *Spec-to-binary drift*. A hybrid’s guarantee is stated on a specification, sometimes proved on a model or source, but *run* as a compiled artifact—and real deployments compile it through many language bindings, operating systems, and instruction sets, each an opportunity for the proven object and the executed object to diverge. These three are one underlying problem: *a security property holds only of a realization, and a hybrid has many*. This is distinct from—and orthogonal to—*auditability* and *crypto-agility*, which describe *what* crypto is running; the question here is whether the hybrid that actually runs is *safe*.

Approach: We present Q-Periapt, an assurance suite for PQ/T hybrid KEMs built around one principle: a security property is only as trustworthy as its weakest realization. Figure 1 is the organising idea—a single machine-checked binding result whose proven object is realized byte-identically across the whole deployment surface (“proof-to-byte”), guarded by reproducible gates (configured in CI; we report constant-

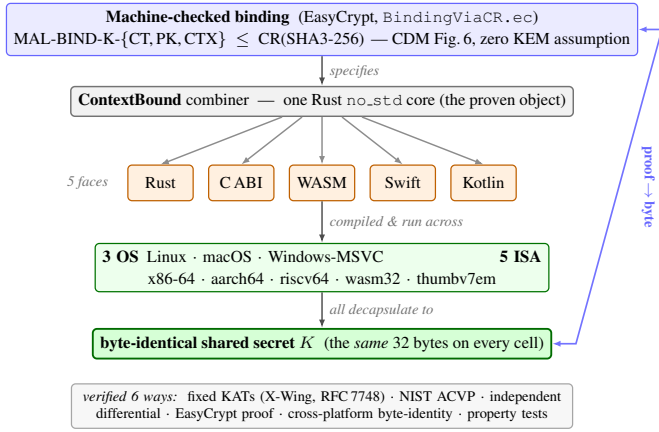


Fig. 1. Proof-to-byte. One machine-checked binding result (top) is realized by a single Rust `no_std` core, exposed through five language faces and run across three operating systems and five instruction-set architectures, all converging on a byte-identical shared secret K (bottom); verified six ways.

time results here as local runs) that target the specific failure classes above.¹

Contributions:

- **C1 (Section III).** An *assumption-minimal, machine-checked commitment / transcript-binding* result: ContextBound’s session key commits to the full transcript, reducing MAL-BIND-K-CT, -K-PK, and a (non-standard) context extension -K-CTX to collision-resistance of SHA3 *alone*, with the injective encoding proved and the Cremers–Dax–Medinger Figure-6 game discharged for both rejection styles in EasyCrypt. We are explicit (Section III) that the cryptographic content is this injective-encoding + CR(SHA3) argument—the component Decaps is *inert*, which is exactly why no KEM assumption is needed—and that C1’s value is its completeness, mechanization, and role as the proven object tied to C2, not proof difficulty. On the same kernel we additionally prove a machine-checked *separation map* (Theorem 1, Table IV): deleting each combiner input in turn, we show that omitting pk_{pq} loses MAL-BIND-K-PK over the expanded-dk serialization (concrete $\Pr=1$ adversary) but not over seed-dk, omitting the context loses MAL-BIND-K-CTX unconditionally, and omitting ct_{pq} is harmless (the shared secret binds it)—so public-key and context absorption are *necessary*, ciphertext absorption redundant, and the X-Wing shape provably loses K-PK+K-CTX while keeping K-CT. This is a stated, proven characterization, not an integration claim.
- **C2 (Section IV).** *Proof-to-byte cross-substrate equivalence*: the same proven combiner is byte-identical across 5 language faces and 3 OSes, *executed* on 4 ISAs (3 native + riscv64 under qemu-user) and cross-compiled

with by-construction identity on 1 more, verified by a six-method conformance matrix.

- **C3 (Section V).** Reproducible assurance *instruments*, configured as CI gates: a self-validating source→binary constant-time probe that *discriminates* clean (ML-KEM) from leaky (HQC) (HQC’s leak is known; the contribution is the discriminating, gated oracle); an executable *regression oracle* for the known dk-format binding coupling; and two independent symbolic provers. We report the constant-time measurements honestly as local runs.
- **C4 (Section VI).** A production TLS 1.3 path (rustls CryptoProvider), evaluated under real `tc netem`.

What is and is not novel: We state the boundary plainly.

Not novel, and attributed: that a “hash-everything” combiner is binding from collision-resistance alone while a lean combiner inherits the omitted component’s binding is the Chempat result [4]; ContextBound *is* that construction, not a new primitive. That ML-KEM’s expanded-dk form breaks binding (and the seed form fixes it) is Schmiege’s [3]. The binding-notion lattice is CDM’s [2]. That a constant-time harness must mark only secret sub-fields is standard practice [6]. *Novel:* first, a *machine-checked shape × format separation* (Theorem 1)—where Schmiege shows a *single* KEM’s binding depends on its dk format, we prove, in one model, the orthogonal statement that the *combiner shape* determines whether that format-dependence survives composition (lean: yes, broken over expanded-dk; full: no, binding over both), a result no prior work states; second, the *conjunction* C1∧C2—a CR-only, machine-checked binding result whose proven object is demonstrably byte-identical across a heterogeneous substrate; third, the self-validating source→binary *discriminator* as a reusable CI artifact. The design invariant that binding the ciphertext and public key in the KDF makes hybrid binding robust to the component KEM’s key-serialization format is thus no longer merely operationalized but *proven*. We do **not** claim stronger binding than X-Wing (Section III), nor any live exploited deployment. Table I makes the conjunction concrete: the delta is that no prior effort ties a mechanized hybrid-combiner binding theorem to a byte-identical multi-substrate realization with source→binary and TLS gates—and we are equally explicit that an axis we do *not* hold (a verified specification-to-binary implementation) is held by formosa-mlkem.

II. BACKGROUND AND THREAT MODEL

A. PQ/T hybrid KEMs and combiners

ML-KEM [1] is the standardized form of CRYSTALS-Kyber [11]; like most lattice KEMs it applies a Fujisaki–Okamoto transform [13], [14] with *implicit rejection* (a malformed ciphertext yields a pseudorandom key rather than $\perp$). The companion signature standards are ML-DSA [9] and SLH-DSA [10], and the broader process is surveyed in [12]. Hybrid key exchange in TLS 1.3 [16] follows the IETF design [17]—e.g. the X25519MLKEM768 group [18] over X25519 [15]—and deployed PQ/T protocols include Signal’s PQXDH [19] and Apple’s PQ3 [20].

¹To be precise about the link: the EasyCrypt proof is over an abstract injective `encode`, and the shipped Rust `encode` is tied to it by *conformance* (shared ContextBound KATs + the six-method matrix of Section IV), not by a mechanized refinement proof. “Proof-to-byte” thus names a conformance-established link, not a verified compiler; the byte-identity *across substrates* (Section IV) is what is mechanically and exhaustively checked.

TABLE I

POSITIONING. ✓ HOLDS; ~ PARTIAL; BLANK: NOT PROVIDED. THE CONTRIBUTION IS THE CONJUNCTION (LAST COLUMN), NOT ANY SINGLE ROW.

	X-Wing	formosa	SandboxAQ	Chempat	this work
Mechanized binding (combiner)			~		✓
Both rejection styles, $K \neq \perp$				~	✓
Zero KEM-binding assumption			✓		✓
Byte-identical multi-substrate object					✓
Source→binary CT gate		~			✓
Production TLS 1.3 path	~				✓
Verified spec→binary impl. (we do not)		✓	~		

A hybrid KEM combines component shared secrets ss_{pq} and ss_{tr} into a session key K . The design space ranges from simple *concatenation* KDFs (X25519MLKEM768) to combiners that additionally absorb ciphertexts and public keys. X-Wing [5] uses a fixed “lean” combiner $K = \text{SHA3}(ss_{pq} \parallel ss_{tr} \parallel ct_{tr} \parallel pk_{tr} \parallel \text{label})$, omitting the ML-KEM ciphertext and public key.

B. KEM binding

We use the Cremers–Dax–Medinger (CDM) X-BIND- P - Q framework [2]. A *transcript* is a tuple (pk, ct, K) produced by a (possibly malicious) execution; the two “observable” quantities a binding notion may constrain are the public key PK, the ciphertext CT, and the derived key K . For $P, Q \in \{K, PK, CT\}$, a KEM is X-BIND- P - Q if no adversary in the game below wins:

Game $\text{Bind}_{\mathcal{A}}^{P,Q}$: the adversary $\mathcal{A}$ outputs two transcripts (pk_0, ct_0, K_0) and (pk_1, ct_1, K_1) . $\mathcal{A}$ wins iff $P_0 = P_1 \wedge Q_0 \neq Q_1 \wedge K_0 \neq \perp \wedge K_1 \neq \perp$.

Three adversary classes parameterize how the key material in each transcript is produced, in increasing strength: **HON** (honest KeyGen), **LEAK** (honest keys, secret keys leaked to $\mathcal{A}$), and **MAL** (adversarially generated keys). Thus $\text{MAL} \supseteq \text{LEAK} \supseteq \text{HON}$, and CDM prove monotonicity across the (P, Q) lattice, so MAL-BIND-K-CT and MAL-BIND-K-PK jointly yield the whole K-binds- $\{PK, CT\}$ sub-lattice. For implicitly-rejecting KEMs such as ML-KEM, these two are the achievable ceiling on those axes; the dual X-BIND-CT-* notions (a ciphertext binding the key or public key) are *structurally unachievable*, since decapsulation never returns $\perp$ and a mutated ciphertext always yields *some* key. The binding of implicitly-rejecting KEMs—and how it fails under malicious, expanded-form keys—is studied in detail by Schmieg and others [3], [21].

C. Threat model

Reflecting the deployment-safety framing, we consider three realization-level adversaries, each matched to one assurance instrument.

The malicious-key (MAL) binding adversary (Section III) is the strongest CDM class: it generates *all* key material, including mutually inconsistent or maliciously-serialized keys, and wins if it produces two accepting executions whose hybrid keys collide while a ciphertext, public key, or context differs. This is the adversary behind re-encapsulation attacks (PQXDH) and the ML-KEM expanded-key binding failures; it is the one our combiner’s binding theorem targets, and the one the dk-format demonstration of Section V exercises concretely.

The binary-level constant-time adversary (Section V) observes wall-clock or microarchitectural timing and seeks any secret-dependent conditional branch, memory index, or division in the *compiled* code—not the source. KyberSlash is the canonical instance. Source-level constant-time arguments do not discharge this adversary on their own, which is why we probe the emitted binary.

The downgrade/agility adversary attacks profile and policy negotiation, attempting to force a weaker suite or combiner. We bind a signed, versioned policy and a domain-separating `policy_version` into the transcript (Algorithm 1) and fail closed on an unsatisfiable policy; a full treatment of agility is out of scope here.

We *assume*, rather than re-prove, the confidentiality of the component primitives—IND-CCA2 of ML-KEM and the Diffie–Hellman assumption for X25519—and the collision-resistance of SHA3; our contribution is orthogonal, concerning whether the *realization* preserves the binding and constant-time properties the primitives are chosen for.

III. CONTEXTBOUND AND THE MACHINE-CHECKED KERNEL

A. Construction

ContextBound derives

$$K = \text{SHA3-256}(\text{encode}[\text{LABEL}, \text{suite_id}, \text{policy_version}, ss_{pq}, ss_{tr}, ct_{pq}, pk_{pq}, ct_{tr}, pk_{tr}, \text{context}]),$$

where `encode` concatenates each field under a fixed-width 8-byte big-endian length prefix (Algorithm 1). Two design choices are load-bearing. (i) *Injective, length-prefixed encoding*. A naive concatenation $a \parallel b$ is ambiguous— $(“ab”, “”)$ and $(“a”, “b”)$ collide—which is precisely the structure a binding adversary exploits. Prefixing every field with its fixed-width length makes the encoding self-delimiting and therefore injective (proved, Section III), so distinct transcripts map to distinct byte strings and the binding question reduces cleanly to collision-resistance of the hash. (ii) *Absorb everything*. ContextBound feeds *all* component ciphertexts and public keys (and an application context, and a domain-separating label, suite identifier, and policy version) into the hash. By the combiner-binding theorem of Chempat [4], the binding advantage is then bounded by Adv^{CR} alone, with no residual term requiring any component KEM to be binding—trading a KEM-self-binding assumption for a single, well-studied collision-resistance assumption. The dual profile `CompatXWing` reproduces X-Wing’s lean combiner byte-for-byte for interoperability; ContextBound is the hash-everything profile and is policy-selected when an application

Algorithm 1: CONTEXTBOUND combiner.
CompatXWing omits $ct_{pq}, pk_{pq}, S, v, x$ and uses X-Wing’s label.

Input: label L , suite id S , version v , secrets ss_{pq}, ss_{tr} , ciphertexts ct_{pq}, ct_{tr} , public keys pk_{pq}, pk_{tr} , context x

Output: 32-byte shared secret K
 $lp(f) \leftarrow \text{be8}(|f|) \| f$ // 8-byte big-endian length prefix
 $T \leftarrow [L, S, \text{be4}(v), ss_{pq}, ss_{tr}, ct_{pq}, pk_{pq}, ct_{tr}, pk_{tr}, x]$
 $e \leftarrow lp(T_0) \| lp(T_1) \| \dots \| lp(T_{n-1})$ // injective encode
return SHA3-256(e)

needs context binding or wishes to avoid the serialization-format dependence of Section V. The suite is open source.²

B. The EasyCrypt development

Figure 2 shows the mechanized reduction. The encoding’s injectivity, `encode_inj`, is a *proved lemma*, not an assumed one. The proof has two steps. A field-level lemma, `lp_cat_inj`, shows that a single length-prefixed field is self-delimiting: if $lp(a) \| x = lp(b) \| y$ then $a = b$ and $x = y$, because the two 8-byte prefixes have equal width and (being injective) force $|a| = |b|$, after which list-concatenation cancellation (`eqseq_cat`) splits off equal field bodies and equal remainders. A structural induction over the field list then lifts this to whole encodings: a non-empty encoding has size ≥ 8 and so can never equal the empty one, ruling out boundary-shift collisions, and the inductive step peels one field with `lp_cat_inj`. The whole argument bottoms out at two elementary facts about the length field—that an 8-byte big-endian integer occupies 8 bytes (`be8_size`) and is injective (`be8_inj`). On top of `encode_inj` we discharge (i) a generic transcript-collision reduction `bind_le_cr`: a *byequiv* argument that maps any adversary winning the binding game to one finding an H -collision, since a differing observable forces a differing field list (hence, by injectivity, a differing encoding) while the colliding key forces equal hashes; and (ii) the *KEM-aware game*, in which the MAL adversary supplies the keypairs and K is *derived* via the component Decaps functions and the combiner. We mechanize both the implicit-rejection specialization (`malbind*_le_cr`) and the fully general CDM Figure-6 game with explicit rejection (`malbind*_xrej_le_cr`), where the $K \neq \perp$ side condition is *present and load-bearing*—removing it makes the proof fail. Every theorem reduces only to CR(SHA3); the development compiles under the EasyCrypt checker (we pin the checker and SMT-solver versions in the artifact), and each theorem is verified load-bearing on `encode_inj` by an adversarial negative control.

What this is, precisely. We state the cryptographic content plainly, because a CDM-literate reader should not over-read the label. Since K is *defined* as $H(\text{encode}[\dots])$ and the tuple contains the ciphertexts and public keys *directly*, the proof is in substance a *commitment / transcript-binding* statement: it

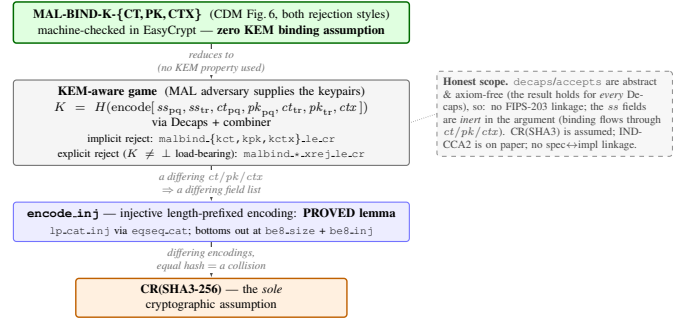


Fig. 2. The kernel as a reduction tower: MAL-BIND-K-{CT,PK,CTX} (CDM Figure 6, both rejection styles) reduces—using no property of Decaps—through the proved `encode_inj` to collision-resistance of SHA3, the sole cryptographic assumption.

shows that one cannot collide K while changing a ciphertext, public key, or context, by reducing to $CR(H)$ over the proved injective encoding. The component Decaps—and hence the adversary’s freedom over key material that CDM binding is fundamentally about—is *inert* in the argument: the theorems would hold for *any* Decaps, which is exactly what “zero KEM assumption” means and why it is achievable. The value of C1 is therefore its *completeness* (both rejection styles, with $K \neq \perp$ discharged), its *mechanization*, and its role as the *proven object* carried, byte-identically, across the substrate of Section IV—not proof difficulty. The concrete attack-resistance arguments (re-encapsulation, the dk-format coupling of Section V) are *design arguments* consistent with this commitment property, not claims the EasyCrypt artifact witnesses at the Decaps level.

C. Honest binding position

Figure 3 states our position precisely. On the standard $\{CT,PK\}$ axes, ContextBound and a correctly-implemented seed-dk X-Wing both reach the MAL-BIND-K-{CT,PK} ceiling; we are *not* “stronger” there. Our edge is (a) *assumption-minimality*: ContextBound’s binding reduces to CR(SHA3) alone, whereas X-Wing’s relies on ML-KEM’s Fujisaki–Okamoto self-binding together with the seed-dk format; (b) *mechanization*; and (c) *provable dk-format robustness* (Theorem 1)—and here the distinction is sharp: while we are not stronger on the $\{CT,PK\}$ axis *at a fixed* serialization, the lean shape’s MAL-BIND-K-PK *depends on* the format (it is broken over expanded-dk), whereas ContextBound’s holds over *both*. That is a genuine, machine-checked robustness separation, not a stronger notion on a shared axis. Honest scope of the mechanization: Decaps is modelled as an abstract function (which is exactly why “zero KEM assumption” is literal), so there is no FIPS-203 linkage and the shared-secret fields are inert in the binding argument; CR(SHA3) is assumed; IND-CCA2 robustness is argued on paper; and there is no specification-to-implementation linkage proof.

IV. PROOF-TO-BYTE CROSS-SUBSTRATE REALIZATION

A. Architecture

The combiner and its encoding live in a single dependency-free Rust `no_std` core; the component primitives are *injected*

²Anonymized artifact repository provided to reviewers via the editor.

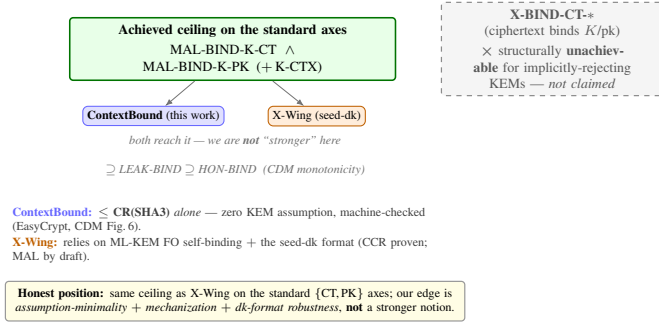


Fig. 3. Honest binding position. Both schemes reach the standard ceiling at a fixed serialization; our edge is assumption-minimality, mechanization, and a machine-checked dk-format robustness separation (Theorem 1)—not a stronger notion on a shared axis. X-BIND-CT* is structurally unachievable for implicitly-rejecting KEMs and is not claimed.

through narrow traits (a `Kem` trait for ML-KEM and X25519, an extendable-output `Xof` trait for the hash), so the proven object—the encoding and combiner—is the same code regardless of which vetted backend (libcrux ML-KEM, x25519-dalek, libcrux SHA3) is plugged in. This trait boundary is also where *crypto-agility* lives: a `Kem`’s `C2PRI` flag (ciphertext second-preimage resistance) gates which combiner profiles it may use, and a signed, versioned *policy* object selects the suite and profile and binds a downgrade-resistant `policy_version` into the transcript (Algorithm 1). Shared secrets are zeroized on drop, and the composition code is written in a branchless, mask-based constant-time style (Section V).

The core is exposed through five language faces—Rust, a C ABI (via `cbindgen`), WebAssembly (via `wasm-bindgen`), Swift (over the C ABI), and Kotlin (via the Java Foreign Function & Memory API)—each a thin shim that marshals bytes in and out of the same core. Because the proven object is one piece of code and every face is a marshalling layer over it, the cross-substrate question reduces to a single, checkable property: does *every* realization produce the same bytes?

B. The six-method conformance matrix

Table II lists the six orthogonal methods (NIST ACVP vectors [31] among them) that answer this. They are deliberately heterogeneous in their oracle and in what they catch, so a defect that escapes one is caught by another. Table III reports the substrate coverage with honest per-cell status: the byte-identical shared secret is *run* natively on x86-64 and aarch64, on a WebAssembly runtime (`wasm32`), and under `qemu-user` emulation on riscv64—there the `ContextBound` and enhanced reference-vector KATs pass on the emulated riscv64 binary, reproducing the pinned bytes exactly. Only the bare-metal `no_std thumbv7em` target is *not* executed; there byte-identity follows by construction from the FIPS-deterministic encodings and identical source. So four of the five faces are *run* and one is cross-compiled—this honest accounting is what makes the cross-substrate claim defensible.

Footprint. The deployable surface is small: the stripped C-ABI shared library—the default hybrid build (libcrux ML-KEM-768, X25519, and the combiner; the optional signature and




Fig. 4. The source→binary CT probe discriminates: the genuine ML-KEM secret yields zero flags; the same harness flags HQC’s constant-weight sampler (193) and validates itself with two controls.

HQC backends are feature-gated and dead-stripped from this build—is 441 KB; the WebAssembly module is ≈ 195 KB via `wasm-pack --release` (further reducible with `wasm-opt -Oz`); and the dependency-free core cross-compiled to the bare-metal `thumbv7em` target, so the same proven combiner runs from a server down to an embedded MCU.

V. CI-GATED ASSURANCE

A. A source→binary constant-time discriminator

libcrux proves its ML-KEM secret-independent at the *source* level (a typed discipline, machine-checked). The orthogonal *binary*-level question—does the compiler reintroduce a secret-dependent branch on the genuine secret path?—is answered by a probe that marks only the genuinely-secret decapsulation-key sub-fields ($\hat{s}$, z) and runs the real libcrux `decapsulate` under Memcheck. The probe is *self-validating*: a planted secret-indexed load is caught (harness sanity), and marking the embedded *public* key flags 5696 reports (Memcheck demonstrably sees real branches in this binary). With only the genuine secret marked, the probe reports **zero** flags on aarch64: source-level secret-independence survives compilation. Run against the suite’s unaudited HQC backend, the same probe flags **193** reports localized to `vect_set_random_fixed_weight`—so the probe is a *discriminator* (Fig. 4): clean ML-KEM versus leaky HQC in one framework. (HQC’s leak is known; the contribution is the discriminating, gated artifact and the corollary that per-backend constant-time gating is necessary, not decorative.) We additionally ran the probe natively on the bare-metal x86-64 host of Table VI: again `probe` = 0 for ML-KEM and 22,849 flags for HQC, confirming the discrimination holds *cross-architecture*. The *absolute* counts differ markedly between ISAs (HQC 193 on aarch64 vs. 22,849 on x86-64; the public-key positive control 5696 vs. 1778)—exactly why the load-bearing signal is the 0-versus-positive *discrimination*, not the raw count.

Why marking granularity matters. A naive harness that marks the *whole* serialized decapsulation key secret produces, on the aarch64 NEON path, 5696 reports across 60 branches comparing 12-bit coefficients to q —which look exactly like a KyberSlash-class leak. They are not: the FIPS-203 `dk embeds` the public key ($dk = dk_pke \parallel ek \parallel H(ek) \parallel z$), and those branches are the compiler’s scalar lowering of the public-key

TABLE II
THE SIX ORTHOGONAL VERIFICATION METHODS.

Method	Reference / oracle	Independence	What it catches
Fixed KATs	X-Wing draft vectors; RFC 7748	published vectors	combiner and X25519 spec-conformance regressions
NIST ACVP	<i>usnistgov</i> ACVP, full FIPS family	NIST ground truth	primitive non-conformance (FIPS 203/204/205)
Multi-backend differential	RustCrypto ml-kem/ml-dsa; orion X25519	independent re-implementations	implementation bugs (output must stay byte-identical)
EasyCrypt proof	machine-checked; CR(SHA3)	formal (computational)	binding-design flaws in the combiner
Cross-platform byte-identity	shared test vector	5 faces $\times$ substrates	per-substrate divergence (same K everywhere)
Generative property tests	proptest: 8 properties $\times$ 128 cases	randomized	edge-case violations (injectivity, domain separation, K-CTX)

TABLE III
CROSS-SUBSTRATE COVERAGE ($\checkmark$ VERIFIED, CI OR LOCAL).

(a) ISA targets (Rust core; the byte-level claim)

ISA	runner	build	byte-id. K	binary-CT
x86-64	Linux, Win-MSVC	$\checkmark$	$\checkmark$ (run)	$\checkmark$ Memcheck
aarch64	Linux, macOS	$\checkmark$	$\checkmark$ (run)	$\checkmark$ Memcheck
wasm32	node	$\checkmark$	$\checkmark$ (run)	source-CT
riscv64	qemu-user	$\checkmark$	$\checkmark$ (run) [‡]	source-CT
thumbv7em	bare-metal	$\checkmark$	by constr. [†]	n/a

[‡] executed under *qemu-user* emulation (reference-vector KATs reproduce pinned bytes); [†] cross-compiled, identity by FIPS-deterministic encodings + identical source.

(b) language face $\times$ OS (shared vector $\rightarrow$ same K)

Face	Linux	macOS	Windows
Rust	$\checkmark$	$\checkmark$	$\checkmark$
C ABI	$\checkmark$	$\checkmark$	$\checkmark$ (MSVC)
WASM (node)	$\checkmark$	$\checkmark$	$\checkmark$
Swift	—	$\checkmark$	—
Kotlin (JVM/FFM)	$\checkmark$	$\checkmark$	—

reduction run during FO re-encryption—operating on *public* bytes. Marking only the genuinely-secret sub-fields ($\hat{s}$ and the implicit-rejection seed z) yields zero reports. This is standard practice [6], but the embedded-public-key pitfall is a real trap; our probe encodes the correct granularity and gates it.

B. A mechanized separation map across the binding sub-lattice

Section III shows ContextBound *attains* the binding ceiling. Here we prove the paper’s distinctive cryptographic contribution: a *machine-checked map* pinning down, field by field, *which inputs the combiner must absorb*—and *why*—across the CDM K-binds- $\{CT, PK, CTX\}$ sub-lattice. All of it lives in *one* EasyCrypt model (a single execution type, a single shared-secret wiring $ss = J(z, ct)$), so every verdict is attributable to combiner *structure*, not to a difference in the modeled KEM. Starting from the full ContextBound field list (which absorbs ct_{pq} , pk_{pq} , and the context) we delete *exactly one* field and ask which binding notion the deletion costs—*three distinct combiners*, not one combiner read under three notions.

Theorem 1 (separation map, machine-checked). *Model $ss = J(z, ct)$ (FIPS-203 §6.3). In one binding model:*

- 1) Omit $pk_{pq} \Rightarrow$ lose K-PK, serialization-conditionally. *Over expanded-dk (z adversary-substitutable) a concrete deterministic adversary wins MAL-BIND-K-PK with probability 1 (*lean_kpk_broken*); over seed-dk ($z=zof(ek)$) binding is restored, $\leq CR(H)$ under injectivity of J and *zof* (*lean_kpk_seed_le_cr*).*
- 2) Omit $ct_{pq} \Rightarrow$ K-CT survives. *The retained shared secret $ss=J(z, ct_{pq})$ transitively binds ct_{pq} , so $Adv^{MAL-BIND-K-CT} \leq CR(H)$ under injectivity of J (*omit_ct_kct_le_cr*).*
- 3) Omit context $\Rightarrow$ lose K-CTX, unconditionally. *A concrete adversary wins MAL-BIND-K-CTX with probability 1, for any dk format (*omit_ctx_kctx_broken*).*
- 4) Full combiner. *Absorbing all three, ContextBound is MAL-BIND-K- $\{CT, PK, CTX\} \leq CR(H)$, with no J/zof assumption (*full_k{pk, ct, ctx}_le_cr*).*
- 5) Joint X-Wing shape. *The deployed X-Wing combiner omits ct_{pq} , pk_{pq} , and context at once; discharged directly, it loses K-PK (*expanded-dk*) and K-CTX but keeps K-CT (*xwing_kpk_broken*, *xwing_kctx_broken*, *xwing_kct_le_cr*).*

The two safety-of-omission verdicts—item 2 and the seed-dk half of item 1—hold under the idealization that the implicit-rejection function J (FIPS-203’s $J=SHAKE(z \parallel H(ct))$) is injective; under a realistic collision-resistant-only J they degrade to an additional CR term over J ’s inner hash rather than holding outright. The breaks (the $Pr=1$ adversaries) and the full combiner (item 4) need no assumption on J .

Reading the map.: The content is the *dichotomy* between rows 2 and 3, not any cell alone. Omitting ct_{pq} is harmless because the shared secret already encodes it; omitting context is fatal because *nothing* else does. So context absorption is necessary *and* sufficient—necessity is admittedly the easy half of an iff (no other field mentions context), sufficiency is *full_kctx_le_cr*; what makes it a result rather than a tautology is its contrast with ct_{pq} , where the shared secret *does* provide the missing binding. Conversely ct_{pq} absorption is *redundant with the shared secret*—but only under the J -injectivity idealization; the full combiner achieves K-CT with no such assumption. Row 1 is the subtle one (Schmieg’s mechanism): one might *hope* $ss = J(z, ct)$ rescues pk_{pq} as it rescues

TABLE IV

THE MACHINE-CHECKED SEPARATION MAP (THEOREM 1): DELETE ONE FIELD FROM CONTEXTBOUND AND ASK WHICH BINDING NOTION SURVIVES. EVERY CELL IS AN EASYCRYPT LEMMA; J IS THE MODELED IMPLICIT-REJECTION FUNCTION.

omitted field	notion at risk	verdict
pk_{pq}	MAL-BIND-K-PK	broken / expanded-dk ($\Pr=1$); safe / seed-dk
ct_{pq}	MAL-BIND-K-CT	safe $\leq CR(H)$ (ss binds ct ; needs J -inj.)
context	MAL-BIND-K-CTX	broken unconditionally ($\Pr=1$)
<i>none</i> (full)	all three	safe $\leq CR(H)$, no J/zof
<i>all three</i> (X-Wing)	—	loses K-PK + K-CTX, keeps K-CT

ct_{pq} , but over expanded-dk the adversary equalizes the stored seed z across two distinct keypairs, so a garbage ciphertext implicit-rejects to the *same* $J(z, ct)$ under $ek_0 \neq ek_1$ [3]; the key is *independent* of pk_{pq} , so the break is genuine—not an artifact of omission (it vanishes the moment ss depends on the key, i.e. under seed-dk). Every verdict is load-bearing under deleted-hypothesis negative controls (the K-PK break needs two distinct keys; the K-CT and K-CTX results need J -injectivity / a distinct context respectively; all reductions need the proved encoding injectivity). **Moral, and the formal design rationale for ContextBound:** pk_{pq} (over expanded-dk) and context absorption are *necessary*; ct_{pq} absorption is *redundant insurance*. This is the multi-level edge the construction needed: not a stronger notion than X-Wing on a shared axis, but a proven, field-resolved account of *exactly* where the X-Wing shape is exposed and ContextBound is not.

Scope, stated precisely.: (i) J, zof, H are abstract; $J(z, ct)$ is *modeled* after FIPS-203 implicit rejection, not *derived* from a mechanized Decaps, and the component shared secrets are otherwise inert. (ii) The K-PK and K-CTX breaks need *no* cryptographic assumption (concrete $\Pr=1$ adversaries); the safety bounds are conditional on $CR(H)$, and the K-CT and seed-dk-K-PK results additionally on J/zof injectivity—a strictly stronger idealization than the full combiner needs. (iii) The K-PK break is *not* a break of X-Wing in deployment, which mandates seed-dk; the map characterizes the *shape*’s exposure field by field. We complement the proofs with an executable regression oracle that exercises the dk-format coupling against real libcrux as a CI-configured gate.

C. Multi-prover assurance

Beyond the computational EasyCrypt proof, the server-authenticated hybrid handshake is modelled in two independent symbolic provers. The Tamarin model captures the ML-KEM and X25519 key exchanges, the combiner as an opaque one-way function, and an ML-DSA server signature, with rules that let the adversary compromise either KEM or reveal the signing key; it proves executability, server authentication, and—the headline lemma—*hybrid robustness*: under an honest server identity the session key remains secret unless *both* KEMs are broken. A ProVerif model establishes the analogous queries under a different abstraction (the classical leg as a second

TABLE V

PRIMITIVE AND HYBRID COST (μs) AND KEY/CIPHERTEXT SIZES (BYTES). ARM64 HOST.

Scheme	time (μs)			size (B)		
	keygen	encaps	decaps	ek	dk	ct
ML-KEM-512	6.6	6.1	6.8	800	1632	768
ML-KEM-768	11.3	11.0	11.8	1184	2400	1088
ML-KEM-1024	18.5	16.5	16.7	1568	3168	1568
X25519	44.1	89.6	45.0	32	32	32
Hybrid, ContextBound	—	108.4	65.4	1216	—	1120
Hybrid, CompatXWing	—	102.2	59.2	1216	—	1120

idealized KEM), giving an independent cross-check. All three formal artifacts are configured as CI gates (EasyCrypt `make check`, Tamarin, ProVerif), with checker and solver versions pinned in the artifact.

VI. PRODUCTION INTEGRATION AND EVALUATION

A. Microbenchmarks

Table V reports primitive and hybrid costs (Criterion, point estimates, on an Apple-Silicon arm64 host; the artifact reproduces them and an x86-64 run). Two findings matter for deployment. First, *the post-quantum leg is not the bottleneck*: ML-KEM-768 encapsulation is 11.0 μs , roughly $8\times$ cheaper than X25519 encapsulation (89.6 μs , an ephemeral keygen plus a Diffie–Hellman), so the classical half dominates hybrid CPU—a useful corrective to the assumption that PQ is the expensive part. Second, *the hash-everything combiner is cheap*: ContextBound costs only +6.2 μs over the byte-exact X-Wing combiner CompatXWing (encapsulation 108.4 vs. 102.2 μs ; decapsulation 65.4 vs. 59.2 μs)—the price of binding the full transcript is sub-10 μs , well under the $\sim 190 \mu s$ X25519 cost and far below any network RTT.

B. Handshake latency under emulated networks

We expose ContextBound (and CompatXWing) as a TLS 1.3 key-exchange group via a rustls [32] `CryptoProvider`, using private-use group codes; a real loopback TLS 1.3 handshake completes over this provider. Figure 5 shows the server-authenticated hybrid flow: the client offers a combined keyshare, the server encapsulates to both components and derives the session key with the combiner, authenticates with an ML-DSA signature, and the client decapsulates and re-derives the same key—which stays secret unless *both* component KEMs are broken (the robustness property our symbolic models establish, Section V). We evaluate handshake *time-to-session* (TCP connect plus full TLS 1.3 handshake) over a real loopback socket shaped by `tc netem`, comparing a classical X25519 baseline against the two PQ/T profiles.

Figure 6 shows the qualitative picture—the curves coincide because latency is RTT-dominated—and Table VI gives the quantitative camera-ready measurement on quiesced bare-metal hardware (an 8-core AMD Ryzen 7 7700, Ubuntu 24.04, kernel 6.8, the benchmark pinned to two isolated cores under the performance governor; medians over 20 repetitions). At




Fig. 5. Server-authenticated PQ/T hybrid TLS 1.3 handshake over the rustls provider. Both peers derive the same session key K with the ContextBound combiner; K remains secret unless *both* component KEMs are broken.

RTT ≈ 0 the per-group medians are tight to $\pm 2\text{--}3\mu\text{s}$ and the ordering is resolved at the median (ContextBound > CompatXWing in 18 of 20 reps): classical X25519 238 μs , the IANA-standard X25519MLKEM768 263 μs , CompatXWing 318 μs , ContextBound 326 μs . The reproducible cost of the PQ/T upgrade is thus $\approx 88\mu\text{s}$ of ML-KEM CPU over the classical baseline, and—critically—**ContextBound costs only +8.1 μs over CompatXWing**, the price of the stronger hash-everything binding, in close agreement with the +6.2 μs combiner cost measured in isolation (Table V). This quiesced measurement supersedes our earlier virtualized-host run, whose $2\times$ run-to-run swings exceeded the few- μs combiner difference; on bare metal it resolves cleanly.

As RTT grows this fixed cost vanishes into the network: at RTT = 20 ms all four groups sit at 41.1 ms within a 90 μs spread, and at RTT = 50 ms within 15 μs of one another—the combiner difference is two to three orders of magnitude below the link latency. The hybrid’s +2.27 KB (one ML-KEM-768 keyshare each way) fits within the existing TLS flights, adds no round-trip, and—being one ML-KEM-768 keyshare each way—is byte-identical in wire budget to the IANA standard. The two hybrids differ in *two* independent ways, which the table separates: the combiner (ContextBound vs. concatenation), worth +8.1 μs (the bound-compat delta, same ML-KEM implementation), and the ML-KEM implementation itself (+62 μs between the standard row, on aws-lc-rs, and the bound row, on libcrux/ring)—an implementation-not-combiner cost we do not optimize and that an aws-lc-rs ContextBound build would erase. The binding upgrade ContextBound provides is the +8.1 μs , not the 62.

Under jitter and loss. Adding 3 ms of jitter leaves the picture unchanged (all three groups ≈ 41 ms median, ≈ 50 ms p99.9). Under 1–3% packet loss the median is likewise unaffected—loss is rare relative to a handshake—but the *tail* jumps to ~ 1 s, dominated by a TCP retransmission timeout per lost packet rather than by crypto or wire size; this RTO-driven tail falls on the classical and PQ/T handshakes alike, and at our sample size the loss-event noise precludes a precise per-group tail

TABLE VI
CAMERA-READY TIME-TO-SESSION ON QUIESCED BARE-METAL
HARDWARE (AMD RYZEN 7 7700, BENCHMARK PINNED TO TWO ISOLATED
CORES, PERFORMANCE GOVERNOR; MEDIAN P50 OVER 20 REPETITIONS).
AT RTT 0 THE PER-GROUP ORDER RESOLVES AT THE MEDIAN AND
CONTEXTBOUND IS +8.1 μs OVER COMPATXWING; AT RTT ≥ 20 MS THE
COMBINER COST IS DWARFED BY THE LINK. RAW TRANSCRIPT IN THE
ARTIFACT.

key-exchange group	RTT 0	RTT 20 ms	RTT 50 ms
X25519 (classical)	238.3 μs	41.06 ms	101.16 ms
X25519MLKEM768 (std.)	263.5 μs	41.15 ms	101.17 ms
CompatXWing	317.8 μs	41.14 ms	101.17 ms
ContextBound	325.9 μs	41.13 ms	101.17 ms

comparison. The honest takeaway is that the hybrid’s extra keyshare does not *systematically* worsen loss resilience: the median is byte-insensitive and the tail is governed by transport, not by the combiner.

VII. DISCUSSION

When to use which profile: CompatXWing exists for interoperability: it is byte-exact X-Wing and should be selected when talking to an X-Wing peer. ContextBound is the default for first-party deployments and is preferable whenever (i) the application has a meaningful context (a protocol/role/version label, a channel binding) that the session key should commit to; or (ii) the deployment cannot guarantee that its component KEM will only ever be instantiated in the seed-dk form—in which case binding pk_{pq} and ct_{pq} directly removes the latent serialization dependence of Section V. The cost of that robustness, measured in Section VI, is sub-10 μs and invisible end-to-end.

The assurance philosophy: Our position is deliberately modest about cryptographic novelty and deliberately ambitious about *dependability*: a hybrid is trustworthy only to the extent that its weakest realization is. The recurring failures of the last two years—KyberSlash, PQXDH, the ML-KEM malicious-key binding gaps—were not failures of the underlying hard problems but of *realizations*: a compiled branch, a serialization




Fig. 6. Time-to-session under real `tc netem` on the *virtualized* host (mean of two repetitions); the quiesced bare-metal numbers of Table VI supersede it quantitatively, and this figure is retained for the qualitative shape. (a) latency is RTT-dominated—the three curves coincide. (b) the overhead is the fixed ML-KEM CPU cost at RTT 0 ($\sim 190 \mu\text{s}$ on this virtualized host; $\approx 88 \mu\text{s}$ on bare metal, Table VI); at realistic RTT it is within run-to-run noise (error bars span the two repetitions and cross zero at RTT 20 ms).




Fig. 7. Handshake wire budget. The PQ/T cost is one ML-KEM-768 keyshare each way (+2.27 KB), which fits inside the existing flights.

format, a missing transcript field. The contribution of this paper is a discipline that makes those realization-level properties *checkable and reproducible*: a binding result whose proven object is carried byte-for-byte across the substrate, a constant-time probe that discriminates clean from leaky, and a regression oracle for the key-format coupling. We expect the individual instruments to be reused independently of ContextBound.

Generality: Nothing in the approach is specific to ML-KEM + X25519. The combiner and its proof are generic over the component KEMs; the trait-injected backends already include the full ML-KEM and ML-DSA families and an (unaudited) HQC backend used here as the CT discriminator’s leaky reference. As new PQ primitives are standardized, the same proof-to-byte discipline applies unchanged.

VIII. RELATED WORK

KEM binding and committing security: The systematic study of KEM binding is recent. CDM [2] introduced the X-BIND-*P-Q* lattice, the HON/LEAK/MAL adversary hierarchy, and its monotonicity, and showed where standard KEMs sit; Schmieg [3] then proved that ML-KEM in its FIPS-203 *expanded* decapsulation-key form is neither MAL-BIND-K-CT nor MAL-BIND-K-PK, with seed-form keys as the fix, and [21] analyzes implicitly-rejecting KEMs more broadly. For *combined* KEMs, Chempat [4] establishes the dichotomy we rely on—a hash-everything combiner is binding from collision-resistance alone, while a lean combiner inherits the binding of whatever component it omits; ContextBound is an instance of the former, which is exactly why we attribute the binding *fact* rather than claim it. At the symmetric layer, the same idea appears as context-committing AEAD (CMT) [7], of which our context axis is the KEM-output-layer analogue; the role of public contexts specifically in hybrid KEMs is studied in [8]. Our separation map (Theorem 1) sits one level finer than each of these: where Chempat states a *binary* dichotomy (the lean shape inherits the omitted component’s binding) and Schmieg fixes a *single* KEM’s format-dependence, the map is a *field-resolved necessity* statement—per binding notion and per dk-format, exactly which combiner input is load-bearing and which is made redundant by the shared secret. That the public key and context are necessary while the PQ ciphertext is not is a fact about combiner *structure* that neither prior work isolates, and the deleted-hypothesis negative controls (Section B-A) are what upgrade it from a set of sufficient conditions to a necessity claim. Our contribution is thus twofold: this separation map, and the conjunction of Table I—a mechanized, assumption-minimal binding result whose proven object is byte-identical across the deployment surface.

Side-channel assurance: KyberSlash [6] showed that secret-dependent timings in widely-used Kyber/ML-KEM implementations were exploitable, underscoring that source-level constant-time arguments do not by themselves settle the question for shipped binaries. The ductect/ctgrind discipline of marking secrets and checking for secret-dependent control flow [28], [30], realized over Valgrind/Memcheck [29], is the basis of our binary-CT gate. We build on, rather than re-do, the source-level guarantees of HACL* [22] and the libcrux ML-KEM verification [23]: our probe *corroborates* them dynamically and, crucially, *discriminates* a verified primitive from an unverified one in the same framework.

Hybrids in deployment: Hybrid key exchange is being standardized and shipped: the IETF hybrid design [17] and the X25519MLKEM768 group [18] for TLS 1.3 [16], [15], the X-Wing KEM [5], Signal’s PQXDH (formally analyzed in [19]), and Apple’s PQ3 [20]. These are the baselines we position against; our combiner interoperates with X-Wing (the CompatXWing profile) and our rustls integration targets the same TLS 1.3 deployment path.

Formal verification of PQ cryptography: formosamlkem [24] delivers a verified ML-KEM *implementation* (Jasmin/EasyCrypt) with IND-CCA and constant-time guarantees—an axis (specification-to-binary) we explicitly do not hold

(Table I). The SandboxAQ EasyCrypt-KEMs library mechanizes LEAK-BIND-K-PK for ML-KEM. Our EasyCrypt [25] development is complementary: it targets the *hybrid combiner’s* binding (the commitment-style MAL Figure-6 instantiation), and is cross-checked by independent symbolic models in Tamarin [26] and ProVerif [27], all CI-gated.

IX. LIMITATIONS

The mechanized binding result is over an *abstract* Decaps: it is exactly the “zero KEM assumption” that makes the result general, but it carries no FIPS-203 linkage, the shared-secret fields are inert in the argument, and CR(SHA3) and IND-CCA2 remain assumptions argued respectively in the model and on paper; there is no specification-to-implementation linkage proof. Binary-level constant-time tooling is mature only on x86-64 and aarch64; riscv64 and wasm32 are covered at source level plus upstream attestation. ACVP byte-identity is conformance, not CMVP certification. The suite is research-grade and has not had an independent third-party audit. The primary netem latency numbers (Table VI) are from a quiesced bare-metal host; the earlier virtualized-host run (Fig. 6) and the jitter/loss study (Table IX) are retained as supporting context and labeled as such.

X. CONCLUSION

A PQ/T hybrid is only as safe as its weakest realization. We closed that gap with an assumption-minimal, machine-checked binding result whose proven object is byte-identical across a heterogeneous deployment surface, guarded by reproducible gates (configured in CI)—a source→binary constant-time discriminator, a binding key-format coupling demonstration, and multi-prover symbolic models—and validated on a production TLS 1.3 path. The genuinely new element is the conjunction, not any single fact; we have been explicit about the boundary throughout.

REFERENCES

- [1] National Institute of Standards and Technology, “FIPS 203: Module-Lattice-Based Key-Encapsulation Mechanism Standard,” 2024.
- [2] C. Cremers, A. Dax, and N. Medinger, “Keeping Up with the KEMs: Binding Security of Key Encapsulation Mechanisms,” *ACM CCS*, 2024. (ePrint 2023/1933.)
- [3] S. Schmieg, “Unbindable Kemmy Schmidt: ML-KEM is neither MAL-BIND-K-CT nor MAL-BIND-K-PK,” *IACR ePrint* 2024/523, 2024.
- [4] J. Güneysu, K. Hövelmanns, *et al.*, “Binding Security of Combined KEMs / Chempat,” *IACR ePrint* 2025/1416, 2025.
- [5] D. Connolly, P. Schwabe, and B. Westerbaan, “X-Wing: The Hybrid KEM You’ve Been Looking For,” *draft-connolly-cfrg-xwing-kem* (CFRG), 2024.
- [6] D. J. Bernstein, K. Bhargavan, *et al.*, “KyberSlash: Exploiting Secret-Dependent Division Timings in Kyber Implementations,” *IACR TCHES*, 2025.
- [7] M. Bellare and V. T. Hoang, “Efficient Schemes for Committing Authenticated Encryption,” *EUROCRYPT*, 2022.
- [8] “On the Role of Public Contexts in Hybrid KEMs,” *IACR ePrint* 2026/140.
- [9] National Institute of Standards and Technology, “FIPS 204: Module-Lattice-Based Digital Signature Standard,” 2024.
- [10] National Institute of Standards and Technology, “FIPS 205: Stateless Hash-Based Digital Signature Standard,” 2024.
- [11] J. Bos *et al.*, “CRYSTALS-Kyber: A CCA-Secure Module-Lattice-Based KEM,” *IEEE EuroS&P*, 2018.

- [12] G. Alagic *et al.*, “Status Report on the Fourth Round of the NIST Post-Quantum Cryptography Standardization Process,” NIST IR 8545, 2025.
- [13] E. Fujisaki and T. Okamoto, “Secure Integration of Asymmetric and Symmetric Encryption Schemes,” *CRYPTO*, 1999.
- [14] D. Hofheinz, K. Hövelmanns, and E. Kiltz, “A Modular Analysis of the Fujisaki–Okamoto Transformation,” *TCC*, 2017.
- [15] A. Langley, M. Hamburg, and S. Turner, “Elliptic Curves for Security,” RFC 7748, IETF, 2016.
- [16] E. Rescorla, “The Transport Layer Security (TLS) Protocol Version 1.3,” RFC 8446, IETF, 2018.
- [17] D. Stebila, S. Fluhrer, and S. Gueron, “Hybrid Key Exchange in TLS 1.3,” *draft-ietf-tls-hybrid-design*, IETF, 2023.
- [18] K. Kwiatkowski *et al.*, “Post-Quantum Key Exchange X25519MLKEM768 for TLS 1.3,” *draft-kwiatkowski-tls-ecdhe-mlkem*, IETF, 2024.
- [19] K. Bhargavan, C. Jacomme, F. Kiefer, and R. Schmidt, “Formal Verification of the PQXDH Post-Quantum Key Agreement Protocol,” *IEEE S&P*, 2024.
- [20] Apple Security Engineering, “iMessage with PQ3: The New State of the Art in Quantum-Secure Messaging,” Apple Security Research, 2024.
- [21] K. Hövelmanns *et al.*, “Binding Security of Implicitly-Rejecting KEMs and Application to BIKE and HQC,” *IACR ePrint* 2024/1233, 2024.
- [22] J.-K. Zinzindohoué, K. Bhargavan, J. Protzenko, and B. Beurdouche, “HACL*: A Verified Modern Cryptographic Library,” *ACM CCS*, 2017.
- [23] Cryspen, “Verified ML-KEM (Kyber) in libcrux,” 2024.
- [24] J. B. Almeida *et al.*, “Formally Verifying Kyber: Episode V,” *CRYPTO*, 2024.
- [25] G. Barthe, B. Grégoire, S. Heraud, and S. Z. Béguelin, “Computer-Aided Security Proofs for the Working Cryptographer,” *CRYPTO*, 2011.
- [26] S. Meier, B. Schmidt, C. Cremers, and D. Basin, “The TAMARIN Prover for the Symbolic Analysis of Security Protocols,” *CAV*, 2013.
- [27] B. Blanchet, “Modeling and Verifying Security Protocols with the Applied Pi Calculus and ProVerif,” *Found. Trends Priv. Secur.*, 2016.
- [28] O. Reparaz, J. Balasch, and I. Verbauwhede, “Dude, is My Code Constant Time?” *DATE*, 2017.
- [29] N. Nethercote and J. Seward, “Valgrind: A Framework for Heavyweight Dynamic Binary Instrumentation,” *ACM PLDI*, 2007.
- [30] A. Langley, “cigrind: Checking that Functions are Constant Time with Valgrind,” 2010.
- [31] National Institute of Standards and Technology, “Automated Cryptographic Validation Protocol (ACVP),” <https://github.com/usnistgov/ACVP>, 2024.
- [32] J. Birr-Pixton *et al.*, “rustls: A Modern TLS Library in Rust,” <https://github.com/rustls/rustls>, 2024.

APPENDIX A ARTIFACT AVAILABILITY

The complete suite—Rust core and backends, the five language faces, the EasyCrypt/Tamarin/ProVerif developments, the constant-time and netem harnesses, and the figure-generation scripts—is open source and provided to reviewers as an anonymized repository via the editor (public URL withheld for double-blind review). All numbers in this paper are reproducible from that repository; the bare-metal netem protocol is provided as a script.

APPENDIX B THE MACHINE-CHECKED KERNEL IN DETAIL

Table VII inventories the EasyCrypt development. Of the named results, all are fully discharged (no admit/conjecture); the only assumptions are the two elementary `be8_*` facts about the 8-byte length field and the modeling of H as collision-resistant. We validate that the proof is non-vacuous by an *adversarial negative control*: for each binding theorem we mechanically delete a single hint from its closing tactic and re-run the checker, confirming the proof then *fails* (cannot prove goal). Deleting `encode_inj` breaks

TABLE VII
EASYCRYPT LEMMA INVENTORY (BINDINGVIACR.EC). ALL PROVED;
ASSUMPTIONS: be8_* , $\text{CR}(H)$.

Lemma	Statement (informal)
<code>lp_cat_inj</code>	length-prefixed fields are self-delimiting
<code>encode_inj</code>	the field encoding is injective
<code>bind_le_cr</code>	generic transcript-collision $\leq \text{CR}(H)$
<code>bind_le_cr_k*</code>	instantiated projections (CT/PK/CTX)
<code>malbind_k*_le_cr</code>	KEM-aware, implicit rejection $\leq \text{CR}(H)$
<code>malbind_*_xrej_le_cr</code>	full Fig. 6, explicit rejection, $K \neq \perp$
<i>Separation map (Theorem 1): single-field deletions</i>	
<code>lean_kpk_broken</code>	omit pk_{pq} , expanded-dk: K-PK adversary $\text{Pr} = 1$
<code>lean_kpk_seed_le_cr</code>	omit pk_{pq} , seed-dk: K-PK $\leq \text{CR}(H)$ ($J/\text{seed inj.}$)
<code>omit_ct_kct_le_cr</code>	omit ct_{pq} : K-CT $\leq \text{CR}(H)$ (ss binds ct , $J\text{-inj.}$)
<code>omit_ctx_kctx_broken</code>	omit context: K-CTX adversary $\text{Pr} = 1$ (any dk)
<code>full_k*_le_cr</code>	full shape: $K\text{-}\{\text{PK}, \text{CT}, \text{CTX}\}$ $\leq \text{CR}(H)$, no J/zof
<code>xwing_k*_broken</code>	X-Wing (omit all 3): loses K-PK, K-CTX, $\text{Pr} = 1$
<code>xwing_kct_le_cr</code>	X-Wing: keeps K-CT $\leq \text{CR}(H)$

every binding theorem; deleting the disagreement lemma (`ct_neq_fields_neq`, etc.) breaks the corresponding axis; and—in the explicit-rejection game—deleting the $K \neq \perp$ conjunct from the predicate makes the theorem unprovable, since two mutually-rejecting executions would otherwise “win” without inducing a hash collision. This last control is what distinguishes a faithful Figure-6 instantiation from a vacuous one.

A. Proof structure of the separation map

All parts of Theorem 1 are discharged by two idiomatic skeletons over a single `lexec` execution type, so the map is genuinely *one* model rather than a family of bespoke arguments. *Safety* bounds (`full_*`, `omit_ct_kct_le_cr`, `lean_kpk_seed_le_cr`, `xwing_kct_le_cr`) are *byequiv* reductions to the collision game: a binding winner—two executions whose combiner keys collide while the observable differs—is mapped to the pair of *encodings* of their field lists, which are distinct yet hash-equal, i.e. an H -collision. The list-difference step is the only non-boilerplate part, and it is exactly where each result’s assumption enters. When the omitted observable still sits in the field list (the `full_*` rows), plain injectivity of the encoding suffices. When it does *not*—the crux of `omit_ct_kct_le_cr`, where ct_{pq} is absent—a differing ciphertext must instead force a differing *shared secret* $J(z, ct_{pq})$, which it does *only* under injectivity of J (`jrej_inj`). That single implication is the formal content of “the shared secret transitively binds the ciphertext,” and isolating it is what keeps the K-CT result honest rather than circular. *Breaks* (`lean_kpk_broken`, `omit_ctx_kctx_broken`,

`xwing_k{pk, ctx}_broken`) are *byphoare* arguments exhibiting a concrete deterministic adversary that outputs two executions agreeing on every *absorbed* field but differing on the omitted observable; the probability-one win is then a structural computation whose sole hypothesis is that two distinct observables exist (`ek_neq`, `lctx_neq`).

Non-vacuity by deleted-hypothesis negative controls. A passing machine proof certifies only that the conclusion follows from the stated hypotheses; it does *not* certify that the hypotheses are needed, and a vacuous or over-strong premise is the classic way a mechanized result misleads. For every cell of the map we therefore re-ran the checker with its load-bearing hypothesis deleted from the `smt` invocation and confirmed the proof *fails* (reproducible via `negative-controls.sh`, whose log is in the artifact): the K-PK break collapses without two distinct keys, the K-CTX break without two distinct contexts, the K-CT safety without J -injectivity, the seed-dk K-PK safety without J/zof injectivity, and every reduction without the proved `encode_inj`. The negative controls are what turn the map from a list of *sufficient* conditions into a *necessity* statement (“ContextBound must absorb pk_{pq} and context”), and they are the same methodology that establishes the explicit $K \neq \perp$ conjunct as load-bearing in Table VII. The whole development is admit-free and rests on six elementary axioms (`be8_size`, `be8_inj`, `ek_neq`, `lctx_neq`, `zof_inj`, `jrej_inj`); H ’s collision resistance is the sole *cryptographic* assumption, and the injective encoding is a proved lemma, not an axiom.

Honest assumption ladder. The map’s verdicts do not all sit at the same assumption level, and the differences are themselves informative. The two breaks and the full-combiner safety bounds need *no* idealization of J : the breaks are unconditional concrete adversaries, and the full combiner reduces to $\text{CR}(H)$ alone because it absorbs every field directly. The two *conditional* verdicts—K-CT safety after omitting ct_{pq} , and seed-dk K-PK safety—are precisely the ones that lean on the shared secret to do binding work the encoding no longer does, and they pay for it with the J/zof injectivity idealizations. So the assumption a reader must grant scales with how much the combiner offloads onto the component KEM’s internals: zero for ContextBound’s direct absorption, more for every shortcut. That monotonicity is the quantitative form of the paper’s thesis that absorbing the full transcript is the assumption-minimal choice.

APPENDIX C EXTENDED EVALUATION

Table VIII gives the microbenchmarks of Table V with Criterion’s 95% confidence intervals (arm64 host). Table IX gives the jitter/loss data of Section VI: the median is loss-insensitive, while the tail is governed by a per-loss TCP retransmission timeout and is, at 150 samples, statistically indistinguishable across the three groups—hence we report it as “RTO-dominated, transport-bound” rather than as a per-group claim. Footprint: the stripped C-ABI library is 441 KB (default hybrid build), the WebAssembly module ≈ 195 KB.

TABLE VIII
MICROBENCHMARKS, μs [95% CI]. ARM64 HOST, CRITERION.

Scheme	keygen	encaps	decaps
ML-KEM-512	6.6 [6.6,6.7]	6.1 [5.9,6.5]	6.8 [6.8,6.9]
ML-KEM-768	11.3 [11.2,11.4]	11.0 [10.9,11.0]	11.8 [11.7,11.8]
ML-KEM-1024	18.5 [18.1,19.2]	16.5 [15.8,17.7]	16.7 [16.5,16.8]
X25519	44.1 [43.8,44.4]	89.6 [89.2,90.0]	45.0 [44.7,45.3]
Hybrid, ContextBound	—	108.4 [107.9,108.9]	65.4 [65.1,65.7]
Hybrid, CompatXWing	—	102.2 [101.4,103.0]	59.2

TABLE IX
TIME-TO-SESSION UNDER JITTER/LOSS, P50/P99 (μs), ON THE VIRTUALIZED HOST (SUPPORTING CONTEXT; THE PRIMARY PER-GROUP LATENCY IS THE BARE-METAL TABLE VI). THE P99 UNDER LOSS IS RTO-DOMINATED AND STATISTICALLY INDISTINGUISHABLE ACROSS GROUPS.

netem (RTT 20 ms +)	X25519	ContextBound	CompatXWing
+3 ms jitter	41.5k / 49.9k	41.6k / 50.1k	41.2k / 49.5k
+1% loss	41.1k / $\sim 1.08\text{M}$	41.3k / (noisy)	41.4k / $\sim 1.09\text{M}$
+3% loss	41.2k / $\sim 1.09\text{M}$	41.6k / $\sim 1.09\text{M}$	41.5k / $\sim 1.10\text{M}$

APPENDIX D REPRODUCTION

Every result is reproducible from the artifact. The EasyCrypt kernel: `make -C formal/easycrypt check`. The constant-time discriminator: `sh ctstats/scripts/ct-gap-probe.sh` (Linux+Valgrind; ML-KEM and HQC). The binding/key-format demonstration: `cargo test -p q-periapt-backends --test binding_keyformat_separation`. Microbenchmarks: `cargo bench -p q-periapt-backends`. The TLS netem evaluation: `paper/netem-camera-ready.sh` (bare metal). Cross-substrate byte-identity: the per-face vector tests; the symbolic models: `make under formal/{tamarin,proverif}`. The figures of this paper rebuild via `make -C paper/figures`.

ACKNOWLEDGMENT

The author thanks the maintainers of libcrux/HACL*, EasyCrypt, Tamarin, and ProVerif, whose tools this work builds upon. *[Funding/advisor acknowledgments to be added.]*