Proof-to-Byte: Assumption-Minimal, Cross-Substrate Binding Assurance for PQ/T Hybrid KEMs

Ziang Li

***Abstract*—Post-quantum hybrid key-encapsulation (PQ/T KEMs, combining a lattice KEM with an elliptic-curve KEM) is being deployed across TLS, Signal, and MLS faster than its *deployment safety* can be assured: the same period has seen side-channel breaks (KyberSlash), binding / re-encapsulation breaks (PQXDH; the ML-KEM malicious-key binding failures of Schmieg), and a widening gap between what a hybrid's *specification* guarantees and what its compiled, multi-platform *implementation* actually does. We present Q-Periapt, an assurance suite for PQ/T hybrid KEMs organised around one principle: *a security property is only as trustworthy as its weakest realization*. Its combiner, ContextBound, achieves binding (MAL-BIND-K-CT and MAL-BIND-K-PK) reducing to collision-resistance of SHA3 *alone*—with no binding assumption on the component KEMs—and we machine-check this in EasyCrypt. The injective encoding is a proved lemma (not an axiom), and—stated honestly—the result is in substance a *commitment / transcript-binding* theorem (the session key commits to the full transcript; the component decapsulation is inert in the argument, which is precisely why no KEM assumption is needed), instantiated in the Cremers–Dax–Medinger Figure 6 game for both rejection styles. Crucially, the *same* proven combiner is realized byte-identically across five language faces and three operating systems—executed natively on three instruction-set architectures and cross-compiled, with by-construction identity, on two more—verified by a six-method conformance matrix and a self-validating source→binary constant-time probe (configured as a CI gate; measured locally) that we show *discriminates* a clean primitive (libcrux ML-KEM: zero secret-dependent branches after compilation) from a leaky one (HQC's PQClean constant-weight sampler). We integrate the combiner into a production TLS 1.3 path (a rustls `CryptoProvider`) and evaluate handshake tail latency under real `tc netem`. We are explicit about what is and is not novel: the binding *fact* is a known consequence of recent results; our contribution is the *conjunction*—an assumption-minimal, machine-checked binding result whose proven object is demonstrably identical across a heterogeneous deployment surface, with reproducible gates that catch the failure classes that have broken deployed hybrids.**

***Index Terms*—Post-quantum cryptography, hybrid key encapsulation, binding/committing security, formal verification, constant-time, cross-platform assurance, TLS.**

I. Introduction

THE migration to post-quantum (PQ) key establishment is being done conservatively, through *hybrids*: a PQ KEM (ML-KEM [1]) run alongside a classical KEM (X25519), with the two shared secrets combined so that the session key is secure if *either* component holds. Hybrid KEMs are now shipping in TLS 1.3 (`X25519MLKEM768`), in Signal's PQXDH, and in MLS, and are specified in CFRG drafts such as X-Wing [5].

Author affiliation omitted for review. Artifact (open source): https://github.com/billlza/q-periapt.

The gap this paper addresses: Hybrids are being deployed faster than their *deployment safety* can be assured, and the failures of the last two years were realization-level, not failures of the underlying hard problems. *(i) Side channels.* KyberSlash [6] showed that secret-dependent timings in widely-used Kyber/ML-KEM implementations leaked the secret key—a property of compiled code that a source-level constant-time argument does not by itself settle. *(ii) Binding.* The PQXDH protocol admitted a re-encapsulation issue, and Schmieg [3] proved that ML-KEM, in the FIPS-203 *expanded* decapsulation-key form many libraries store, is neither MAL-BIND-K-CT nor MAL-BIND-K-PK—so whether a hybrid binds its transcript can depend on a serialization choice made three layers down. *(iii) Spec-to-binary drift.* A hybrid's guarantee is stated on a specification, sometimes proved on a model or source, but *run* as a compiled artifact—and real deployments compile it through many language bindings, operating systems, and instruction sets, each an opportunity for the proven object and the executed object to diverge. These three are one underlying problem: *a security property holds only of a realization, and a hybrid has many*. This is distinct from—and orthogonal to—*auditability* and *crypto-agility*, which describe *what* crypto is running; the question here is whether the hybrid that actually runs is *safe*.

Approach: We present Q-Periapt, an assurance suite for PQ/T hybrid KEMs built around one principle: a security property is only as trustworthy as its weakest realization. Figure 1 is the organising idea—a single machine-checked binding result whose proven object is realized byte-identically across the whole deployment surface ("proof-to-byte"), guarded by reproducible gates (configured in CI; we report constant-time results here as local runs) that target the specific failure classes above.

Contributions:

- **C1 (Section III).** An *assumption-minimal, machine-checked commitment / transcript-binding* result: ContextBound's session key commits to the full transcript, reducing MAL-BIND-K-CT, -K-PK, and a (non-standard) context extension -K-CTX to collision-resistance of SHA3 *alone*, with the injective encoding proved and the Cremers–Dax–Medinger Figure-6 game discharged for both rejection styles in EasyCrypt. We are explicit (Section III) that the cryptographic content is this injective-encoding + CR(SHA3) argument—the component Decaps is *inert*, which is exactly why no KEM assumption is needed—and that C1's value is its completeness, mechanization, and role as the proven object tied to C2, not proof difficulty.

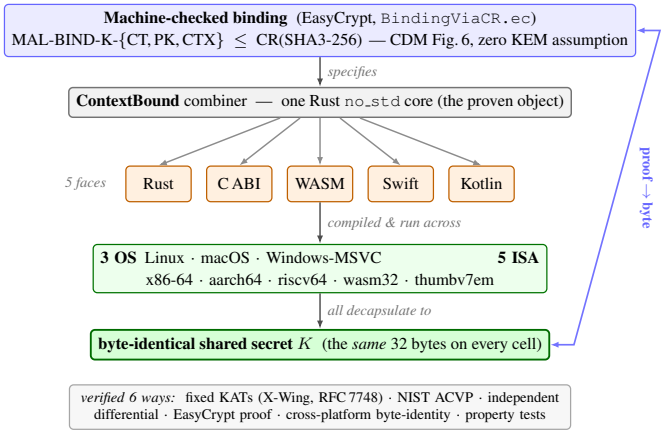

Fig. 1. Proof-to-byte. One machine-checked binding result (top) is realized by a single Rust `no_std` core, exposed through five language faces and run across three operating systems and five instruction-set architectures, all converging on a byte-identical shared secret K (bottom); verified six ways.

- **C2 (Section IV).** *Proof-to-byte cross-substrate equivalence*: the same proven combiner is byte-identical across 5 language faces and 3 OSes, executed natively on 3 ISAs and cross-compiled with by-construction identity on 2 more, verified by a six-method conformance matrix.
- **C3 (Section V).** Reproducible assurance *instruments*, configured as CI gates: a self-validating source→binary constant-time probe that *discriminates* clean (ML-KEM) from leaky (HQC) (HQC's leak is known; the contribution is the discriminating, gated oracle); an executable *regression oracle* for the known dk-format binding coupling; and two independent symbolic provers. We report the constant-time measurements honestly as local runs.
- **C4 (Section VI).** A production TLS 1.3 path (rustls `CryptoProvider`), evaluated under real `tc netem`.

What is and is not novel: We state the boundary plainly. *Not novel, and attributed:* that a "hash-everything" combiner is binding from collision-resistance alone while a lean combiner inherits the omitted component's binding is the Chempat result [4]; ContextBound *is* that construction, not a new primitive. That ML-KEM's expanded-dk form breaks binding (and the seed form fixes it) is Schmieg's [3]. The binding-notion lattice is CDM's [2]. That a constant-time harness must mark only secret sub-fields is standard practice [6]. *Novel:* the *conjunction* C1∧C2—a CR-only, machine-checked binding result whose proven object is demonstrably byte-identical across a heterogeneous substrate, which no prior work provides; the self-validating source→binary *discriminator* as a reusable CI artifact; and the operationalized design invariant that binding the ciphertext and public key in the KDF makes hybrid binding robust to the component KEM's key-serialization format. We do **not** claim stronger binding than X-Wing (Section III), nor any live exploited deployment. Table I makes the conjunction concrete: the delta is that no prior effort ties a mechanized hybrid-combiner binding theorem to a byte-identical multi-substrate realization with source→binary and TLS gates—and we are equally explicit that an axis we do *not* hold (a verified specification-to-binary implementation) is held by formosa-mlkem.

TABLE I
POSITIONING. ✓ HOLDS; ∼ PARTIAL; BLANK: NOT PROVIDED. THE CONTRIBUTION IS THE *conjunction* (LAST COLUMN), NOT ANY SINGLE ROW.

	X-Wing	formosa	SandboxAQ	Chempat	**this work**
Mechanized binding (combiner)			∼		✓
Both rejection styles, $K \neq \perp$				∼	✓
Zero KEM-binding assumption				✓	✓
Byte-identical multi-substrate object					✓
Source→binary CT gate		∼			✓
Production TLS 1.3 path	∼				✓
Verified spec→binary impl. *(we do not)*		✓	∼		

II. Background and Threat Model

A. PQ/T hybrid KEMs and combiners

ML-KEM [1] is the standardized form of CRYSTALS-Kyber [11]; like most lattice KEMs it applies a Fujisaki–Okamoto transform [13], [14] with *implicit rejection* (a malformed ciphertext yields a pseudorandom key rather than $\perp$). The companion signature standards are ML-DSA [9] and SLH-DSA [10], and the broader process is surveyed in [12]. Hybrid key exchange in TLS 1.3 [16] follows the IETF design [17]—e.g. the `X25519MLKEM768` group [18] over X25519 [15]—and deployed PQ/T protocols include Signal's PQXDH [19] and Apple's PQ3 [20].

A hybrid KEM combines component shared secrets ss_{pq} and ss_{tr} into a session key K. The design space ranges from simple *concatenation* KDFs (`X25519MLKEM768`) to combiners that additionally absorb ciphertexts and public keys. X-Wing [5] uses a fixed "lean" combiner $K = \mathrm{SHA3}(ss_{\mathrm{pq}} \,\|\, ss_{\mathrm{tr}} \,\|\, ct_{\mathrm{tr}} \,\|\, pk_{\mathrm{tr}} \,\|\, \mathrm{label})$, omitting the ML-KEM ciphertext and public key.

B. KEM binding

We use the Cremers–Dax–Medinger (CDM) X-BIND-P-Q framework [2]. A *transcript* is a tuple (pk, ct, K) produced by a (possibly malicious) execution; the two "observable" quantities a binding notion may constrain are the public key PK, the ciphertext CT, and the derived key K. For $P, Q \in \{K, \mathsf{PK}, \mathsf{CT}\}$, a KEM is X-BIND-$P$-$Q$ if no adversary in the game below wins:

Game $\mathsf{Bind}^{P,Q}_{\mathcal{A}}$: the adversary $\mathcal{A}$ outputs two transcripts (pk_0, ct_0, K_0) and (pk_1, ct_1, K_1). $\mathcal{A}$ *wins* iff $P_0 = P_1 \;\wedge\; Q_0 \neq Q_1 \;\wedge\; K_0 \neq \perp \;\wedge\; K_1 \neq \perp$.

Three adversary classes parameterize how the key material in each transcript is produced, in increasing strength: **HON** (honest KeyGen), **LEAK** (honest keys, secret keys leaked to $\mathcal{A}$), and **MAL** (adversarially generated keys). Thus MAL $\supseteq$ LEAK $\supseteq$ HON, and CDM prove monotonicity across the (P, Q) lattice, so MAL-BIND-K-CT and

MAL-BIND-K-PK jointly yield the whole K-binds-{PK,CT} sub-lattice. For implicitly-rejecting KEMs such as ML-KEM, these two are the achievable ceiling on those axes; the dual X-BIND-CT-$*$ notions (a ciphertext binding the key or public key) are *structurally unachievable*, since decapsulation never returns $\perp$ and a mutated ciphertext always yields *some* key. The binding of implicitly-rejecting KEMs—and how it fails under malicious, expanded-form keys—is studied in detail by Schmieg and others [3], [21].

C. Threat model

Reflecting the deployment-safety framing, we consider three realization-level adversaries, each matched to one assurance instrument.

The malicious-key (MAL) binding adversary (Section III) is the strongest CDM class: it generates *all* key material, including mutually inconsistent or maliciously-serialized keys, and wins if it produces two accepting executions whose hybrid keys collide while a ciphertext, public key, or context differs. This is the adversary behind re-encapsulation attacks (PQXDH) and the ML-KEM expanded-key binding failures; it is the one our combiner's binding theorem targets, and the one the dk-format demonstration of Section V exercises concretely.

The binary-level constant-time adversary (Section V) observes wall-clock or microarchitectural timing and seeks any secret-dependent conditional branch, memory index, or division in the *compiled* code—not the source. KyberSlash is the canonical instance. Source -level constant-time arguments do not discharge this adversary on their own, which is why we probe the emitted binary.

The downgrade/agility adversary attacks profile and policy negotiation, attempting to force a weaker suite or combiner. We bind a signed, versioned policy and a domain-separating `policy_version` into the transcript (Algorithm 1) and fail closed on an unsatisfiable policy; a full treatment of agility is out of scope here.

We *assume*, rather than re-prove, the confidentiality of the component primitives—IND-CCA2 of ML-KEM and the Diffie–Hellman assumption for X25519—and the collision-resistance of SHA3; our contribution is orthogonal, concerning whether the *realization* preserves the binding and constant-time properties the primitives are chosen for.

III. ContextBound and the Machine-Checked Kernel

A. Construction

ContextBound derives $K = \mathrm{SHA3\text{-}256}\big(\mathrm{encode}[\mathsf{LABEL}, \mathsf{suite\_id}, \mathsf{policy\_version}, ss_{\mathrm{pq}}, ss_{\mathrm{tr}}, ct_{\mathrm{pq}}, pk_{\mathrm{pq}}, ct_{\mathrm{tr}}, pk_{\mathrm{tr}}, \mathsf{context}]\big)$, where `encode` concatenates each field under a fixed-width 8-byte big-endian length prefix (Algorithm 1). Two design choices are load-bearing. *(i) Injective, length-prefixed encoding.* A naive concatenation $a \,\|\, b$ is ambiguous—("ab", "") and ("a", "b") collide—which is precisely the structure a binding adversary exploits. Prefixing every field with its fixed-width length makes the encoding self-delimiting and therefore injective (proved, Section III), so distinct transcripts map to distinct byte strings and the binding question reduces cleanly to collision-resistance of the hash. *(ii) Absorb everything.* ContextBound feeds *all* component ciphertexts and public keys (and an application context, and a domain-separating label, suite identifier, and policy version) into the hash. By the combiner-binding theorem of Chempat [4], the binding advantage is then bounded by $\mathrm{Adv}^{\mathrm{CR}}$ alone, with no residual term requiring any component KEM to be binding—trading a KEM-self-binding assumption for a single, well-studied collision-resistance assumption. The dual profile `CompatXWing` reproduces X-Wing's lean combiner byte-for-byte for interoperability; `ContextBound` is the hash-everything profile and is policy-selected when an application needs context binding or wishes to avoid the serialization-format dependence of Section V. The suite is open source.[1]

Algorithm 1: ContextBound combiner. `CompatXWing` omits $ct_{\mathrm{pq}}, pk_{\mathrm{pq}}, S, v, x$ and uses X-Wing's label.

Input: label L, suite id S, version v, secrets $ss_{\mathrm{pq}}, ss_{\mathrm{tr}}$, ciphertexts $ct_{\mathrm{pq}}, ct_{\mathrm{tr}}$, public keys $pk_{\mathrm{pq}}, pk_{\mathrm{tr}}$, context x
Output: 32-byte shared secret K
$\mathsf{lp}(f) \leftarrow \mathsf{be8}(|f|) \,\|\, f$ `// 8-byte big-endian length prefix`
$T \leftarrow [\,L,\ S,\ \mathsf{be4}(v),\ ss_{\mathrm{pq}},\ ss_{\mathrm{tr}},\ ct_{\mathrm{pq}},\ pk_{\mathrm{pq}},\ ct_{\mathrm{tr}},\ pk_{\mathrm{tr}},\ x\,]$
$e \leftarrow \mathsf{lp}(T_0) \,\|\, \mathsf{lp}(T_1) \,\|\, \cdots \,\|\, \mathsf{lp}(T_{n-1})$ `// injective encode`
return SHA3-256(e)

B. The EasyCrypt development

Figure 2 shows the mechanized reduction. The encoding's injectivity, `encode_inj`, is a *proved lemma*, not an assumed one. The proof has two steps. A field-level lemma, `lp_cat_inj`, shows that a single length-prefixed field is self-delimiting: if $\mathsf{lp}(a) \,\|\, x = \mathsf{lp}(b) \,\|\, y$ then $a = b$ and $x = y$, because the two 8-byte prefixes have equal width and (being injective) force $|a| = |b|$, after which list-concatenation cancellation (`eqseq_cat`) splits off equal field bodies and equal remainders. A structural induction over the field list then lifts this to whole encodings: a non-empty encoding has size ≥ 8 and so can never equal the empty one, ruling out boundary-shift collisions, and the inductive step peels one field with `lp_cat_inj`. The whole argument bottoms out at two elementary facts about the length field—that an 8-byte big-endian integer occupies 8 bytes (`be8_size`) and is injective (`be8_inj`). On top of `encode_inj` we discharge (i) a generic transcript-collision reduction `bind_le_cr`: a `byequiv` argument that maps any adversary winning the binding game to one finding an H-collision, since a differing observable forces a differing field list (hence, by injectivity, a differing encoding) while the colliding key forces equal hashes; and (ii) the *KEM-aware* game, in which the MAL adversary supplies the keypairs and K is *derived* via the component Decaps functions and the combiner. We mechanize both the implicit-rejection specialization (`malbind_*_le_cr`) and the

[1] https://github.com/billlza/q-periapt

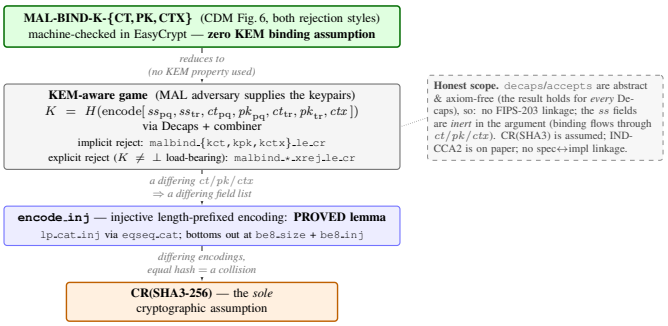

Fig. 2. The kernel as a reduction tower: MAL-BIND-K-{CT,PK,CTX} (CDM Figure 6, both rejection styles) reduces—using no property of Decaps—through the proved `encode_inj` to collision-resistance of SHA3, the sole cryptographic assumption.

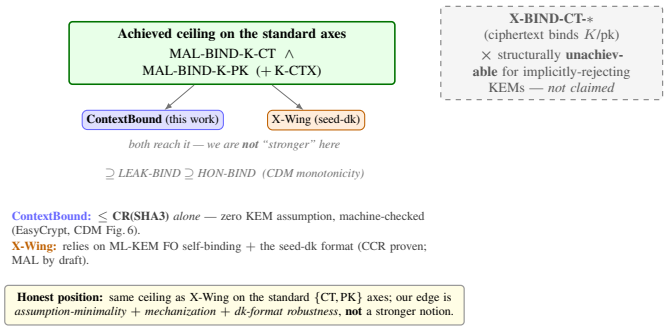

Fig. 3. Honest binding position. Both schemes reach the standard ceiling; our edge is assumption-minimality, mechanization, and dk-format robustness—not a stronger notion. X-BIND-CT-∗ is structurally unachievable for implicitly-rejecting KEMs and is not claimed.

fully general CDM Figure-6 game with explicit rejection (`malbind_*_xrej_le_cr`), where the $K \neq \perp$ side condition is *present and load-bearing*—removing it makes the proof fail. Every theorem reduces only to CR(SHA3); the development compiles under the EasyCrypt checker (we pin the checker and SMT-solver versions in the artifact), and each theorem is verified load-bearing on `encode_inj` by an adversarial negative control.

What this is, precisely. We state the cryptographic content plainly, because a CDM-literate reader should not over-read the label. Since K *is defined* as $H(\text{encode}[\dots])$ and the tuple contains the ciphertexts and public keys *directly*, the proof is in substance a *commitment / transcript-binding* statement: it shows that one cannot collide K while changing a ciphertext, public key, or context, by reducing to $\text{CR}(H)$ over the proved injective encoding. The component Decaps—and hence the adversary's freedom over key material that CDM binding is fundamentally about—is *inert* in the argument: the theorems would hold for *any* Decaps, which is exactly what "zero KEM assumption" means and why it is achievable. The value of C1 is therefore its *completeness* (both rejection styles, with $K \neq \perp$ discharged), its *mechanization*, and its role as the *proven object* carried, byte-identically, across the substrate of Section IV—not proof difficulty. The concrete attack-resistance arguments (re-encapsulation, the dk-format coupling of Section V) are *design arguments* consistent with this commitment property, not claims the EasyCrypt artifact witnesses at the Decaps level.

C. Honest binding position

Figure 3 states our position precisely. On the standard {CT,PK} axes, ContextBound and a correctly-implemented seed-dk X-Wing both reach the MAL-BIND-K-{CT,PK} ceiling; we are *not* "stronger" there. Our edge is (a) *assumption-minimality*: ContextBound's binding reduces to CR(SHA3) alone, whereas X-Wing's relies on ML-KEM's Fujisaki–Okamoto self-binding together with the seed-dk format; (b) *mechanization*; and (c) *dk-format robustness* (Section V). Honest scope of the mechanization: Decaps is modelled as an abstract function (which is exactly why "zero KEM assumption" is literal), so there is no FIPS-203 linkage and the shared-secret fields are inert in the binding argument; CR(SHA3) is assumed; IND-CCA2 robustness is argued on paper; and there is no specification-to-implementation linkage proof.

IV. Proof-to-Byte Cross-Substrate Realization

A. Architecture

The combiner and its encoding live in a single dependency-free Rust `no_std` core; the component primitives are *injected* through narrow traits (a `Kem` trait for ML-KEM and X25519, an extendable-output `Xof` trait for the hash), so the proven object—the encoding and combiner—is the same code regardless of which vetted backend (libcrux ML-KEM, x25519-dalek, libcrux SHA3) is plugged in. This trait boundary is also where *crypto-agility* lives: a `Kem`'s `C2PRI` flag (ciphertext second-preimage resistance) gates which combiner profiles it may use, and a signed, versioned *policy* object selects the suite and profile and binds a downgrade-resistant `policy_version` into the transcript (Algorithm 1). Shared secrets are zeroized on drop, and the composition code is written in a branchless, mask-based constant-time style (Section V).

The core is exposed through five language faces—Rust, a C ABI (via `cbindgen`), WebAssembly (via `wasm-bindgen`), Swift (over the C ABI), and Kotlin (via the Java Foreign Function & Memory API)—each a thin shim that marshals bytes in and out of the same core. Because the proven object is one piece of code and every face is a marshalling layer over it, the cross-substrate question reduces to a single, checkable property: does *every* realization produce the same bytes?

B. The six-method conformance matrix

Table II lists the six orthogonal methods (NIST ACVP vectors [31] among them) that answer this. They are deliberately heterogeneous in their oracle and in what they catch, so a defect that escapes one is caught by another. Table III reports the substrate coverage with honest per-cell status: the byte-identical shared secret is *run* natively on x86-64, aarch64, and wasm32; on riscv64 and the bare-metal `thumbv7em` target the core cross-compiles cleanly and byte-identity follows by construction from the FIPS-deterministic encodings and identical source, but is not executed on a native runner. This honest accounting is what makes the cross-substrate claim defensible.

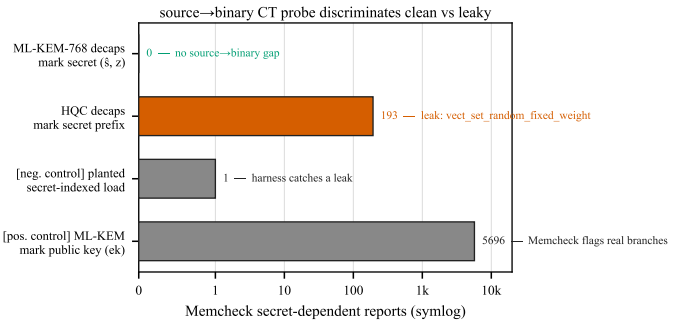

Fig. 4. The source→binary CT probe discriminates: the genuine ML-KEM secret yields zero flags; the same harness flags HQC's constant-weight sampler (193) and validates itself with two controls.

Footprint. The deployable surface is small: the stripped C-ABI shared library—the *entire* suite, including libcrux ML-KEM and ML-DSA—is 441 KB; the WebAssembly module is 249 KB; and the dependency-free core cross-compiles to the bare-metal `thumbv7em` target, so the same proven combiner runs from a server down to an embedded MCU.

V. CI-Gated Assurance

A. A *source→binary constant-time* discriminator

libcrux proves its ML-KEM secret-independent at the *source* level (a typed discipline, machine-checked). The orthogonal *binary*-level question—does the compiler reintroduce a secret-dependent branch on the genuine secret path?—is answered by a probe that marks only the genuinely-secret decapsulation-key sub-fields ($\hat{s}$, z) and runs the real libcrux `decapsulate` under Memcheck. The probe is *self-validating*: a planted secret-indexed load is caught (harness sanity), and marking the embedded *public* key flags 5696 reports (Memcheck demonstrably sees real branches in this binary). With only the genuine secret marked, the probe reports **zero** flags on aarch64: source-level secret-independence survives compilation. Run against the suite's unaudited HQC backend, the same probe flags **193** reports localized to `vect_set_random_fixed_weight`—so the probe is a *discriminator* (Fig. 4): clean ML-KEM versus leaky HQC in one framework. (HQC's leak is known; the contribution is the discriminating, gated artifact and the corollary that per-backend constant-time gating is necessary, not decorative.)

Why marking granularity matters. A naive harness that marks the *whole* serialized decapsulation key secret produces, on the NEON path, 2848 reports across 30 branches comparing 12-bit coefficients to q—which look exactly like a KyberSlash-class leak. They are not: the FIPS-203 `dk` *embeds* the public key ($\texttt{dk} = \texttt{dk\_pke} \,\|\, \texttt{ek} \,\|\, H(\texttt{ek}) \,\|\, z$), and those branches are the compiler's scalar lowering of the public-key reduction run during FO re-encryption—operating on *public* bytes. Marking only the genuinely-secret sub-fields ($\hat{s}$ and the implicit-rejection seed z) yields zero reports. This is standard practice [6], but the embedded-public-key pitfall is a real trap; our probe encodes the correct granularity and gates it.

B. Binding is contingent on the component dk format

We demonstrate, executably against real libcrux, that a lean (X-Wing-shaped) combiner's MAL-BIND-K-PK is *contingent on* the component KEM's key-serialization format. Schmieg's construction exploits that the FIPS-203 *expanded* `dk` stores the implicit-rejection seed z as a substitutable field: an adversary equalizes z across two otherwise-distinct keypairs, so a garbage ciphertext implicit-rejects to the *same* $J(z \,\|\, c)$ under two different public keys. The lean combiner—which omits pk_{pq}—then yields the same hybrid key for $\texttt{ek}_0 \neq \texttt{ek}_1$, breaking MAL-BIND-K-PK; ContextBound, which absorbs pk_{pq}, does not. We stress that this is *not* a break of X-Wing, which mandates the seed-dk format that defeats the attack (KeyGen re-derives z); it is a robustness separation of the combiner *shape*, and the operationalized lesson that binding ct/pk in the KDF removes a latent dependence on the component library's serialization choice—a dependence an implementation that caches expanded keys for performance could silently reintroduce. Our test executes this against real libcrux as a CI-configured regression oracle.

C. Multi-prover assurance

Beyond the computational EasyCrypt proof, the server-authenticated hybrid handshake is modelled in two independent symbolic provers. The Tamarin model captures the ML-KEM and X25519 key exchanges, the combiner as an opaque one-way function, and an ML-DSA server signature, with rules that let the adversary compromise either KEM or reveal the signing key; it proves executability, server authentication, and—the headline lemma—*hybrid robustness*: under an honest server identity the session key remains secret unless *both* KEMs are broken. A ProVerif model establishes the analogous queries under a different abstraction (the classical leg as a second idealized KEM), giving an independent cross-check. All three formal artifacts are configured as CI gates (EasyCrypt `make check`, Tamarin, ProVerif), with checker and solver versions pinned in the artifact.

VI. Production Integration and Evaluation

A. Microbenchmarks

Table IV reports primitive and hybrid costs (Criterion, point estimates, on an Apple-Silicon arm64 host; the artifact reproduces them and an x86-64 run). Two findings matter for deployment. First, *the post-quantum leg is not the bottleneck*: ML-KEM-768 encapsulation is 11.0 μs, roughly 8× cheaper than X25519 encapsulation (89.6 μs, an ephemeral keygen plus a Diffie–Hellman), so the classical half dominates hybrid CPU—a useful corrective to the assumption that PQ is the expensive part. Second, *the hash-everything combiner is cheap*: ContextBound costs only +6.2 μs over the byte-exact X-Wing combiner CompatXWing (encapsulation 108.4 vs. 102.2 μs; decapsulation 65.4 vs. 59.2 μs)—the price of binding the full transcript is sub-10 μs, well under the ∼190 μs X25519 cost and far below any network RTT.

TABLE II
THE SIX ORTHOGONAL VERIFICATION METHODS.

Method	Reference / oracle	Independence	What it catches
Fixed KATs	X-Wing draft vectors; RFC 7748	published vectors	combiner and X25519 spec-conformance regressions
NIST ACVP	*usnistgov* ACVP, full FIPS family	NIST ground truth	primitive non-conformance (FIPS 203/204/205)
Multi-backend differential	RustCrypto `ml-kem`/`ml-dsa`; orion X25519	independent re-implementations	implementation bugs (output must stay byte-identical)
EasyCrypt proof	machine-checked; CR(SHA3)	formal (computational)	binding-design flaws in the combiner
Cross-platform byte-identity	shared test vector	5 faces × substrates	per-substrate divergence (same K everywhere)
Generative property tests	proptest: 7 properties × 128 cases	randomized	edge-case violations (injectivity, domain separation, K-CTX)

TABLE III
CROSS-SUBSTRATE COVERAGE (✓ VERIFIED, CI OR LOCAL).

(a) ISA targets (Rust core; the byte-level claim)

ISA	runner	build	byte-id. K	binary-CT
x86-64	Linux, Win-MSVC	✓	✓ (run)	✓ Memcheck
aarch64	Linux, macOS	✓	✓ (run)	✓ Memcheck
wasm32	node	✓	✓ (run)	source-CT
riscv64	Linux (cross)	✓	by constr.†	source-CT
thumbv7em	bare-metal	✓	by constr.†	n/a

† cross-compiled, not run natively; identity follows from FIPS-deterministic encodings + identical source.

(b) language face × OS (shared vector → same K)

Face	Linux	macOS	Windows
Rust	✓	✓	✓
C ABI	✓	✓	✓ (MSVC)
WASM (node)	✓	✓	✓
Swift	–	✓	–
Kotlin (JVM/FFM)	✓	✓	–

TABLE IV
PRIMITIVE AND HYBRID COST (μS) AND KEY/CIPHERTEXT SIZES (BYTES). ARM64 HOST.

	time (μs)			size (B)		
Scheme	keygen	encaps	decaps	ek	dk	ct
ML-KEM-512	6.6	6.1	6.8	800	1632	768
ML-KEM-768	11.3	11.0	11.8	1184	2400	1088
ML-KEM-1024	18.5	16.5	16.7	1568	3168	1568
X25519	44.1	89.6	45.0	32	32	32
Hybrid, ContextBound	—	108.4	65.4	1216	—	1120
Hybrid, CompatXWing	—	102.2	59.2	1216	—	1120

B. Handshake latency under emulated networks

We expose ContextBound (and CompatXWing) as a TLS 1.3 key-exchange group via a rustls [32] `CryptoProvider`, using private-use group codes; a real loopback TLS 1.3 handshake completes over this provider. We evaluate handshake *time-to-session* (TCP connect plus full TLS 1.3 handshake) over a real loopback socket shaped by `tc netem`, comparing a classical X25519 baseline against the two PQ/T profiles. Figure 5 reports the result: the three curves coincide because latency is RTT-dominated, and the PQ/T overhead is a *fixed* cost ($\sim$190 μs of ML-KEM CPU plus the wire bytes of Fig. 6). At realistic RTT this is negligible: at RTT $\geq$ 20 ms the measured overhead is *within run-to-run noise*—its sign varies across repeated runs (Fig. 5b, error bars span two repetitions)—while the only reproducible cost is the $\sim$190 μs of CPU visible at RTT 0. The hybrid's $+2.27$ KB (one ML-KEM-768 keyshare each way) fits within the existing TLS flights and triggers no additional round-trip. The two combiner profiles are indistinguishable: the stronger binding of the hash-everything combiner is free in latency. We report these measurements honestly as obtained on a virtualized host; a turnkey script reproduces the protocol on quiesced bare-metal hardware for the camera-ready (the RTT-dominated points are already stable across repetitions).

Under jitter and loss. Adding 3 ms of jitter leaves the picture unchanged (all three groups ≈ 41 ms median, ≈ 50 ms p99.9). Under 1–3% packet loss the median is likewise unaffected—loss is rare relative to a handshake—but the *tail* jumps to $\sim$1 s, dominated by a TCP retransmission timeout per lost packet rather than by crypto or wire size; this RTO-driven tail falls on the classical and PQ/T handshakes alike, and at our sample size the loss-event noise precludes a precise per-group tail comparison. The honest takeaway is that the hybrid's extra keyshare does not *systematically* worsen loss resilience: the median is byte-insensitive and the tail is governed by transport, not by the combiner.

VII. DISCUSSION

When to use which profile: `CompatXWing` exists for interoperability: it is byte-exact X-Wing and should be selected when talking to an X-Wing peer. `ContextBound` is the default for first-party deployments and is preferable whenever (i) the application has a meaningful context (a protocol/role/version label, a channel binding) that the session key should commit to; or (ii) the deployment cannot guarantee that its component KEM will only ever be instantiated in the seed-dk form—in which case binding pk_{pq} and ct_{pq} directly removes the latent serialization dependence of Section V. The cost of that robustness, measured in Section VI, is sub-10 μs and invisible end-to-end.

The assurance philosophy: Our position is deliberately modest about cryptographic novelty and deliberately ambitious

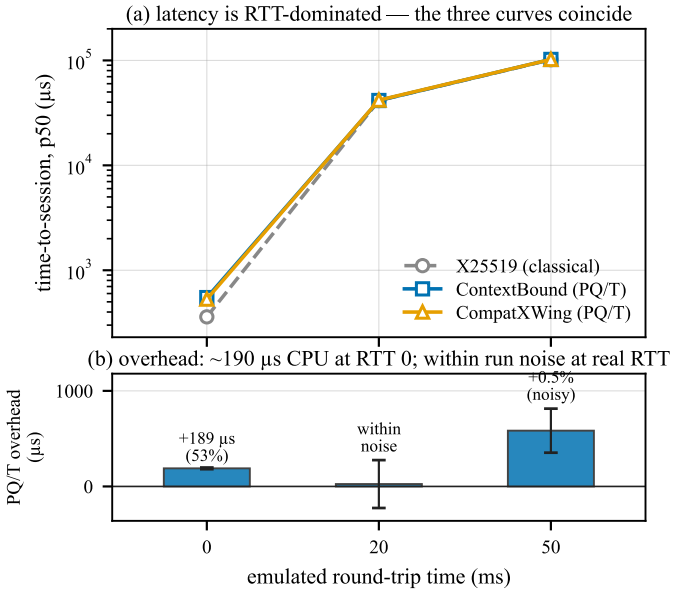

Fig. 5. Time-to-session under real `tc netem` (mean of two repetitions). (a) latency is RTT-dominated—the three curves coincide. (b) the overhead is ∼190 μs of CPU at RTT 0; at realistic RTT it is within run-to-run noise (error bars span the two repetitions and cross zero at RTT 20 ms).

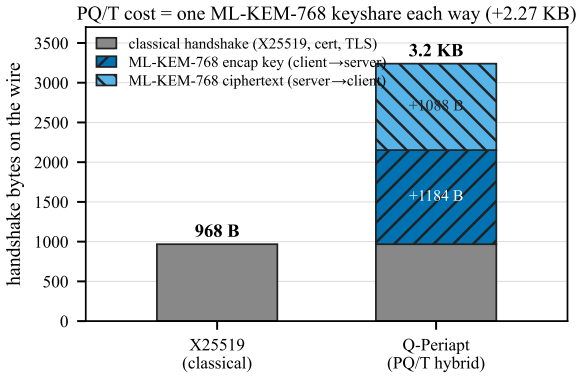

Fig. 6. Handshake wire budget. The PQ/T cost is one ML-KEM-768 keyshare each way (+2.27 KB), which fits inside the existing flights.

about *dependability*: a hybrid is trustworthy only to the extent that its weakest realization is. The recurring failures of the last two years—KyberSlash, PQXDH, the ML-KEM malicious-key binding gaps—were not failures of the underlying hard problems but of *realizations*: a compiled branch, a serialization format, a missing transcript field. The contribution of this paper is a discipline that makes those realization-level properties *checkable and reproducible*: a binding result whose proven object is carried byte-for-byte across the substrate, a constant-time probe that discriminates clean from leaky, and a regression oracle for the key-format coupling. We expect the individual instruments to be reused independently of ContextBound.

Generality: Nothing in the approach is specific to ML-KEM + X25519. The combiner and its proof are generic over the component KEMs; the trait-injected backends already include the full ML-KEM and ML-DSA families and an (unaudited) HQC backend used here as the CT discriminator's leaky reference. As new PQ primitives are standardized, the same proof-to-byte discipline applies unchanged.

VIII. Related Work

KEM binding and committing security: The systematic study of KEM binding is recent. CDM [2] introduced the X-BIND-P-Q lattice, the HON/LEAK/MAL adversary hierarchy, and its monotonicity, and showed where standard KEMs sit; Schmieg [3] then proved that ML-KEM in its FIPS-203 *expanded* decapsulation-key form is neither MAL-BIND-K-CT nor MAL-BIND-K-PK, with seed-form keys as the fix, and [21] analyzes implicitly-rejecting KEMs more broadly. For *combined* KEMs, Chempat [4] establishes the dichotomy we rely on—a hash-everything combiner is binding from collision-resistance alone, while a lean combiner inherits the binding of whatever component it omits; ContextBound is an instance of the former, which is exactly why we attribute the binding *fact* rather than claim it. At the symmetric layer, the same idea appears as context-committing AEAD (CMT) [7], of which our context axis is the KEM-output-layer analogue; the role of public contexts specifically in hybrid KEMs is studied in [8]. Our addition is not a new notion but the conjunction of Table I: a mechanized, assumption-minimal binding result whose proven object is byte-identical across the deployment surface.

Side-channel assurance: KyberSlash [6] showed that secret-dependent timings in widely-used Kyber/ML-KEM implementations were exploitable, underscoring that source-level constant-time arguments do not by themselves settle the question for shipped binaries. The dudect/ctgrind discipline of marking secrets and checking for secret-dependent control flow [28], [30], realized over Valgrind/Memcheck [29], is the basis of our binary-CT gate. We build on, rather than re-do, the source-level guarantees of HACL* [22] and the libcrux ML-KEM verification [23]: our probe *corroborates* them dynamically and, crucially, *discriminates* a verified primitive from an unverified one in the same framework.

Hybrids in deployment: Hybrid key exchange is being standardized and shipped: the IETF hybrid design [17] and the `X25519MLKEM768` group [18] for TLS 1.3 [16], [15], the X-Wing KEM [5], Signal's PQXDH (formally analyzed in [19]), and Apple's PQ3 [20]. These are the baselines we position against; our combiner interoperates with X-Wing (the `CompatXWing` profile) and our rustls integration targets the same TLS 1.3 deployment path.

Formal verification of PQ cryptography: formosa-mlkem [24] delivers a verified ML-KEM *implementation* (Jasmin/EasyCrypt) with IND-CCA and constant-time guarantees—an axis (specification-to-binary) we explicitly do not hold (Table I). The SandboxAQ EasyCrypt-KEMs library mechanizes LEAK-BIND-K-PK for ML-KEM. Our EasyCrypt [25] development is complementary: it targets the *hybrid combiner's* binding (the commitment-style MAL Figure-6 instantiation), and is cross-checked by independent symbolic models in Tamarin [26] and ProVerif [27], all CI-gated.

IX. Limitations

The mechanized binding result is over an *abstract* Decaps: it is exactly the "zero KEM assumption" that makes the result

general, but it carries no FIPS-203 linkage, the shared-secret fields are inert in the argument, and CR(SHA3) and IND-CCA2 remain assumptions argued respectively in the model and on paper; there is no specification-to-implementation linkage proof. Binary-level constant-time tooling is mature only on x86-64 and aarch64; riscv64 and wasm32 are covered at source level plus upstream attestation. ACVP byte-identity is conformance, not CMVP certification. The suite is research-grade and has not had an independent third-party audit. The netem evaluation was obtained on a virtualized host; bare-metal reproduction is provided as a script.

X. Conclusion

A PQ/T hybrid is only as safe as its weakest realization. We closed that gap with an assumption-minimal, machine-checked binding result whose proven object is byte-identical across a heterogeneous deployment surface, guarded by reproducible gates (configured in CI)—a source→binary constant-time discriminator, a binding key-format coupling demonstration, and multi-prover symbolic models—and validated on a production TLS 1.3 path. The genuinely new element is the conjunction, not any single fact; we have been explicit about the boundary throughout.

References

[1] National Institute of Standards and Technology, "FIPS 203: Module-Lattice-Based Key-Encapsulation Mechanism Standard," 2024.

[2] C. Cremers, A. Dax, and N. Medinger, "Keeping Up with the KEMs: Binding Security of Key Encapsulation Mechanisms," *ACM CCS*, 2024. (ePrint 2023/1933.)

[3] S. Schmieg, "Unbindable Kemmy Schmidt: ML-KEM is neither MAL-BIND-K-CT nor MAL-BIND-K-PK," IACR ePrint 2024/523, 2024.

[4] J. Güneysu, K. Hövelmanns, *et al.*, "Binding Security of Combined KEMs / Chempat," IACR ePrint 2025/1416, 2025.

[5] D. Connolly, P. Schwabe, and B. Westerbaan, "X-Wing: The Hybrid KEM You've Been Looking For," draft-connolly-cfrg-xwing-kem (CFRG), 2024.

[6] D. J. Bernstein, K. Bhargavan, *et al.*, "KyberSlash: Exploiting Secret-Dependent Division Timings in Kyber Implementations," *IACR TCHES*, 2025.

[7] M. Bellare and V. T. Hoang, "Efficient Schemes for Committing Authenticated Encryption," *EUROCRYPT*, 2022.

[8] "On the Role of Public Contexts in Hybrid KEMs," IACR ePrint 2026/140.

[9] National Institute of Standards and Technology, "FIPS 204: Module-Lattice-Based Digital Signature Standard," 2024.

[10] National Institute of Standards and Technology, "FIPS 205: Stateless Hash-Based Digital Signature Standard," 2024.

[11] J. Bos *et al.*, "CRYSTALS-Kyber: A CCA-Secure Module-Lattice-Based KEM," *IEEE EuroS&P*, 2018.

[12] G. Alagic *et al.*, "Status Report on the Fourth Round of the NIST Post-Quantum Cryptography Standardization Process," NIST IR 8545, 2025.

[13] E. Fujisaki and T. Okamoto, "Secure Integration of Asymmetric and Symmetric Encryption Schemes," *CRYPTO*, 1999.

[14] D. Hofheinz, K. Hövelmanns, and E. Kiltz, "A Modular Analysis of the Fujisaki–Okamoto Transformation," *TCC*, 2017.

[15] A. Langley, M. Hamburg, and S. Turner, "Elliptic Curves for Security," RFC 7748, IETF, 2016.

[16] E. Rescorla, "The Transport Layer Security (TLS) Protocol Version 1.3," RFC 8446, IETF, 2018.

[17] D. Stebila, S. Fluhrer, and S. Gueron, "Hybrid Key Exchange in TLS 1.3," draft-ietf-tls-hybrid-design, IETF, 2023.

[18] K. Kwiatkowski *et al.*, "Post-Quantum Key Exchange `X25519MLKEM768` for TLS 1.3," draft-kwiatkowski-tls-ecdhe-mlkem, IETF, 2024.

[19] K. Bhargavan, C. Jacomme, F. Kiefer, and R. Schmidt, "Formal Verification of the PQXDH Post-Quantum Key Agreement Protocol," *IEEE S&P*, 2024.

[20] Apple Security Engineering, "iMessage with PQ3: The New State of the Art in Quantum-Secure Messaging," Apple Security Research, 2024.

[21] K. Hövelmanns *et al.*, "Binding Security of Implicitly-Rejecting KEMs and Application to BIKE and HQC," IACR ePrint 2024/1233, 2024.

[22] J.-K. Zinzindohoué, K. Bhargavan, J. Protzenko, and B. Beurdouche, "HACL*: A Verified Modern Cryptographic Library," *ACM CCS*, 2017.

[23] Cryspen, "Verified ML-KEM (Kyber) in `libcrux`," 2024.

[24] J. B. Almeida *et al.*, "Formally Verifying Kyber: Episode V," *CRYPTO*, 2024.

[25] G. Barthe, B. Grégoire, S. Heraud, and S. Z. Béguelin, "Computer-Aided Security Proofs for the Working Cryptographer," *CRYPTO*, 2011.

[26] S. Meier, B. Schmidt, C. Cremers, and D. Basin, "The TAMARIN Prover for the Symbolic Analysis of Security Protocols," *CAV*, 2013.

[27] B. Blanchet, "Modeling and Verifying Security Protocols with the Applied Pi Calculus and ProVerif," *Found. Trends Priv. Secur.*, 2016.

[28] O. Reparaz, J. Balasch, and I. Verbauwhede, "Dude, is My Code Constant Time?" *DATE*, 2017.

[29] N. Nethercote and J. Seward, "Valgrind: A Framework for Heavyweight Dynamic Binary Instrumentation," *ACM PLDI*, 2007.

[30] A. Langley, "ctgrind: Checking that Functions are Constant Time with Valgrind," 2010.

[31] National Institute of Standards and Technology, "Automated Cryptographic Validation Protocol (ACVP)," https://github.com/usnistgov/ACVP, 2024.

[32] J. Birr-Pixton *et al.*, "rustls: A Modern TLS Library in Rust," https://github.com/rustls/rustls, 2024.

Appendix A
Artifact Availability

The complete suite—Rust core and backends, the five language faces, the EasyCrypt/Tamarin/ProVerif developments, the constant-time and netem harnesses, and the figure-generation scripts—is open source at https://github.com/billlza/q-periapt. All numbers in this paper are reproducible from that repository; the bare-metal netem protocol is provided as a script.

Appendix B
The Machine-Checked Kernel in Detail

Table V inventories the EasyCrypt development. Of the named results, all are fully discharged (no `admit`/`conjecture`); the only assumptions are the two elementary `be8_*` facts about the 8-byte length field and the modeling of H as collision-resistant. We validate that the proof is non-vacuous by an *adversarial negative control*: for each binding theorem we mechanically delete a single hint from its closing tactic and re-run the checker, confirming the proof then *fails* (`cannot prove goal`). Deleting `encode_inj` breaks every binding theorem; deleting the disagreement lemma (`ct_neq_fields_neq`, etc.) breaks the corresponding axis; and—in the explicit-rejection game—deleting the $K \neq \perp$ conjunct from the predicate makes the theorem unprovable, since two mutually-rejecting executions would otherwise "win" without inducing a hash collision. This last control is what distinguishes a faithful Figure-6 instantiation from a vacuous one.

Appendix C
Extended Evaluation

Table VI gives the microbenchmarks of Table IV with Criterion's 95% confidence intervals (arm64 host). Table VII

TABLE V
EASYCRYPT LEMMA INVENTORY (BINDINGVIACR.EC). ALL PROVED; ASSUMPTIONS: BE8_*, CR(H).

Lemma	Statement (informal)
lp_cat_inj	length-prefixed fields are self-delimiting
encode_inj	the field encoding is injective
bind_le_cr	generic transcript-collision $\leq$ CR(H)
bind_le_cr_k{ct,pk,ctx}	instantiated projections (CT/PK/CTX)
malbind_k{ct,pk,ctx}_le_cr	KEM-aware, implicit rejection $\leq$ CR(H)
malbind_*_xrej_le_cr	full Fig. 6, explicit rejection, $K \neq \perp$

gives the jitter/loss data of Section VI: the median is loss-insensitive, while the tail is governed by a per-loss TCP retransmission timeout and is, at 150 samples, statistically indistinguishable across the three groups—hence we report it as "RTO-dominated, transport-bound" rather than as a per-group claim. Footprint: the stripped C-ABI library is 441 KB (entire suite), the WebAssembly module 249 KB.

TABLE VI
MICROBENCHMARKS, μS [95% CI]. ARM64 HOST, CRITERION.

Scheme	keygen	encaps	decaps
ML-KEM-512	6.6 [6.6,6.7]	6.1 [5.9,6.5]	6.8 [6.8,6.9]
ML-KEM-768	11.3 [11.2,11.4]	11.0 [10.9,11.0]	11.8 [11.7,11.8]
ML-KEM-1024	18.5 [18.1,19.2]	16.5 [15.8,17.7]	16.7 [16.5,16.8]
X25519	44.1 [43.8,44.4]	89.6 [89.2,90.0]	45.0 [44.7,45.3]
Hybrid, ContextBound	—	108.4 [107.9,108.9]	65.4 [65.1,65.7]
Hybrid, CompatXWing	—	102.2 [101.4,103.0]	59.2

TABLE VII
TIME-TO-SESSION UNDER JITTER/LOSS, P50/P99 (μS). THE P99 UNDER LOSS IS RTO-DOMINATED AND STATISTICALLY INDISTINGUISHABLE ACROSS GROUPS.

netem (RTT 20 ms +)	X25519	ContextBound	CompatXWing
+3 ms jitter	41.5k / 49.9k	41.6k / 50.1k	41.2k / 49.5k
+1% loss	41.1k / ∼1.08M	41.3k / (noisy)	41.4k / ∼1.09M
+3% loss	41.2k / ∼1.09M	41.6k / ∼1.09M	41.5k / ∼1.10M

APPENDIX D
REPRODUCTION

Every result is reproducible from the artifact. The EasyCrypt kernel: `make -C formal/easycrypt check`. The constant-time discriminator: `sh ctstats/scripts/ct-gap-probe.sh` (Linux+Valgrind; ML-KEM and HQC). The binding/key-format demonstration: `cargo test -p q-periapt-backends --test binding_keyformat_separation`. Microbenchmarks: `cargo bench -p q-periapt-backends`. The TLS netem evaluation: `paper/netem-camera-ready.sh` (bare metal). Cross-substrate byte-identity: the per-face vector tests; the symbolic models: `make` under `formal/{tamarin,proverif}`. The figures of this paper rebuild via `make -C paper/figures`.

ACKNOWLEDGMENT

The author thanks the maintainers of libcrux/HACL*, EasyCrypt, Tamarin, and ProVerif, whose tools this work builds upon. *[Funding/advisor acknowledgments to be added.]*